



Disciplined Inconsistency

with Consistency Types

Brandon Holt, James Bornholt, Irene Zhang,
Dan Ports, Mark Oskin, Luis Ceze

W UNIVERSITY *of* WASHINGTON

Building scalable & reliable applications is **difficult**.

The Dress That Broke the Internet, and the Woman Who Started It All

Rishi Iyengar @lyengarish | Feb. 27, 2015



Constantly changing network conditions, traffic spikes...

Caitlin McNeill had no idea she would start something that would be tweeted by Taylor Swift



On Thursday, Scottish musician Caitlin McNeill [posted](#) a picture of a dress on micro-blogging website Tumblr with the caption “guys please help me — is this dress white and gold, or blue and black?”

Building scalable & reliable applications is **difficult**.

Constantly changing network conditions, traffic spikes...

NEWSFEED VIRAL

The Dress That Broke the Internet, and the Woman Who Started It All

Rishi Iyengar @lyengarish | Feb. 27, 2015



Caitlin McNeill had no idea she would start something that would be tweeted by Taylor Swift



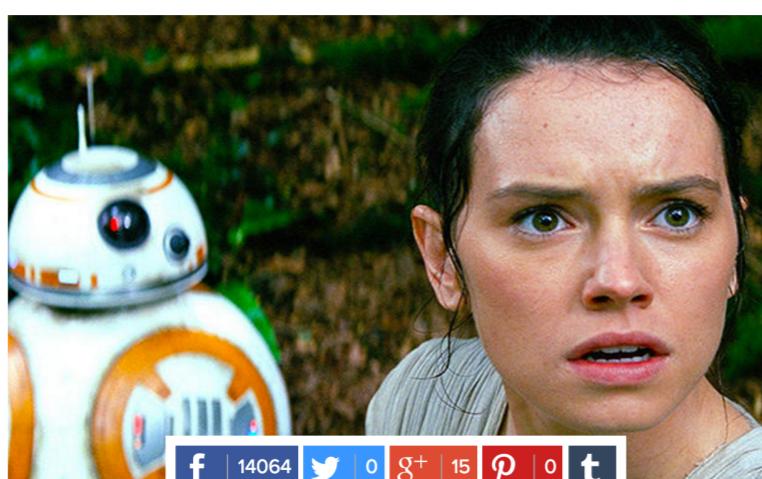
On Thursday, Scottish musician Caitlin McNeill posted a picture of a dress on micro-blogging website Tumblr with the caption “guys please help me — is this dress white and gold, or blue and black?”

Building scalable & reliable applications is difficult.

MOVIES | STAR WARS GALAXY

Star Wars: The Force Awakens presale crashes ticketing websites

BY OLIVER GETTELL · @OGETTELL



(Lucasfilm)

Top Stories

- Tom Hanks braces for impact in first 'Sully' trailer
- 'Lost in Space' reboot ordered by Netflix
- 'Star Trek' fan film guidelines defended by CBS official
- 'Ghostbusters' cast shows off new gadgets
- Emma Watson's Tina Turner ringtone interrupts interview
- What to Watch Wednesday: 'Barely Famous' is fully funny
- 'Night Of': Riz Ahmed held 'onto trauma for 6 months' during filming

Building scalable & reliable applications is **difficult.**

Programming languages can **help:**

- Prevent common mistakes *statically*
- Communicate application requirements

TICKETSLEUTH



MOVIES | STAR WARS GALAXY

Star Wars: The Force Awakens presale crashes ticket websites

BY OLIVER GETTELL · @OGETTELL

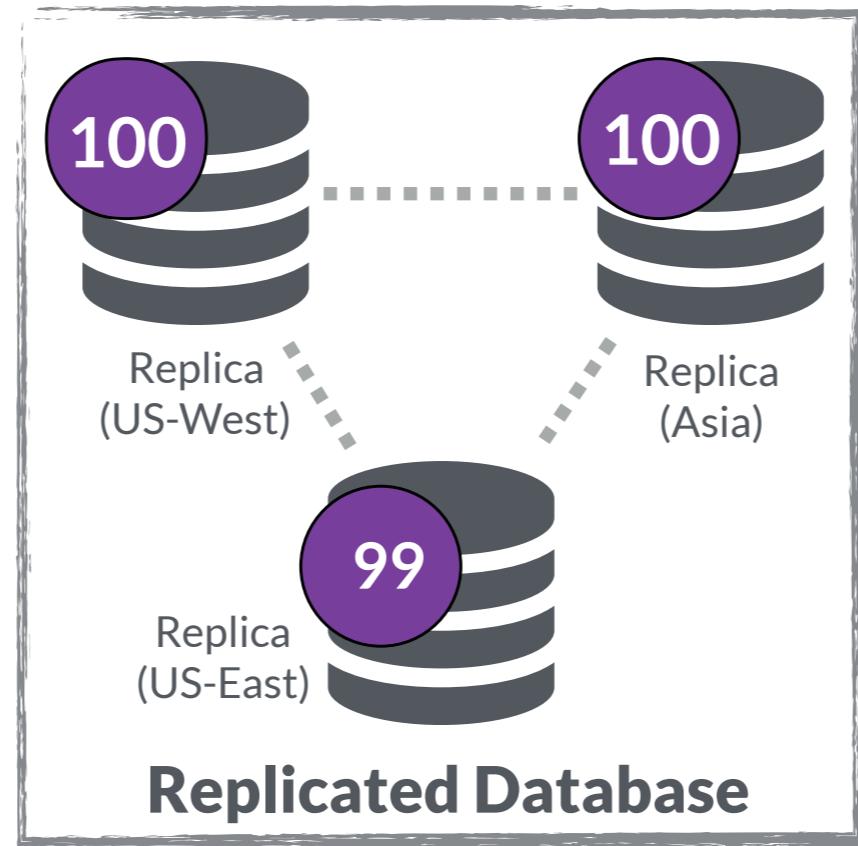
[f Share | 14064](#) [t Tweet | 0](#) [g+ | 15](#) [p | 0](#) [t](#)



Top Stories

- Tom Hanks braces for impact in first 'Sully' trailer
- 'Lost in Space' reboot ordered by Netflix
- 'Star Trek' fan film guidelines defended by CBS official
- 'Ghostbusters' cast shows off new gadgets
- Emma Watson's Tina Turner ringtone interrupts interview
- What to Watch Wednesday: 'Barely Famous' is fully funny
- 'Night Of': Riz Ahmed held 'onto trauma for 6 months' during filming

(Lucasfilm)



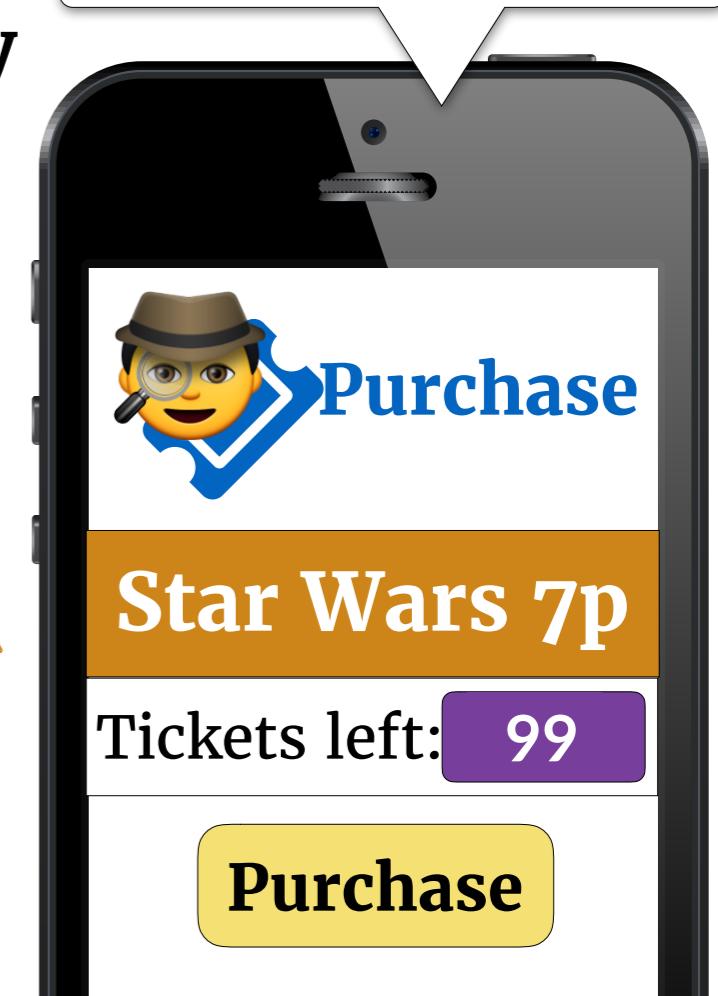
Performance constraint:

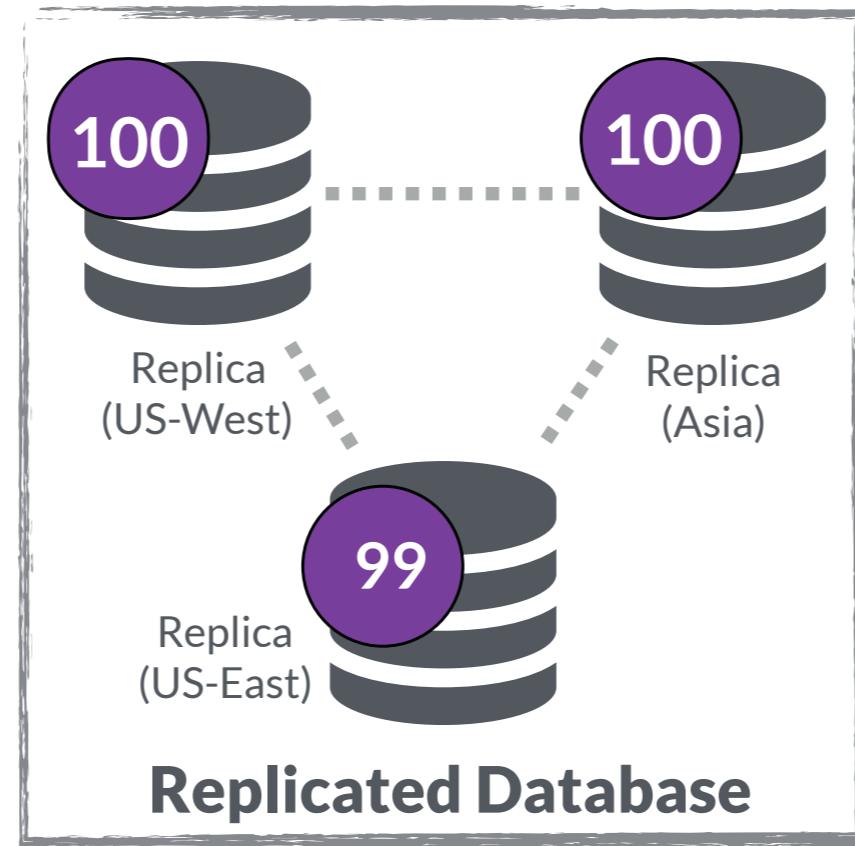
< 50 ms latency



Error tolerance:
count does not need
to be exact...

Hard constraint:
do not over-sell tickets





Performance constraint:

< 50 ms latency

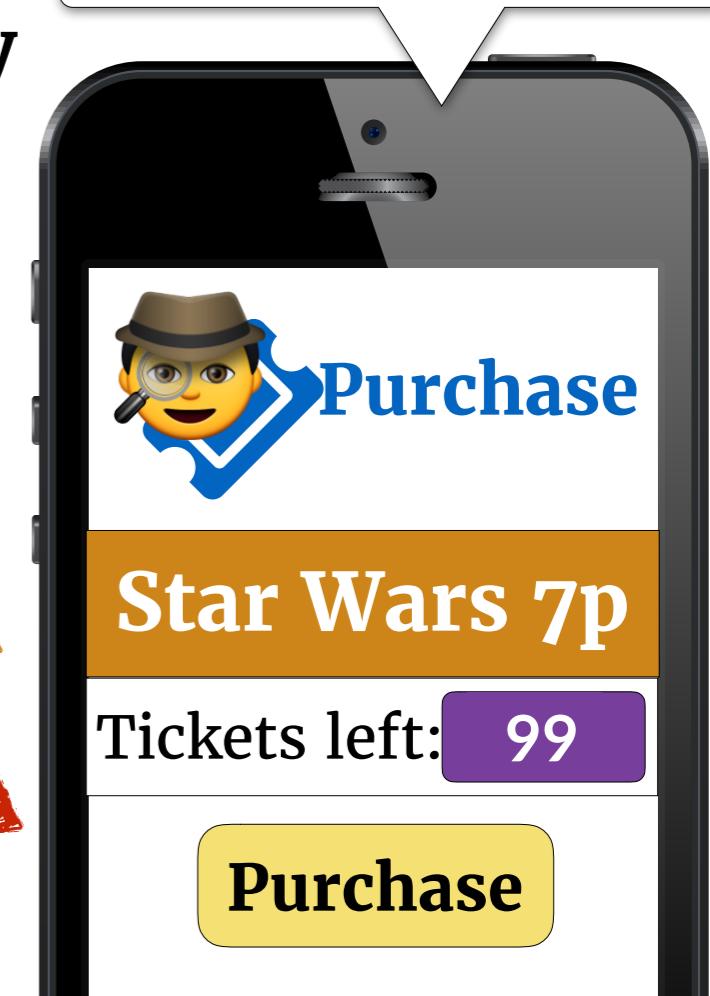


Strong consistency

Weak?

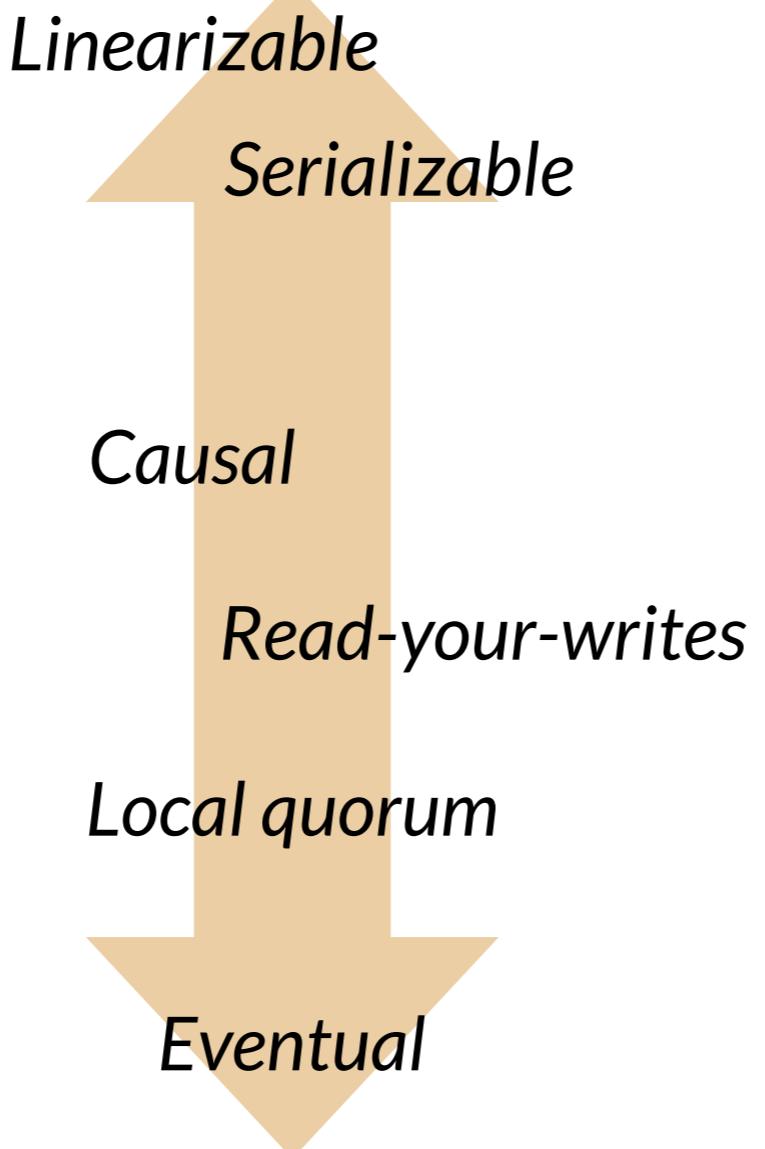
Error tolerance:
count does not need
to be exact...

Hard constraint:
do not over-sell tickets



Trading off consistency

Strong Consistency



easier to use,
slower

Weak Consistency

faster,
more error-prone



I
nconsistent
P
erformance-bound
A
pproximate

} Programming Model

Trading off consistency

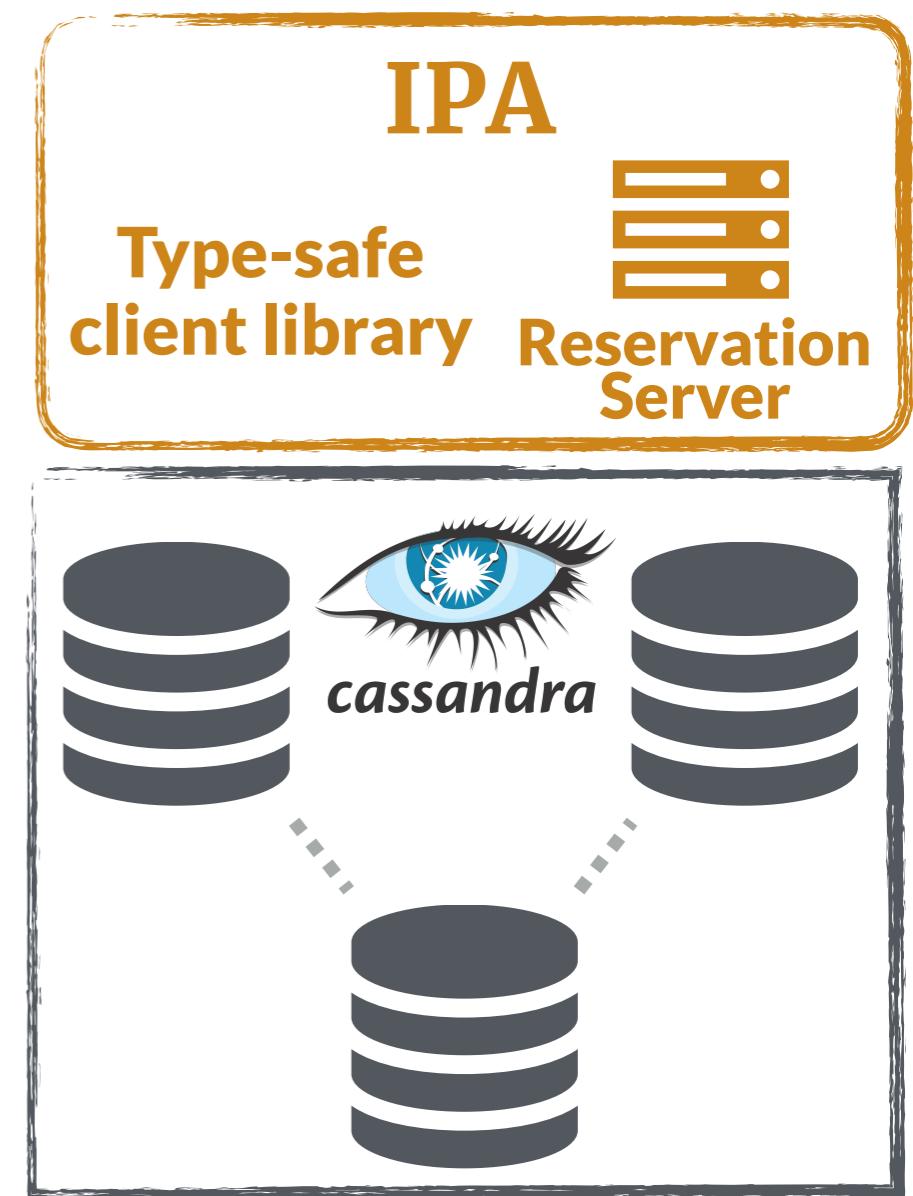
Consistency types

- Statically prevent consistency bugs

Performance & accuracy constraints

- Dynamically adjust consistency

Case study





I
nconsistent
P
erformance-bound
A
pproximate

} Programming Model

Trading off consistency

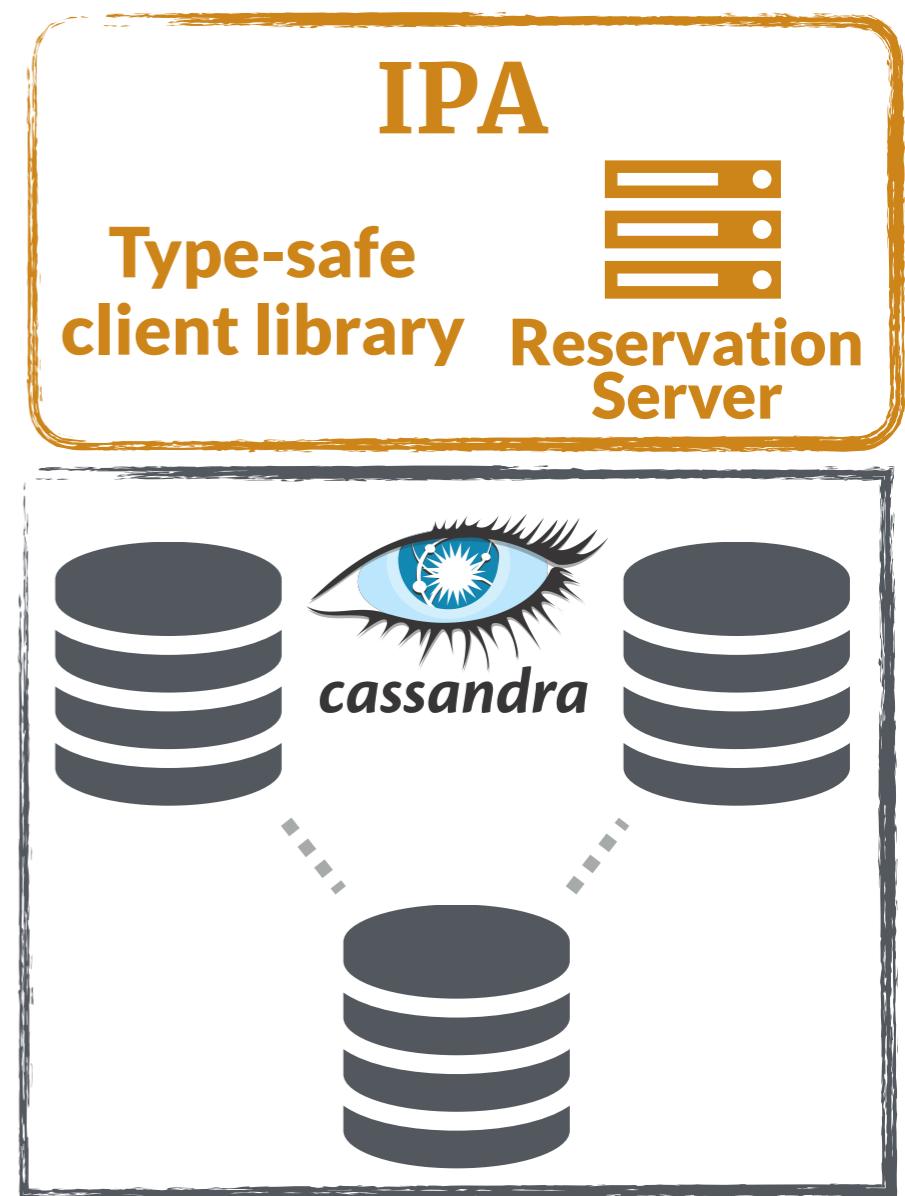
Consistency types

- Statically prevent consistency bugs

Performance & accuracy constraints

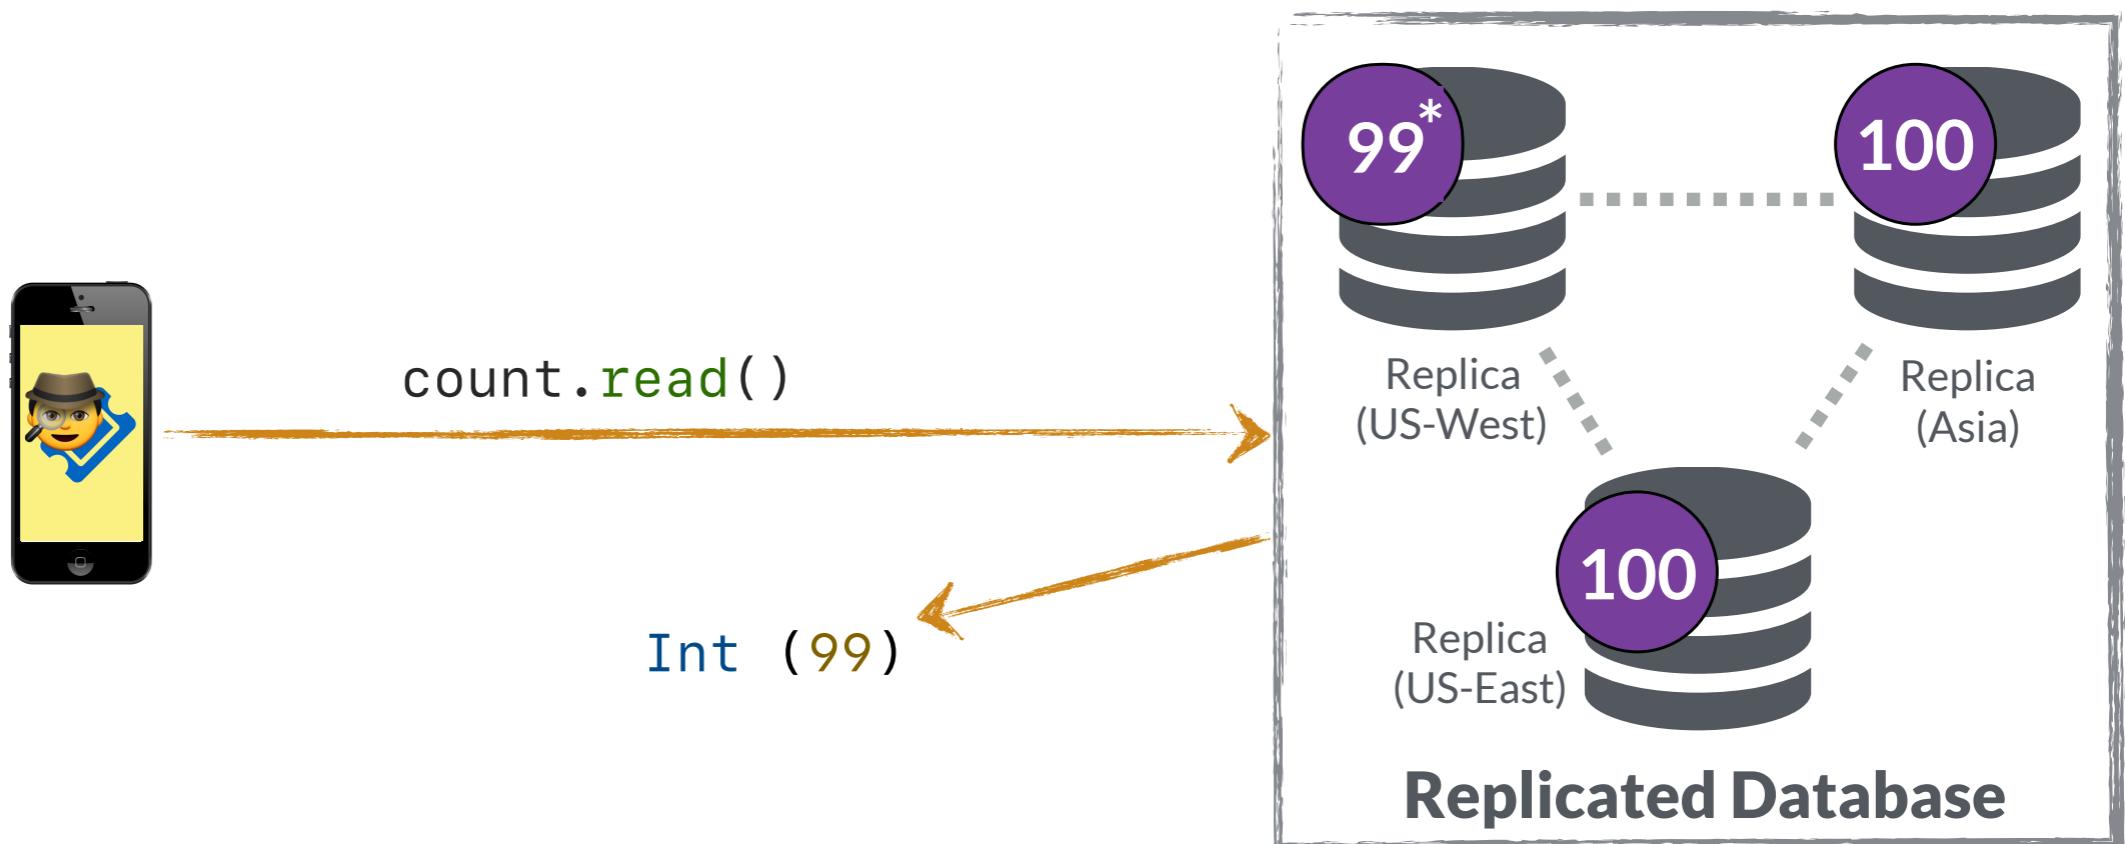
- Dynamically adjust consistency

Case study



Consistency Types

Every value returned from the datastore has a *consistency type*.

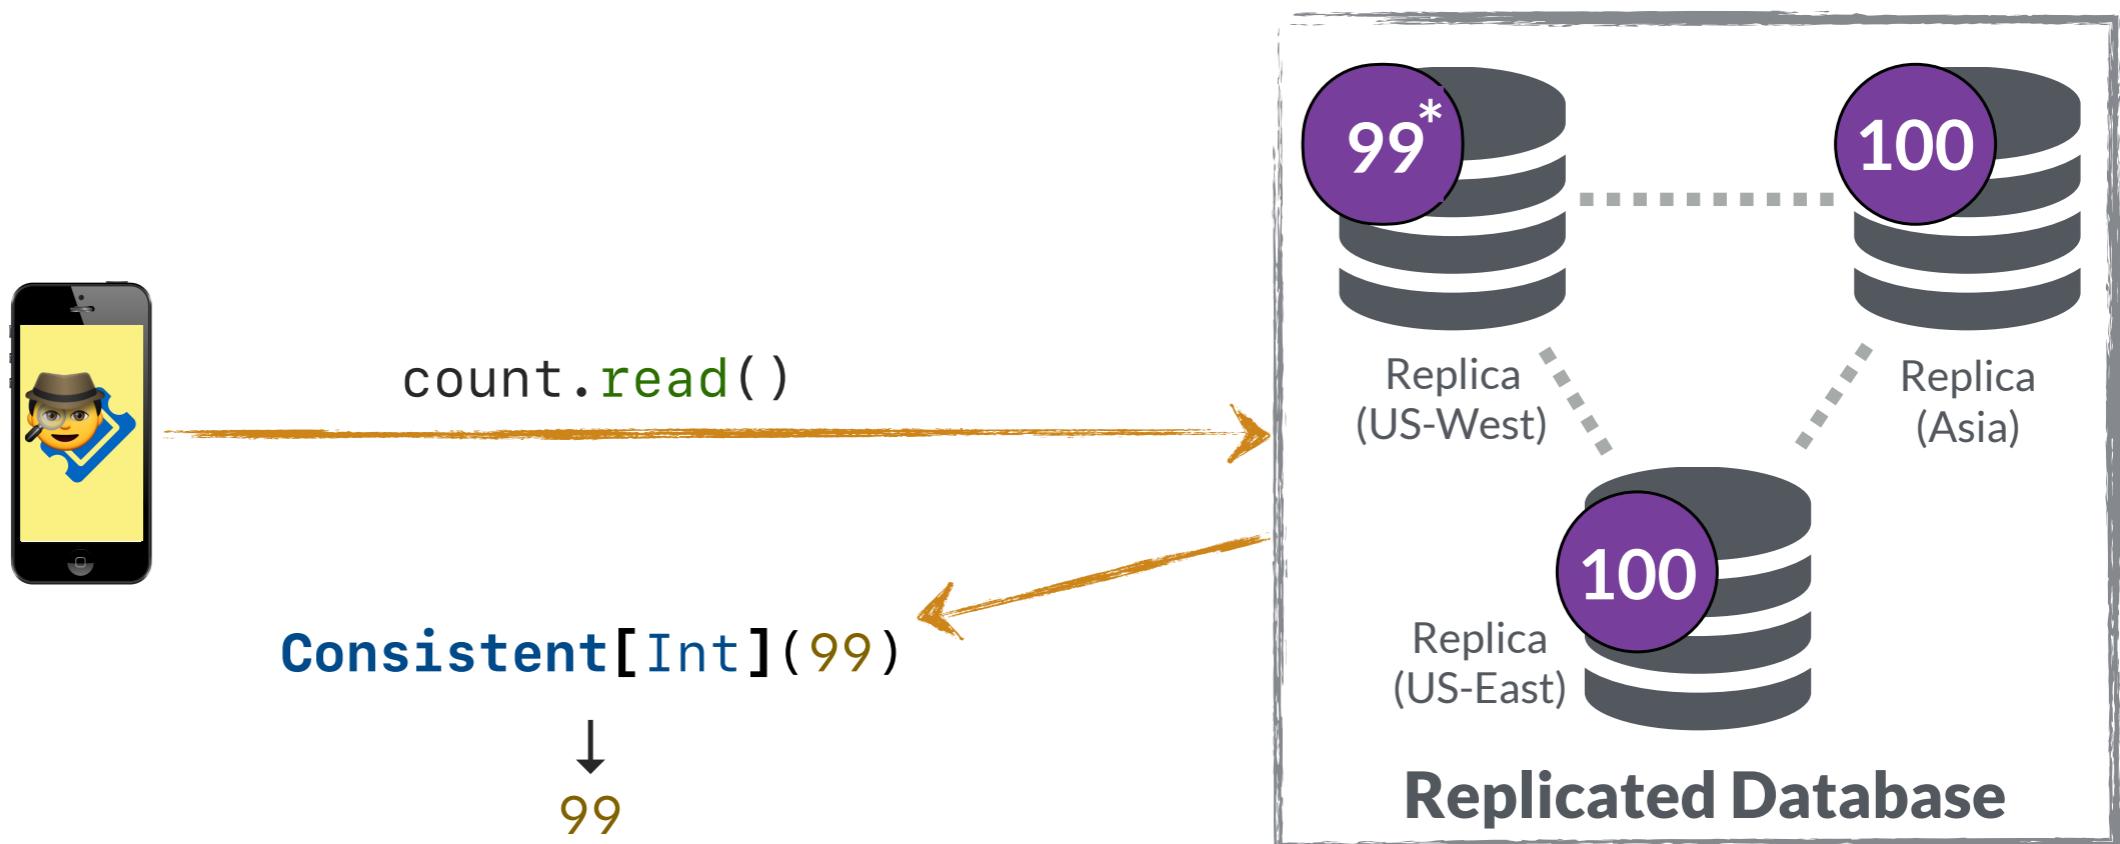


Consistency Types

Every value returned from the datastore has a *consistency type*.

Strong reads return **Consistent[T]**

- Can be implicitly converted to T (treated as a normal value)



Consistency Types

Every value returned from the datastore has a *consistency type*.

Strong reads return **Consistent[T]**

- Can be implicitly converted to T (treated as a normal value)

Weak reads return **Inconsistent[T]**

- Explicitly *endorse* to get the value:

```
endorse(Inconsistent(99)) → 99
```

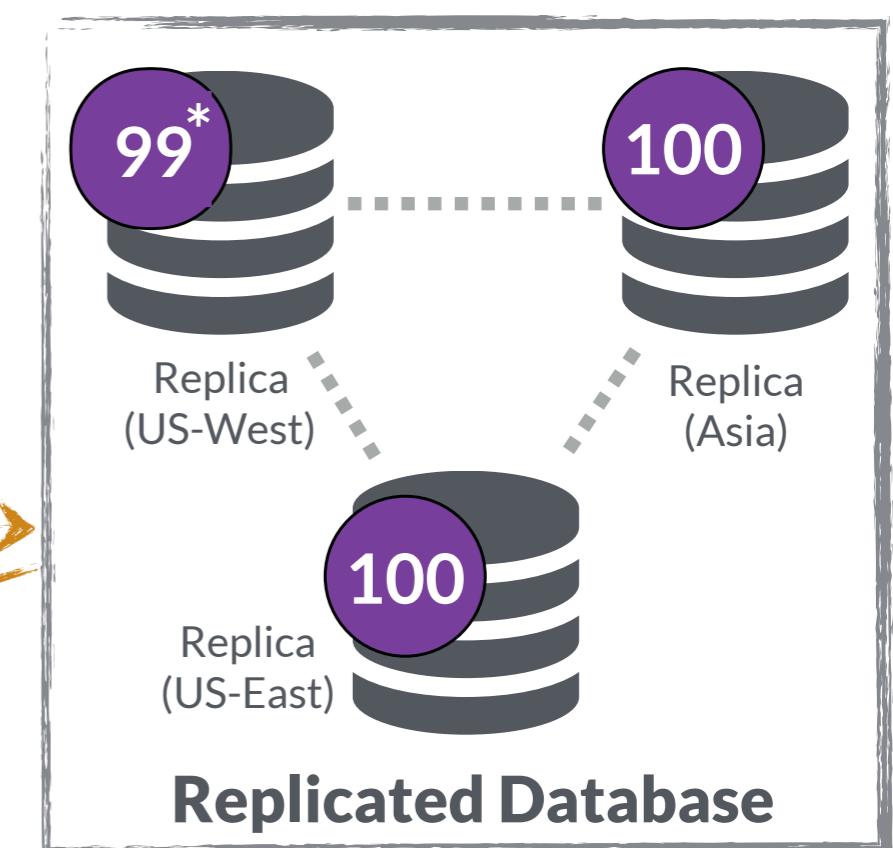
or with a case match:

```
case Inconsistent(x) ⇒ x:Int
```



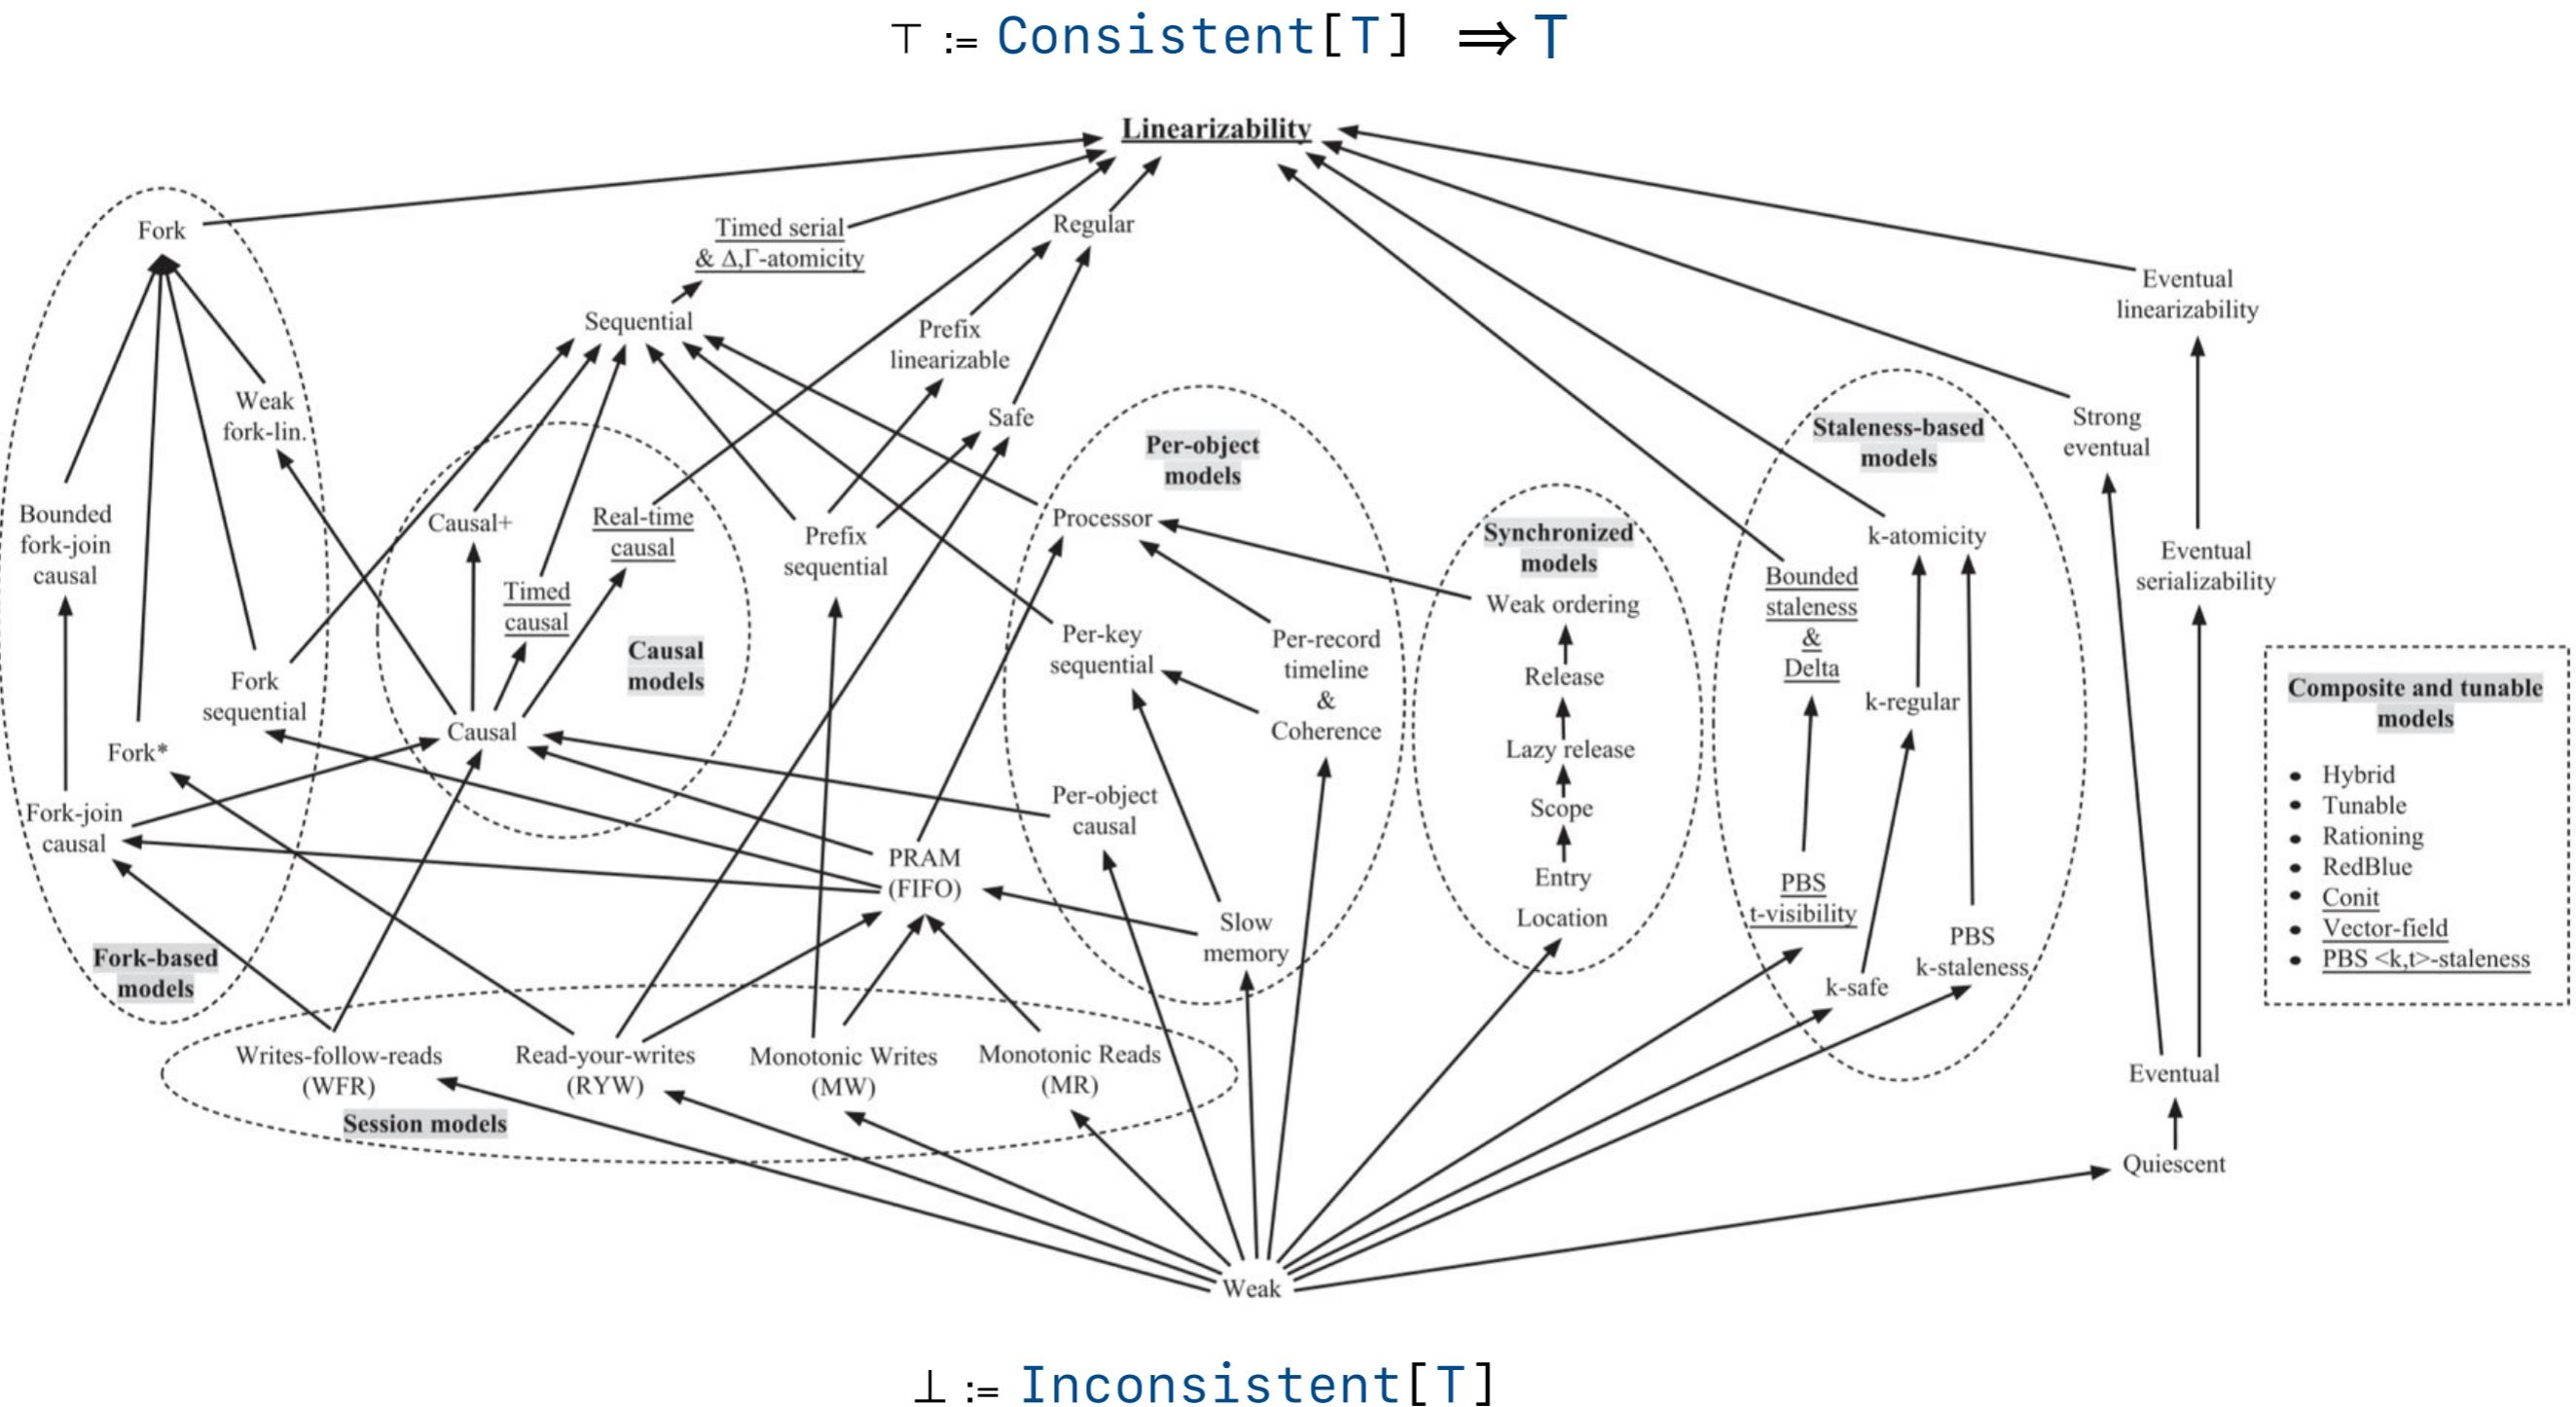
```
count.read_weak()
```

Inconsistent[Int](99)
or... **Inconsistent[Int](100)**



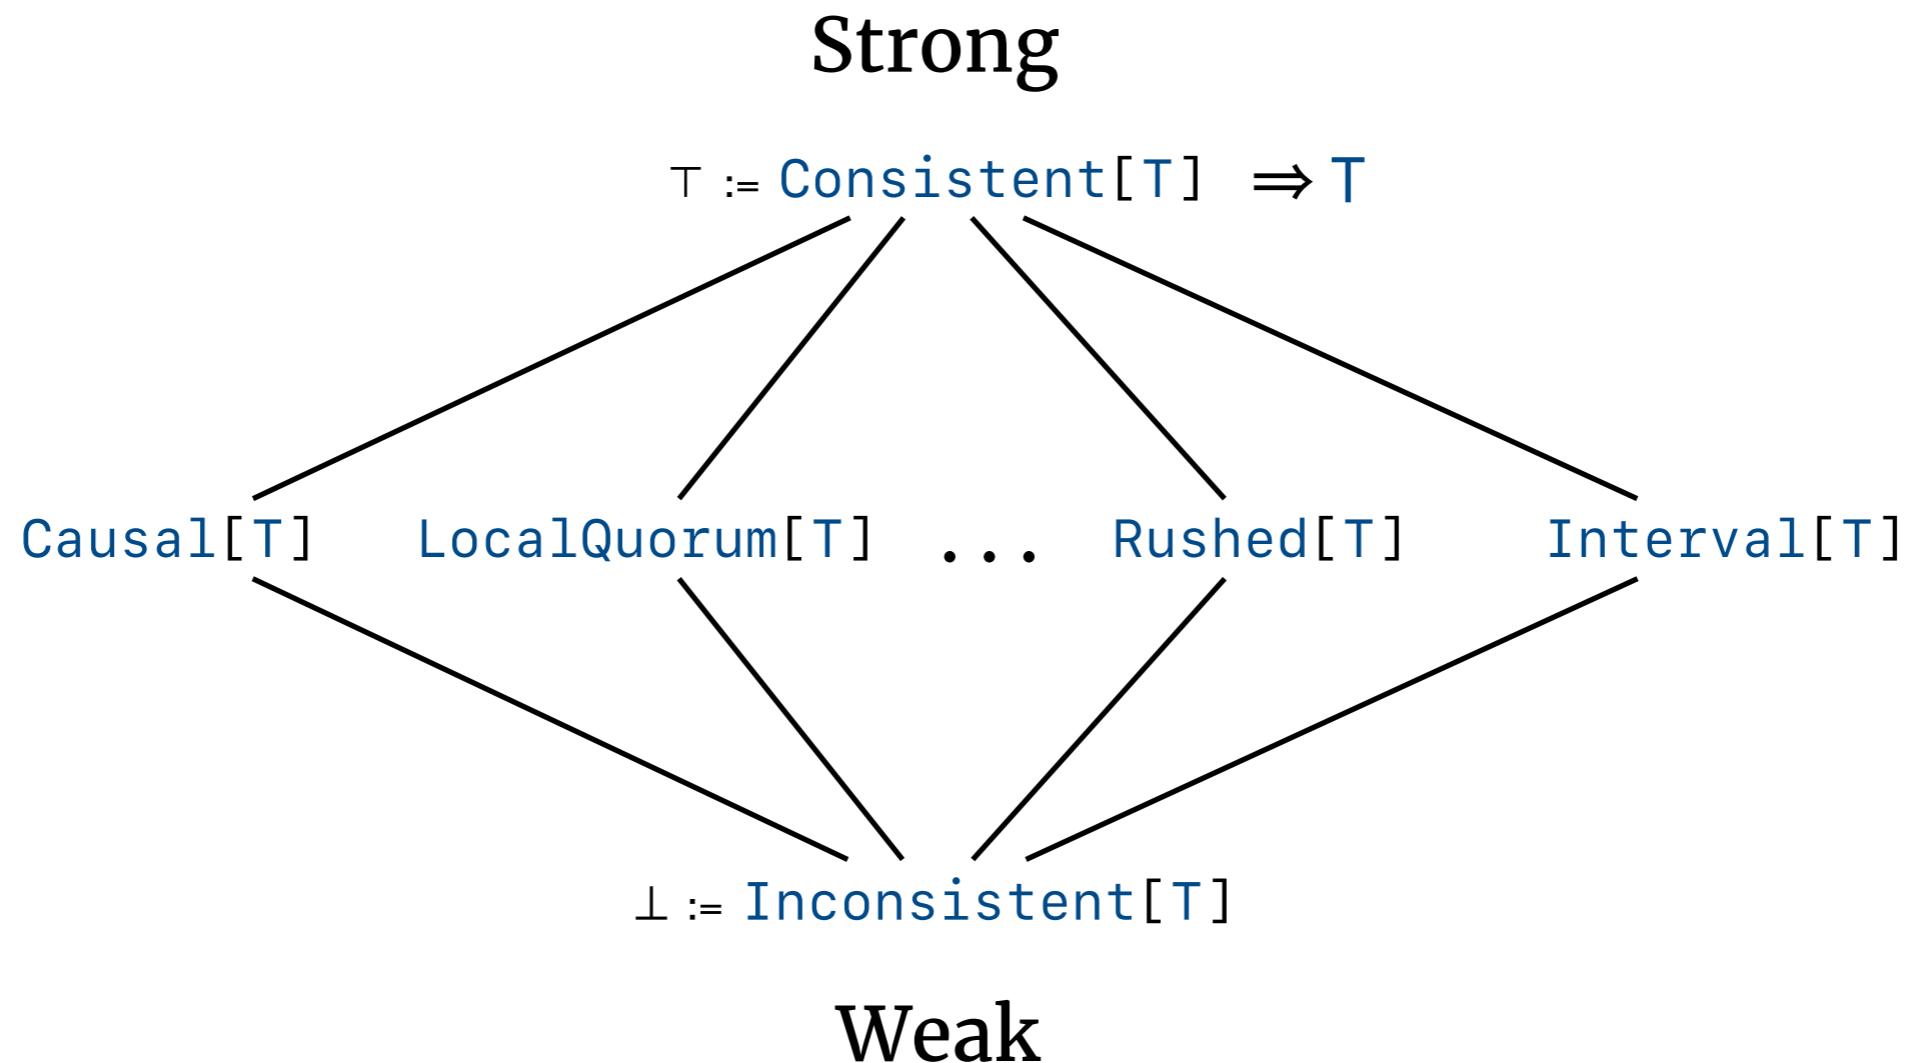
Consistency Types

Partially ordered, forming a lattice:



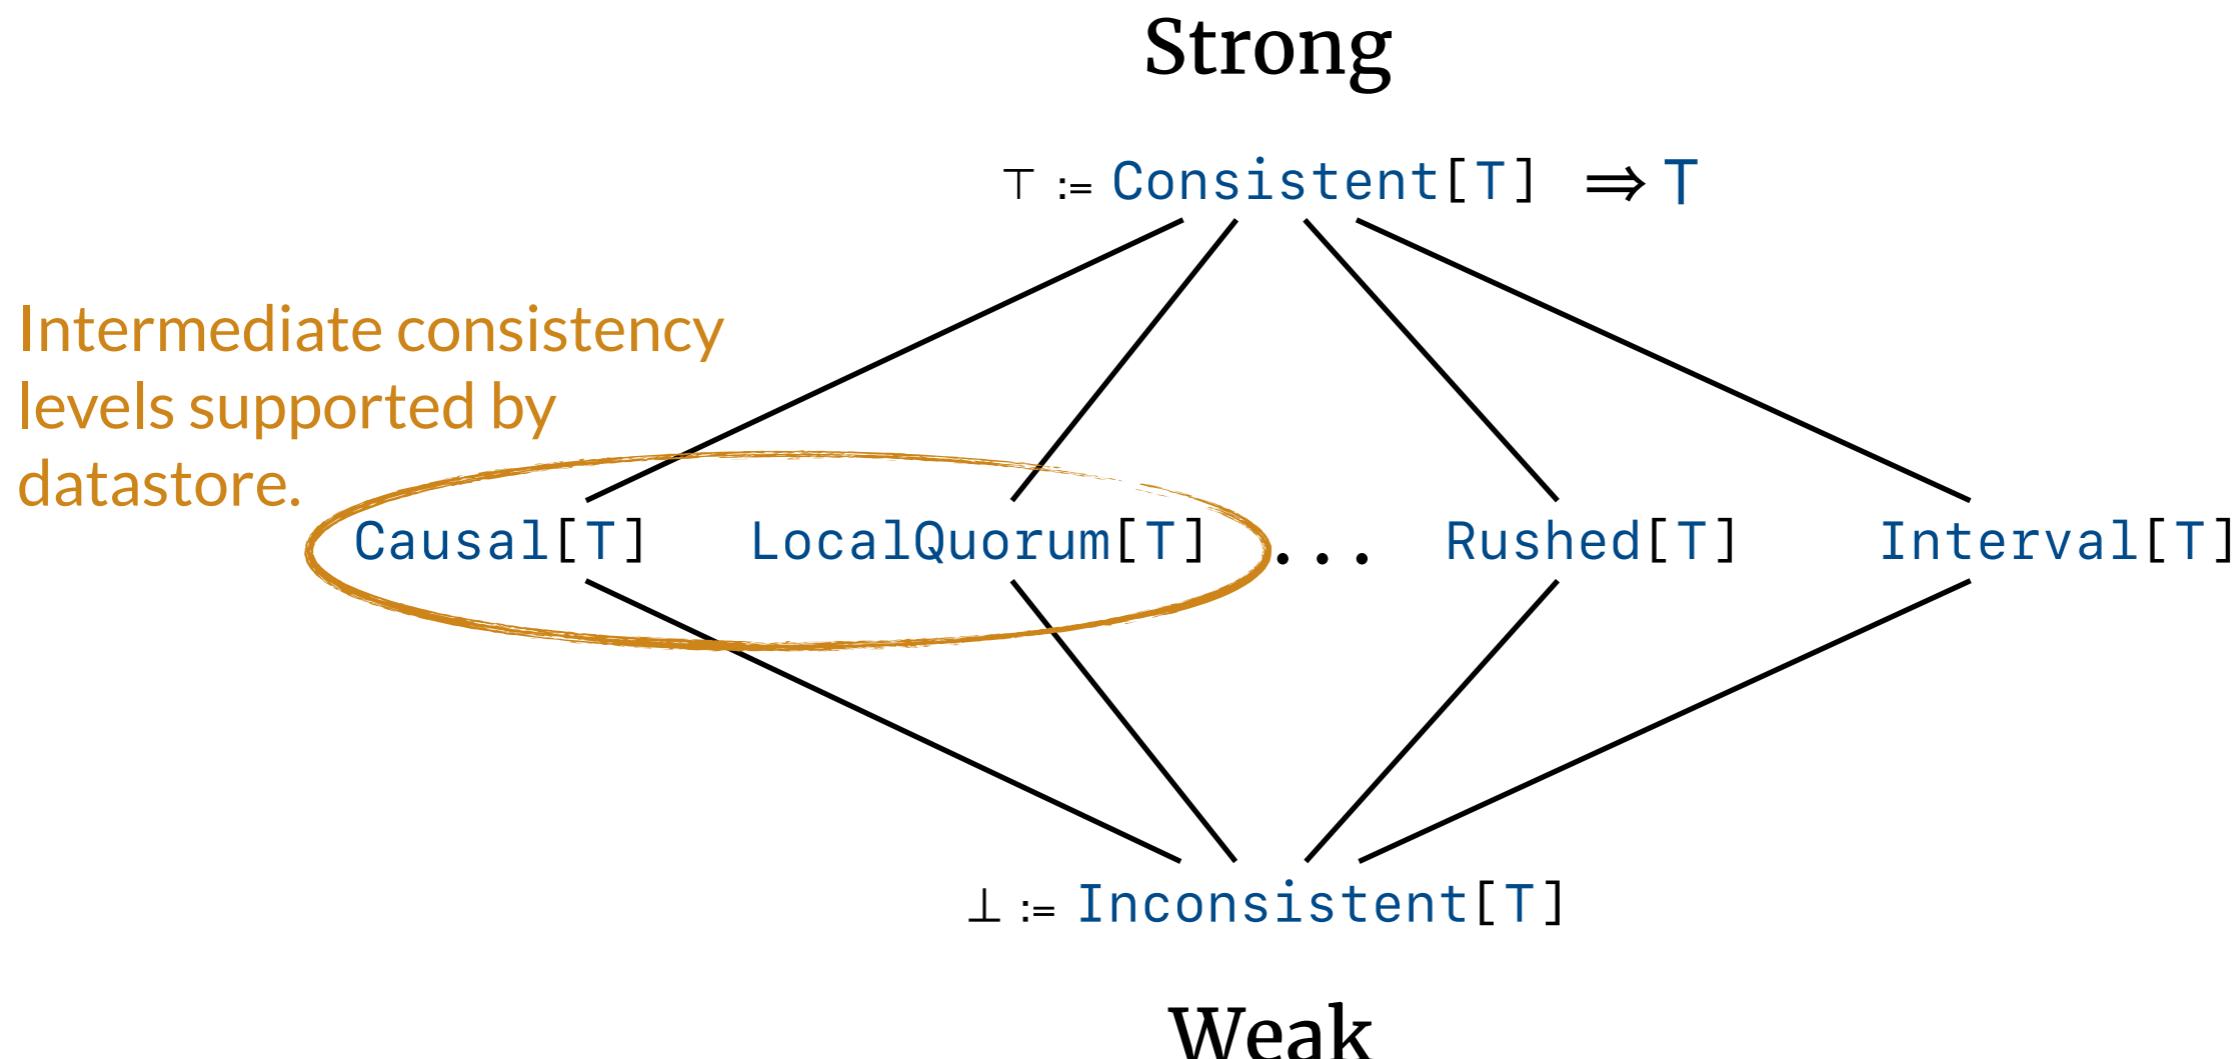
Consistency Types

Partially ordered, forming a lattice:



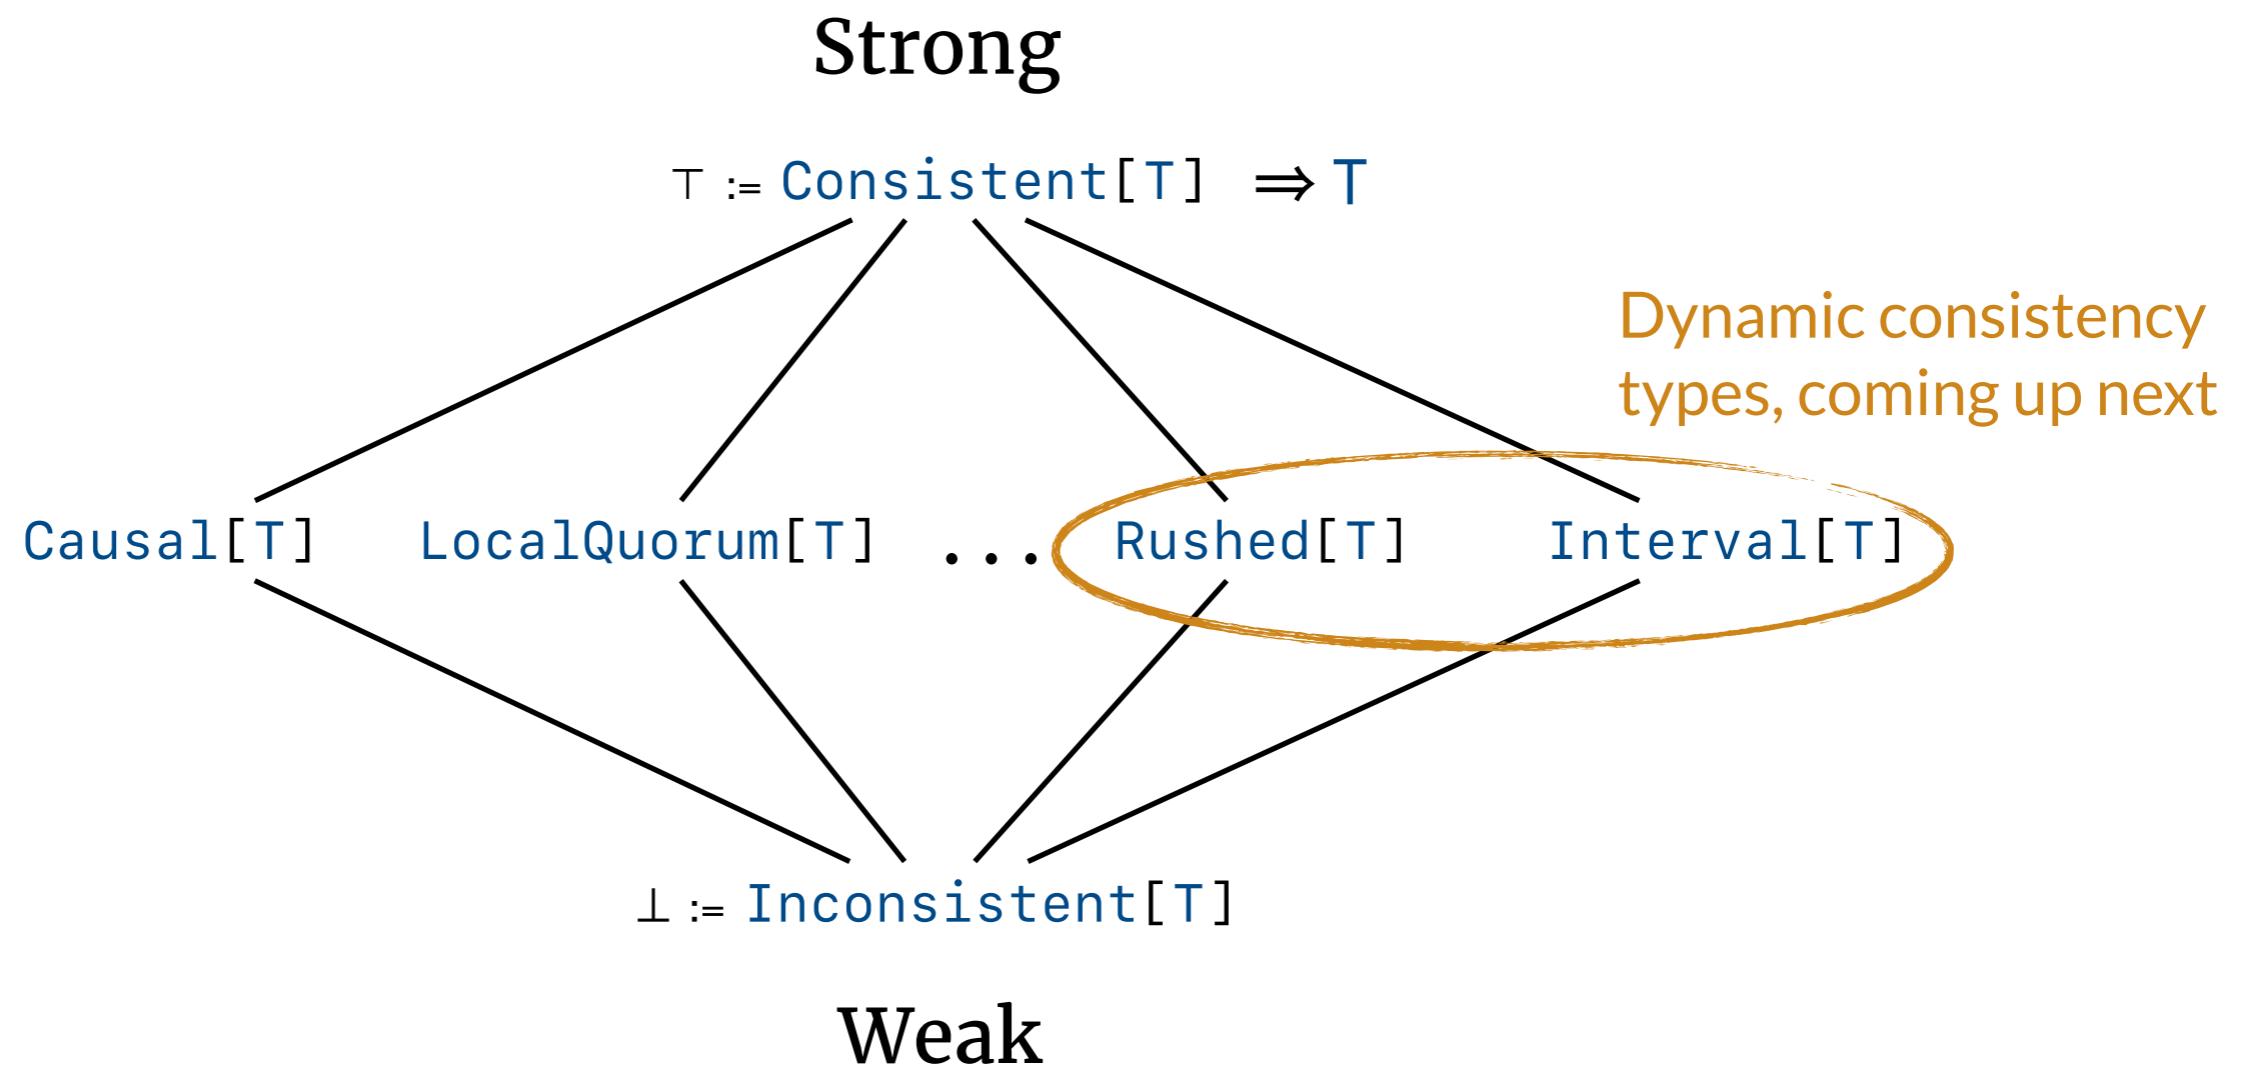
Consistency Types

Partially ordered, forming a lattice:



Consistency Types

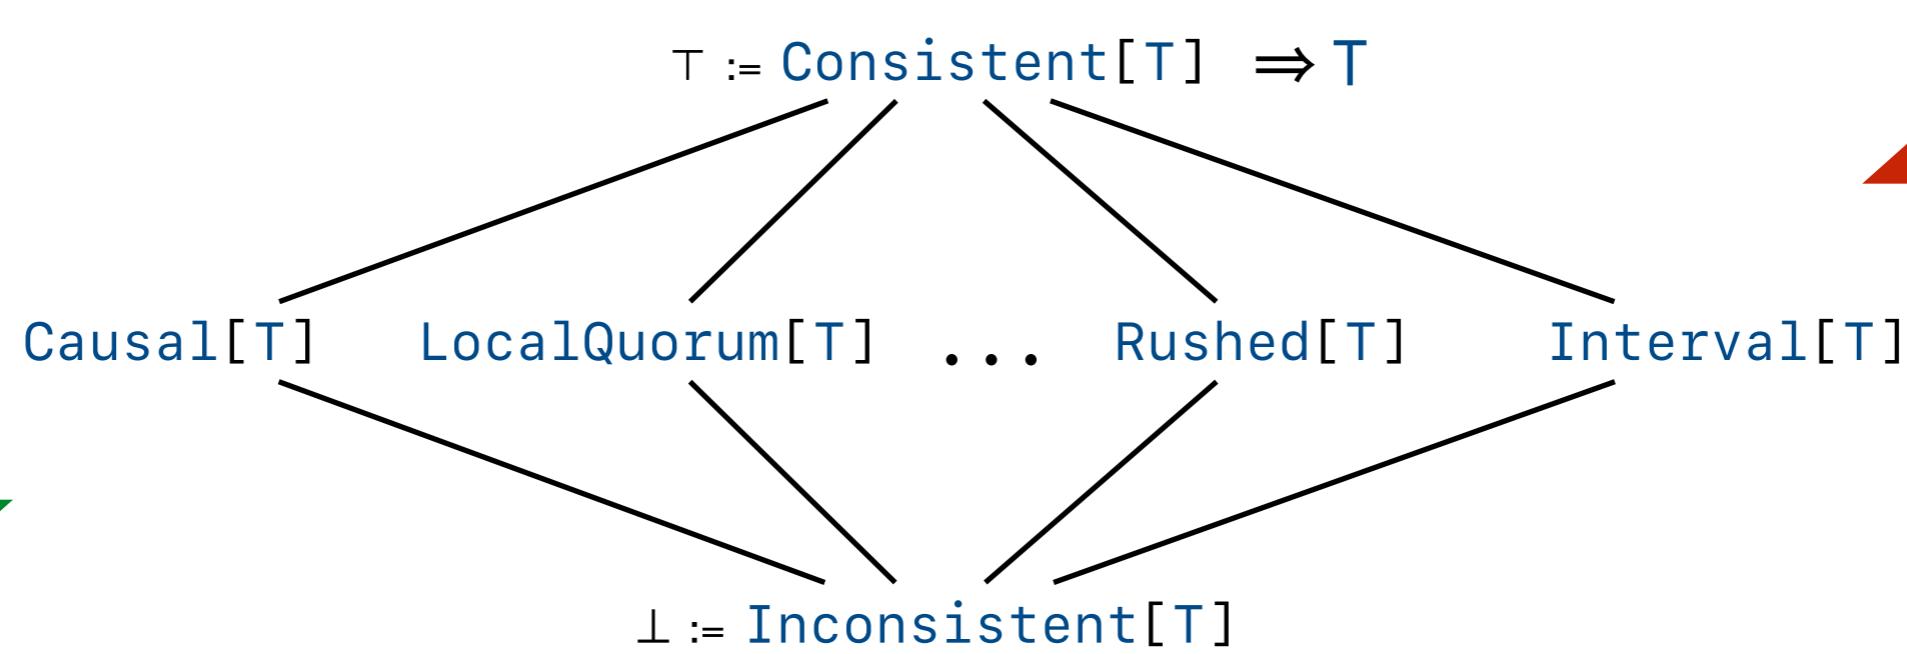
Partially ordered, forming a lattice:



Consistency Types

IPA type system enforces *consistency safety*.

Strong



Weak

Consistency safety:

Weaker values cannot flow into stronger computations.

```
val n: Inconsistent[Int] = count.read()
if n < 100: X
    print("Value is less than 100.")
```

Constraints

Every item in the datastore has a datatype with methods.

Counter:

- `.increment(): Unit`
- `.decrement(): Unit`
- `.read(): Int`

Constraints

Every item in the datastore has a datatype with methods.

Datatypes can be annotated with *consistency constraints*.

- Constraints determine the consistency type of results.

Constraints

Every item in the datastore has a datatype with methods.

Datatypes can be annotated with *consistency constraints*.

- Constraints determine the consistency type of results.

Static constraints:

Counter with Consistency(Strong) :

- .increment(): Unit
- .decrement(): Unit
- .read(): **Consistent[Int]**

Counter with Consistency(Weak) :

- .increment(): Unit
- .decrement(): Unit
- .read(): **Inconsistent[Int]**

Constraints

Every item in the datastore has a datatype with methods.

Datatypes can be annotated with *consistency constraints*.

- Constraints determine the consistency type of results.

Static constraints

- Enforced by the underlying datastore

Consistency(Strong) → Consistent[T]
Consistency(Weak) → Inconsistent[T]

Constraints

Every item in the datastore has a datatype with methods.

Datatypes can be annotated with *consistency constraints*.

- Constraints determine the consistency type of results.

Static constraints

- Enforced by the underlying datastore
 - $\text{Consistency}(\text{Strong}) \rightarrow \text{Consistent}[T]$
 - $\text{Consistency}(\text{Weak}) \rightarrow \text{Inconsistent}[T]$

Dynamic constraints

- Performance targets:
 - $\text{LatencyBound}(50 \text{ ms}) \rightarrow \text{Rushed}[T]$
- Approximation opportunity:
 - $\text{ErrorTolerance}(10\%) \rightarrow \text{Interval}[T]$

Constraints: Latency bounds

Need result in a fixed amount of time.

- Dynamically adjust consistency to meet the target latency.

Counter with LatencyBound(50 ms):

.increment(): Unit

.read(): Rushed[Int]

Consistent[Int]
LocalQuorum[Int]
Inconsistent[Int]

Return value represents achieved consistency

- **Rushed** is a *union* of other consistency types

```
count.read() match {
    case Consistent(x)    => print("consistent")
    case LocalQuorum(x)   => print("locally consistent")
    case Inconsistent(x) => print("??")
}
```

Constraints: Latency bounds

Implementation:

- Send requests to # replicas needed for strong consistency

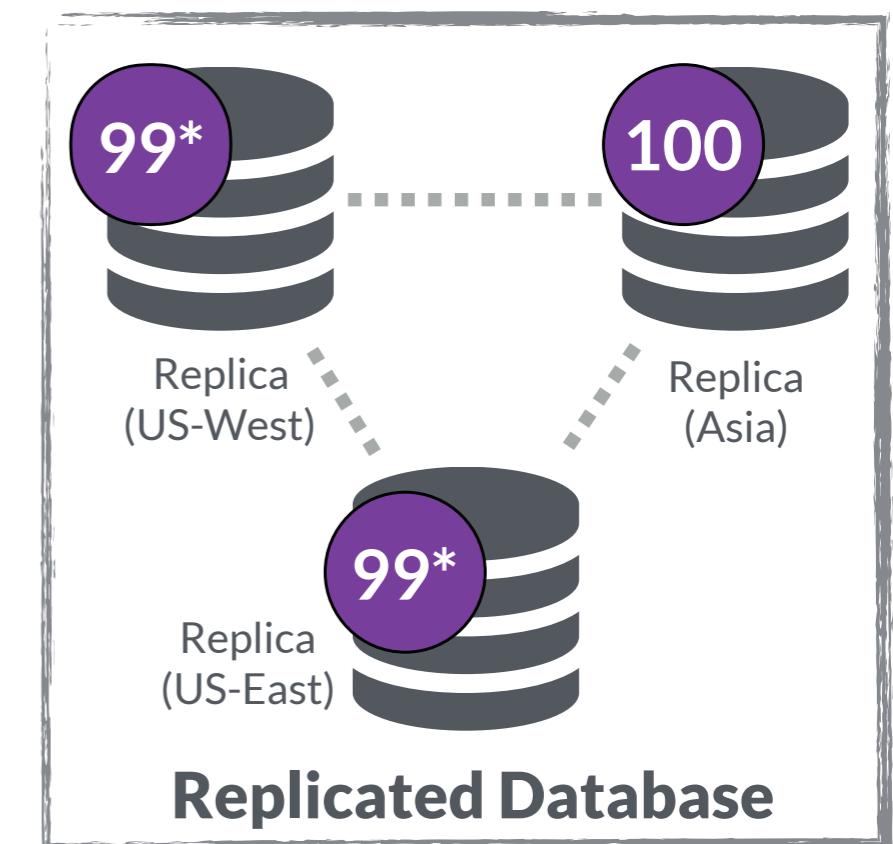
```
Counter with LatencyBound(50 ms):  
    .increment(): Unit  
    .read():     Rushed[Int]
```

Calculating consistency

Strong consistency is achieved if:

$$W + R > N$$

(W = # replicas written to,
 R = # replicas read from,
 N = total # replicas)



Constraints: Latency bounds

Implementation:

- Send requests to # replicas needed for strong consistency

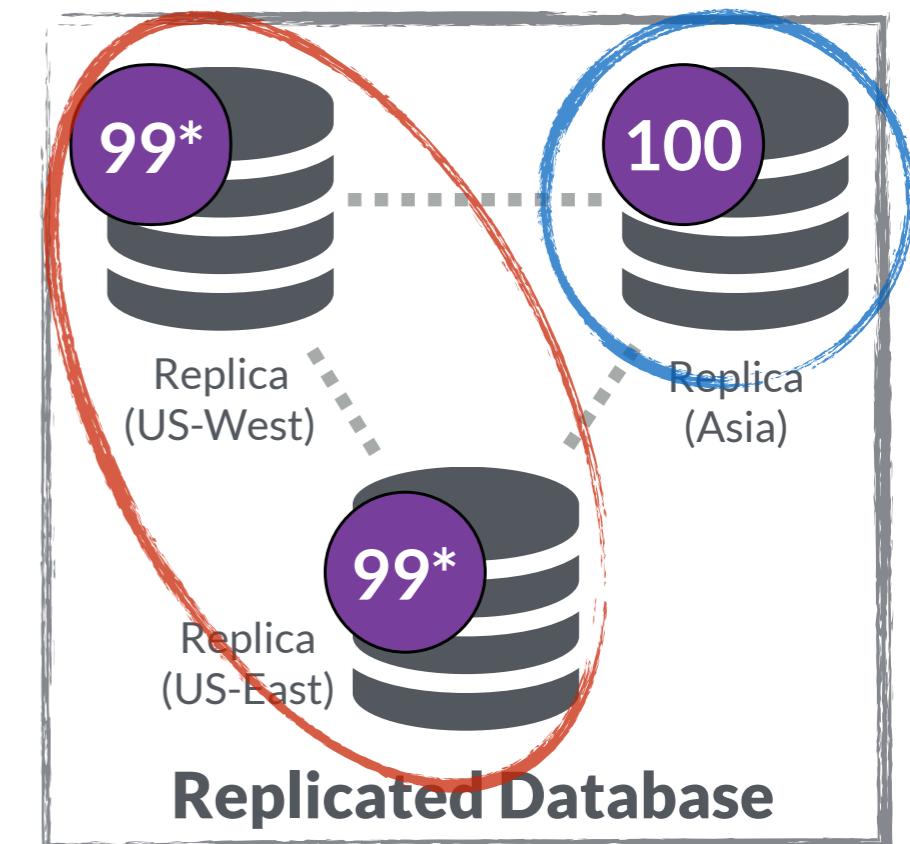
```
Counter with LatencyBound(50 ms):  
    .increment(): Unit  
    .read():     Rushed[Int]
```

Calculating consistency

Strong consistency is achieved if:

$$W + R > N$$

(W = # replicas written to,
 R = # replicas read from,
 N = total # replicas)



Constraints: Latency bounds

Implementation:

- Send requests to # replicas needed for strong consistency

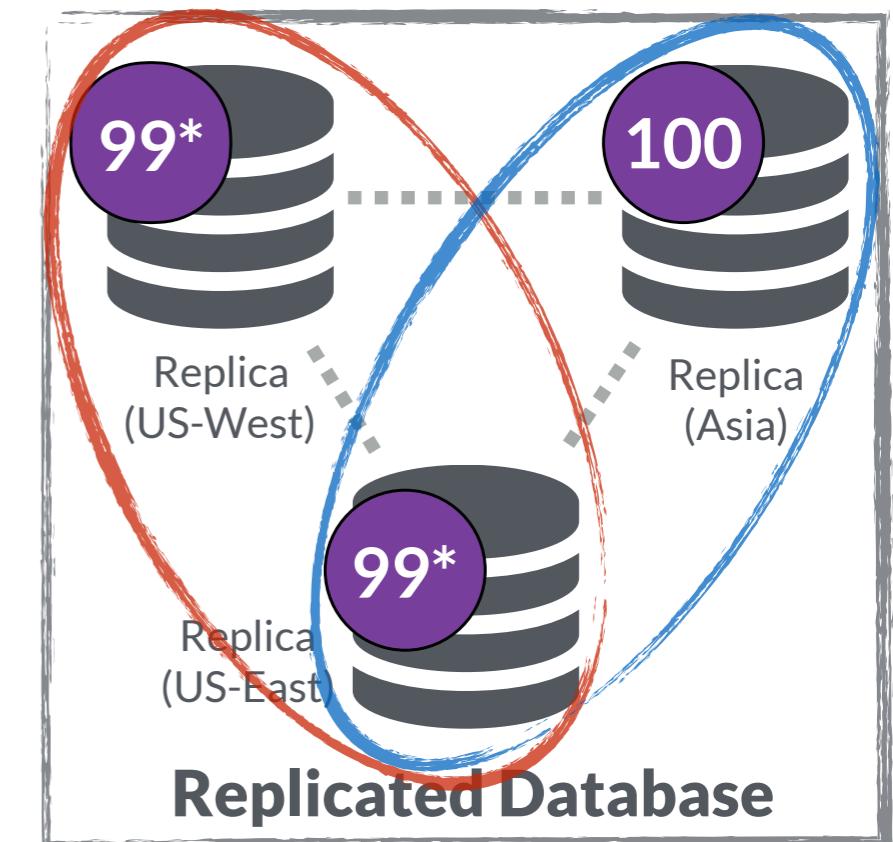
Counter with LatencyBound(50 ms):
.increment(): Unit
.read(): Rushed[Int]

Calculating consistency

Strong consistency is achieved if:

$$W + R > N$$

(W = # replicas written to,
 R = # replicas read from,
 N = total # replicas)



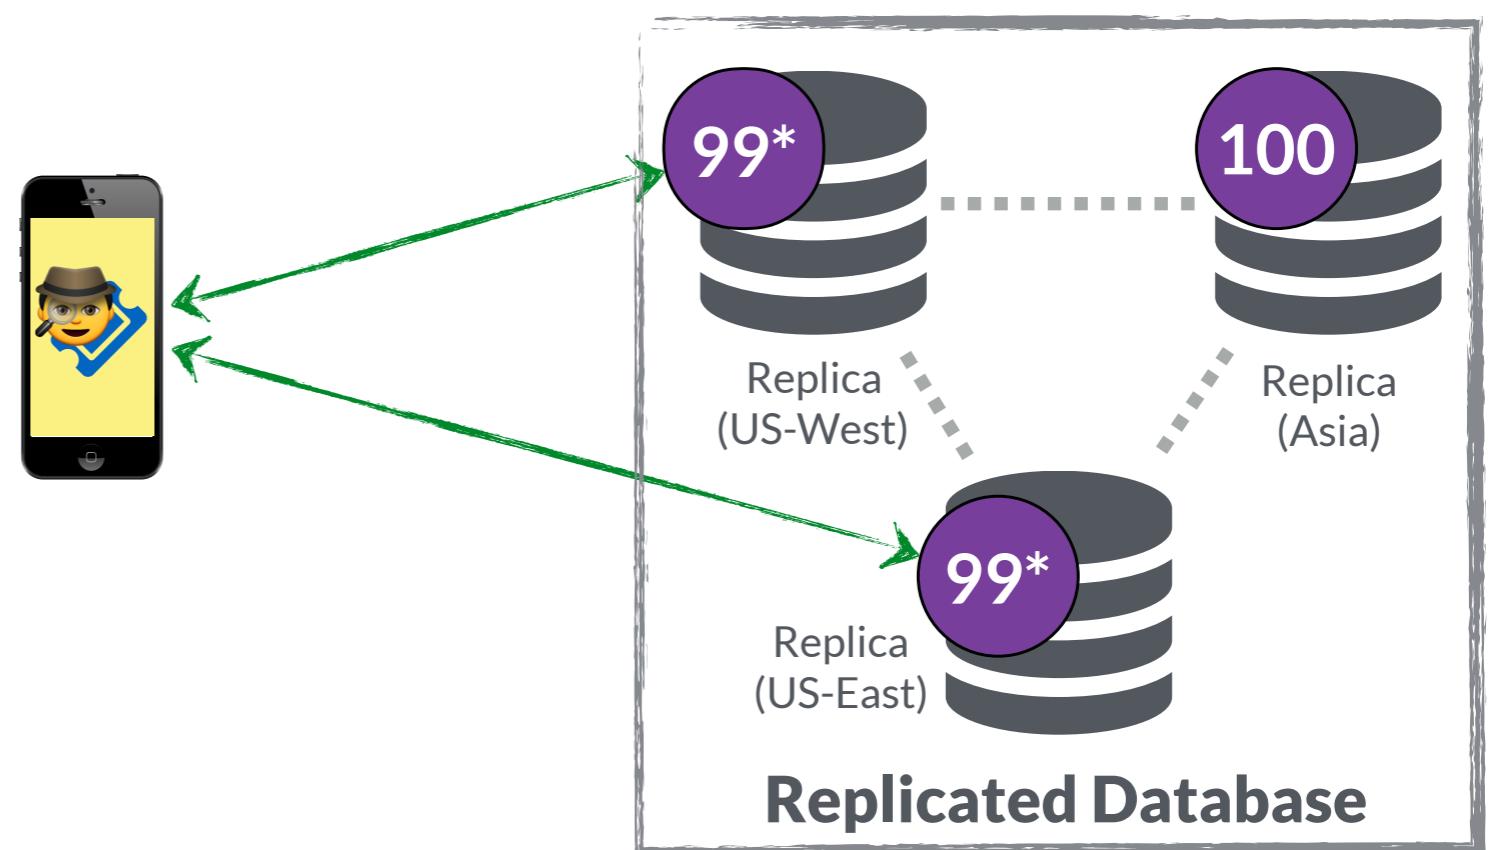
Constraints: Latency bounds

Implementation:

- Send requests to # replicas needed for strong consistency

Counter with LatencyBound(50 ms):

```
.increment(): Unit  
.read(): Rushed[Int]
```



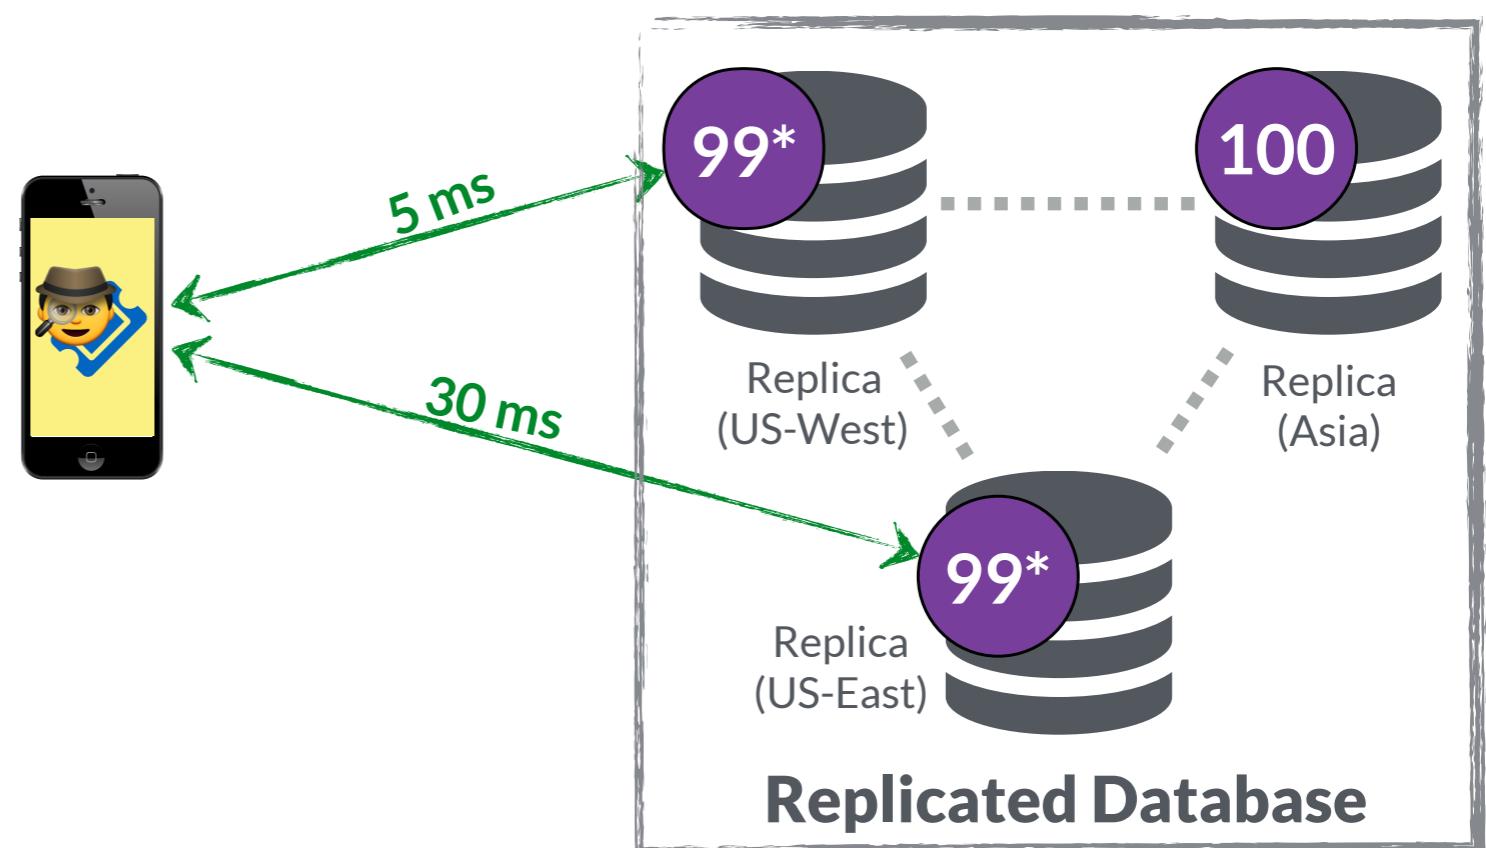
Constraints: Latency bounds

Implementation:

- Send requests to # replicas needed for strong consistency
- Wait for time limit

Counter with LatencyBound(50 ms):

```
.increment(): Unit  
.read(): Rushed[Int]
```

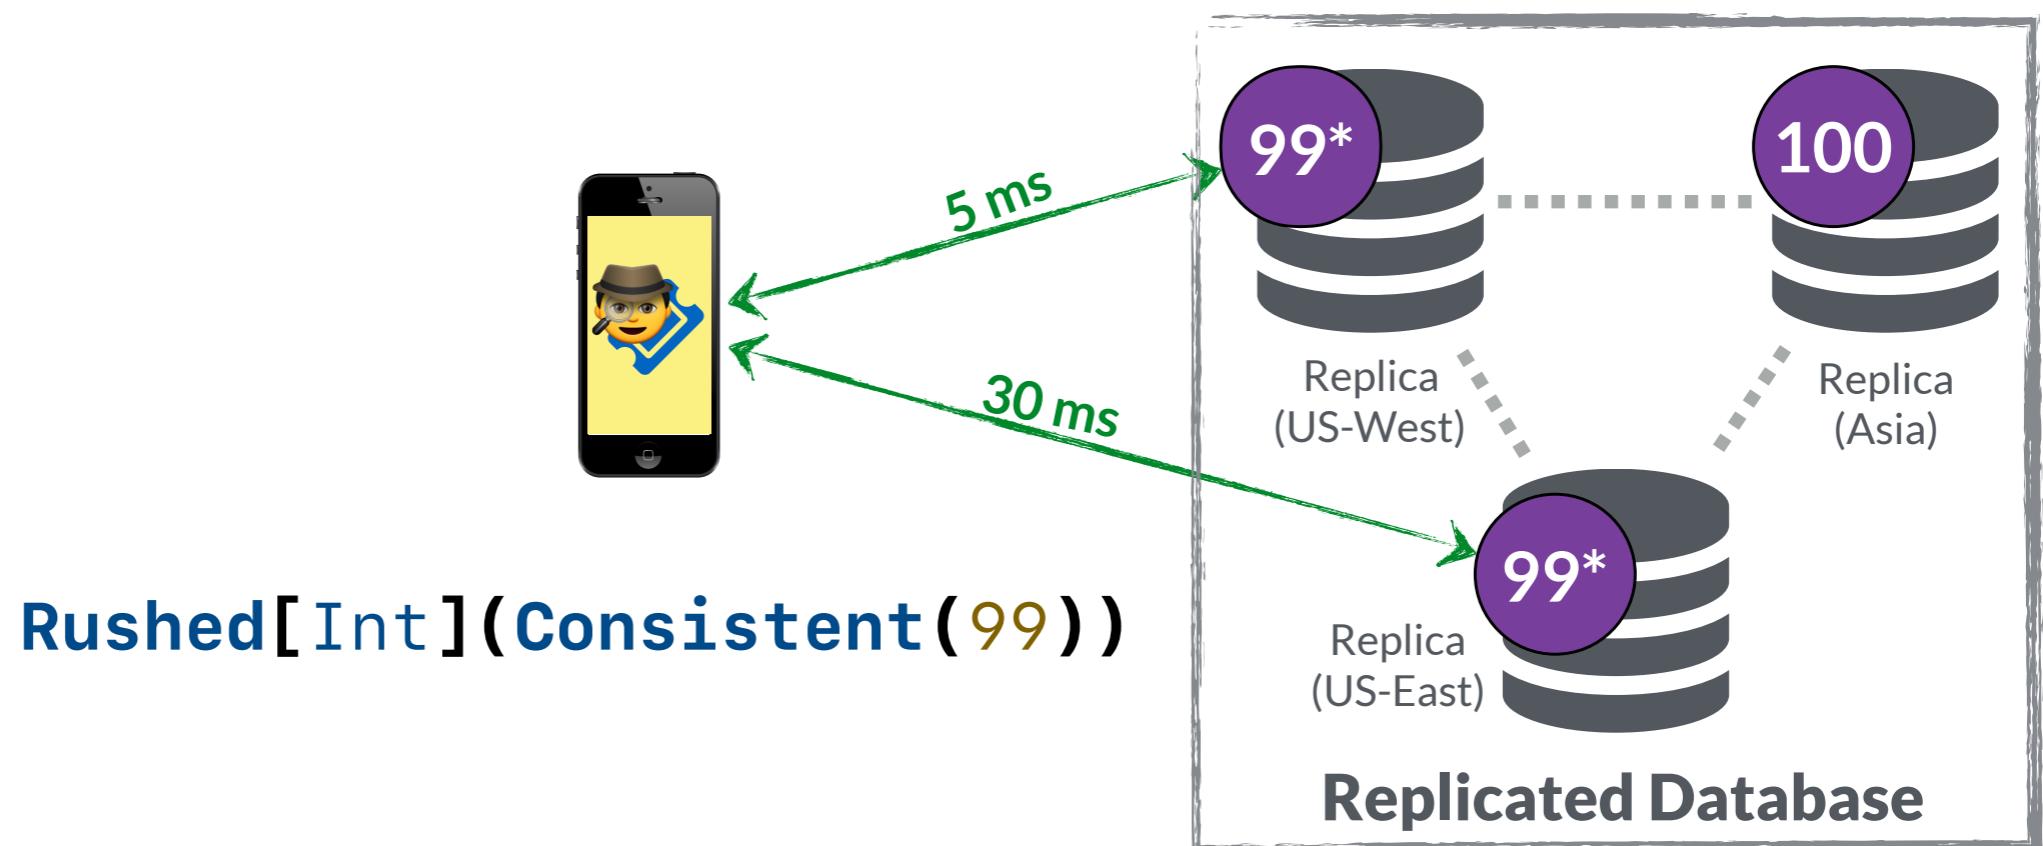


Constraints: Latency bounds

Implementation:

- Send requests to # replicas needed for strong consistency
- Wait for time limit
- Resulting consistency determined by number of on-time responses

Counter with LatencyBound(50 ms):
.increment(): Unit
.read(): Rushed[Int]



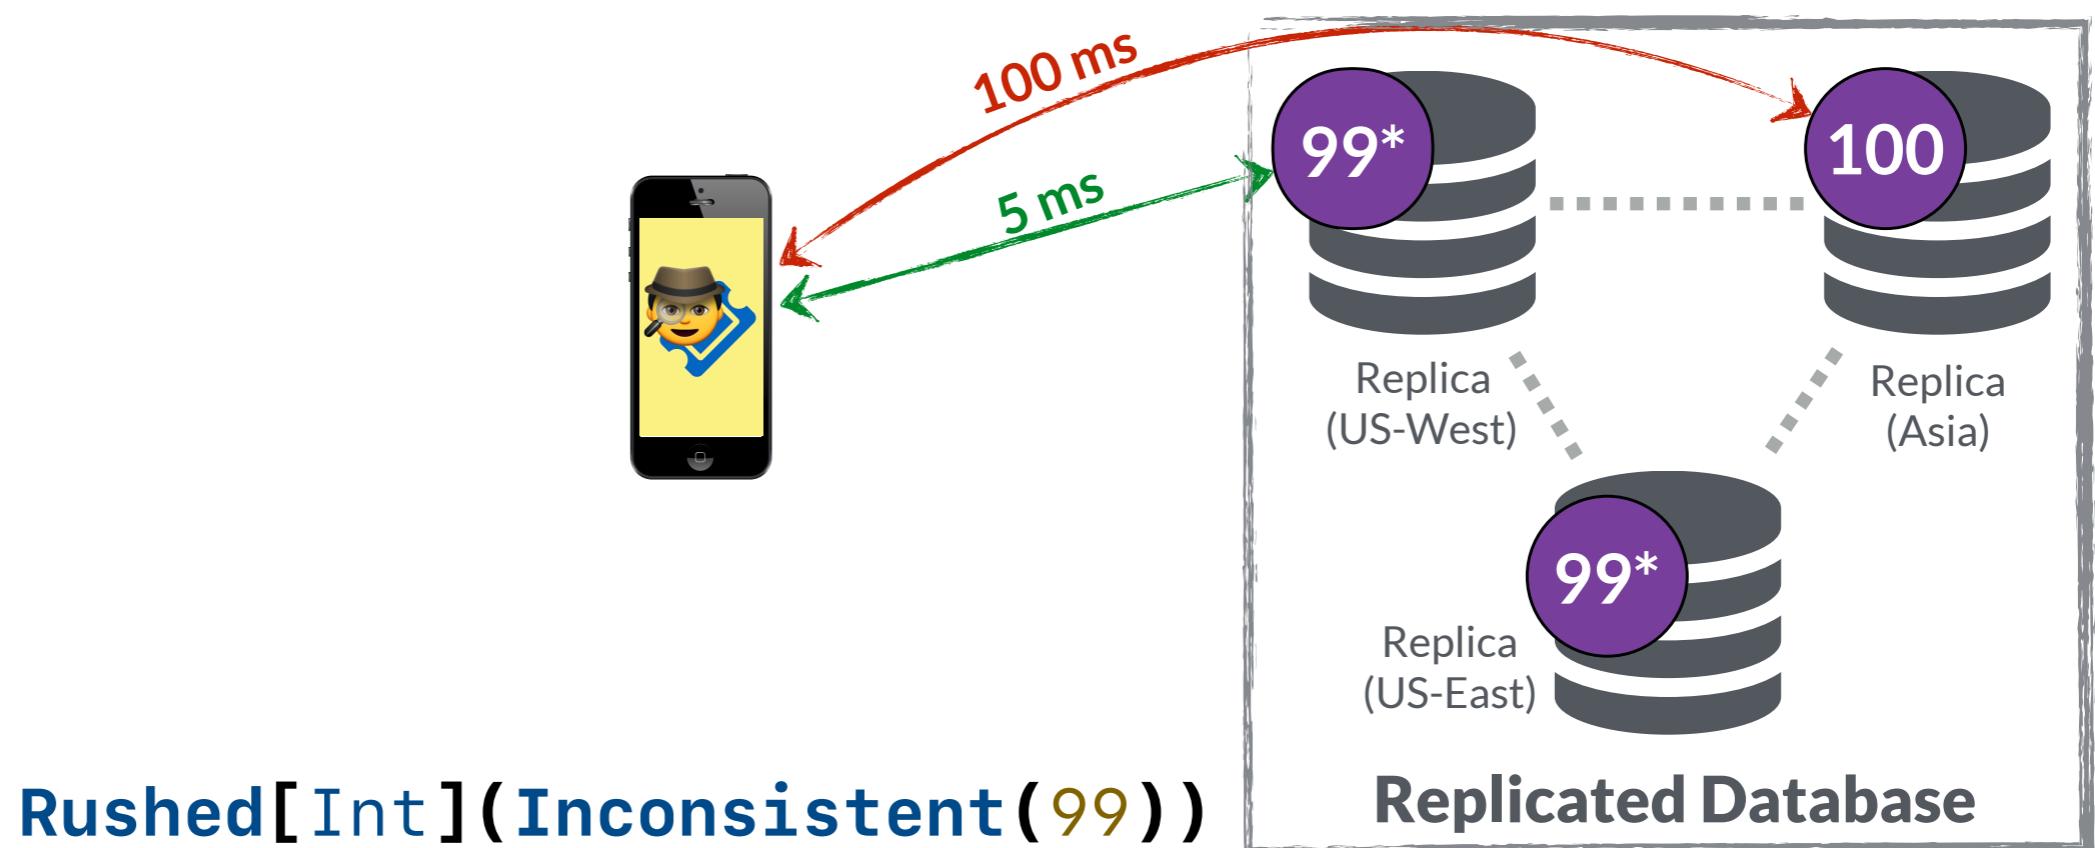
Constraints: Latency bounds

Implementation:

- Send requests to # replicas needed for strong consistency
- Wait for time limit
- Resulting consistency determined by number of on-time responses

Counter with LatencyBound(50 ms):

```
.increment(): Unit  
.read(): Rushed[Int]
```



Constraints: Error bounds

Tolerate a fixed amount of *inaccuracy* in value.

Counter with ErrorTolerance(10%):

- .decrement(): Unit
- .read(): **Interval[Int]**

Return value represents a range of possible consistent values

- Programmer must decide what to do with the **Interval** value

```
val remaining = ticketCount.read()
if (remaining.min > 100) {
    display("At least 100 tickets left.")
}
```

Works by limiting outstanding writes

- Requires knowing the semantics of each operation

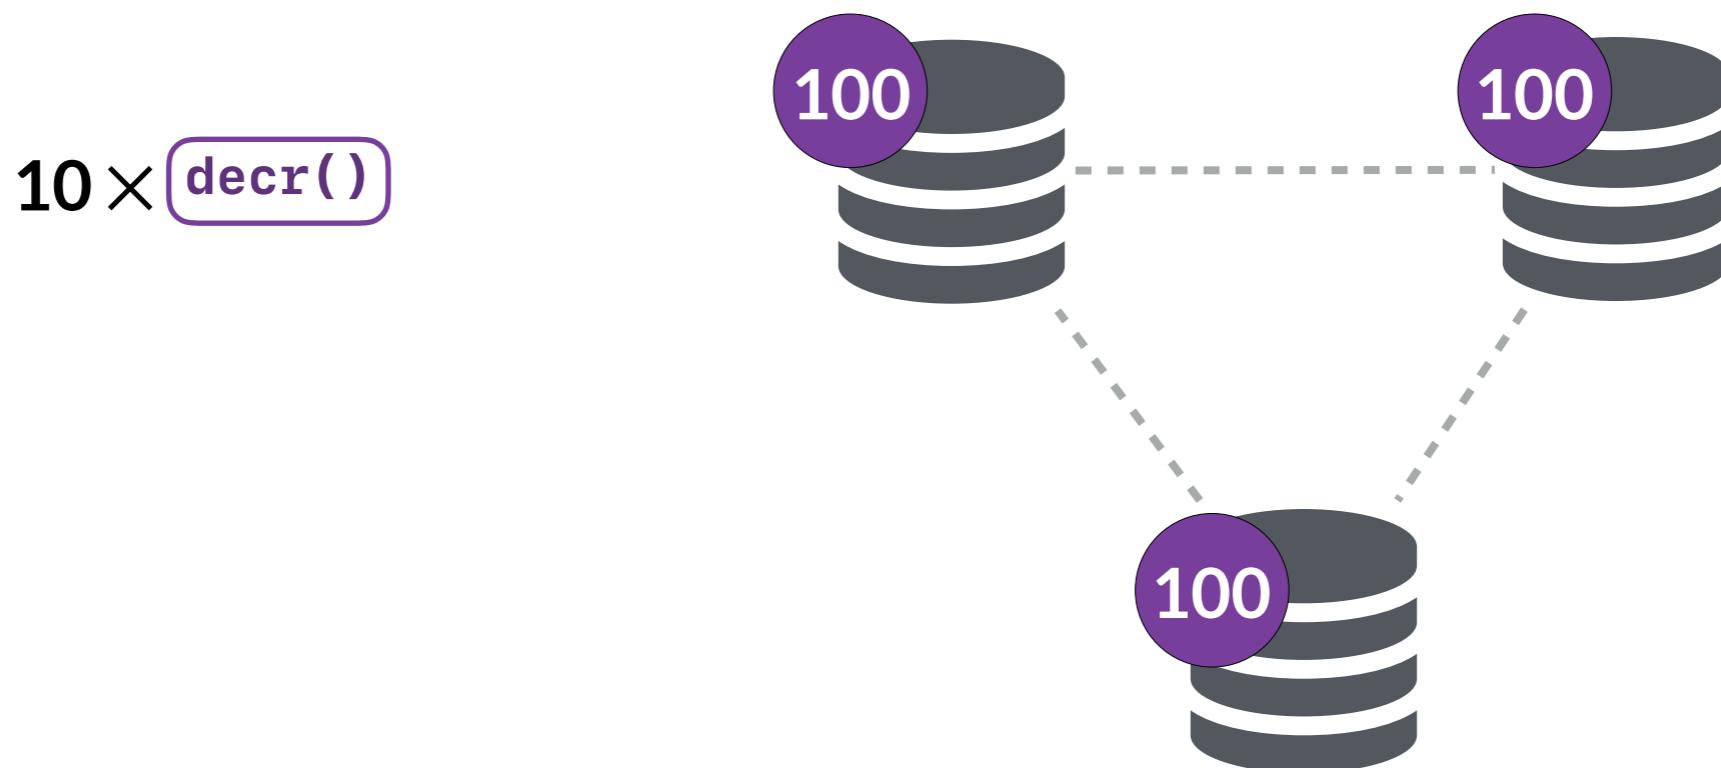
`decrement() × (0..10) ⇒ read() = Interval[Int](90, 100)`

Constraints: Error bounds

Implementation:

- Reservations: pre-allocate permission to do updates, distribute among replicas.

Counter with ErrorTolerance(10%):
.decrement(): Unit
.read(): Interval[Int]

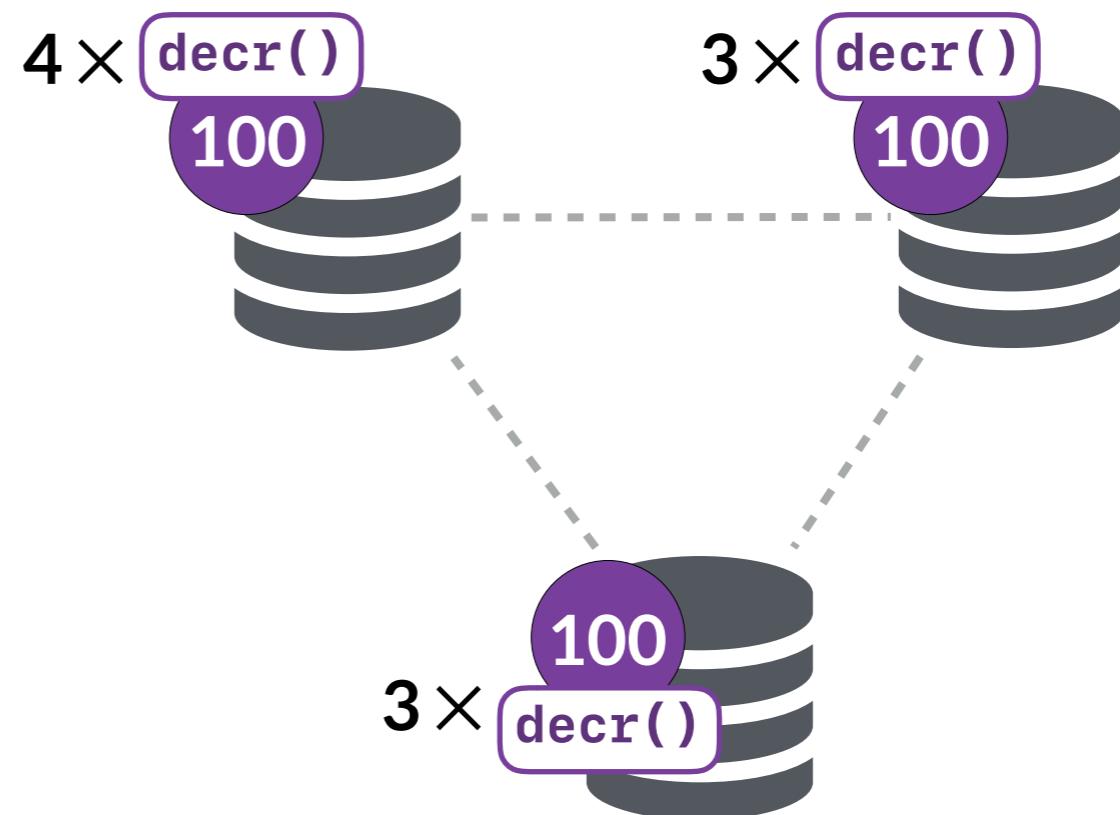


Constraints: Error bounds

Implementation:

- Reservations: pre-allocate permission to do updates, distribute among replicas.

Counter with ErrorTolerance(10%):
.decrement(): Unit
.read(): Interval[Int]

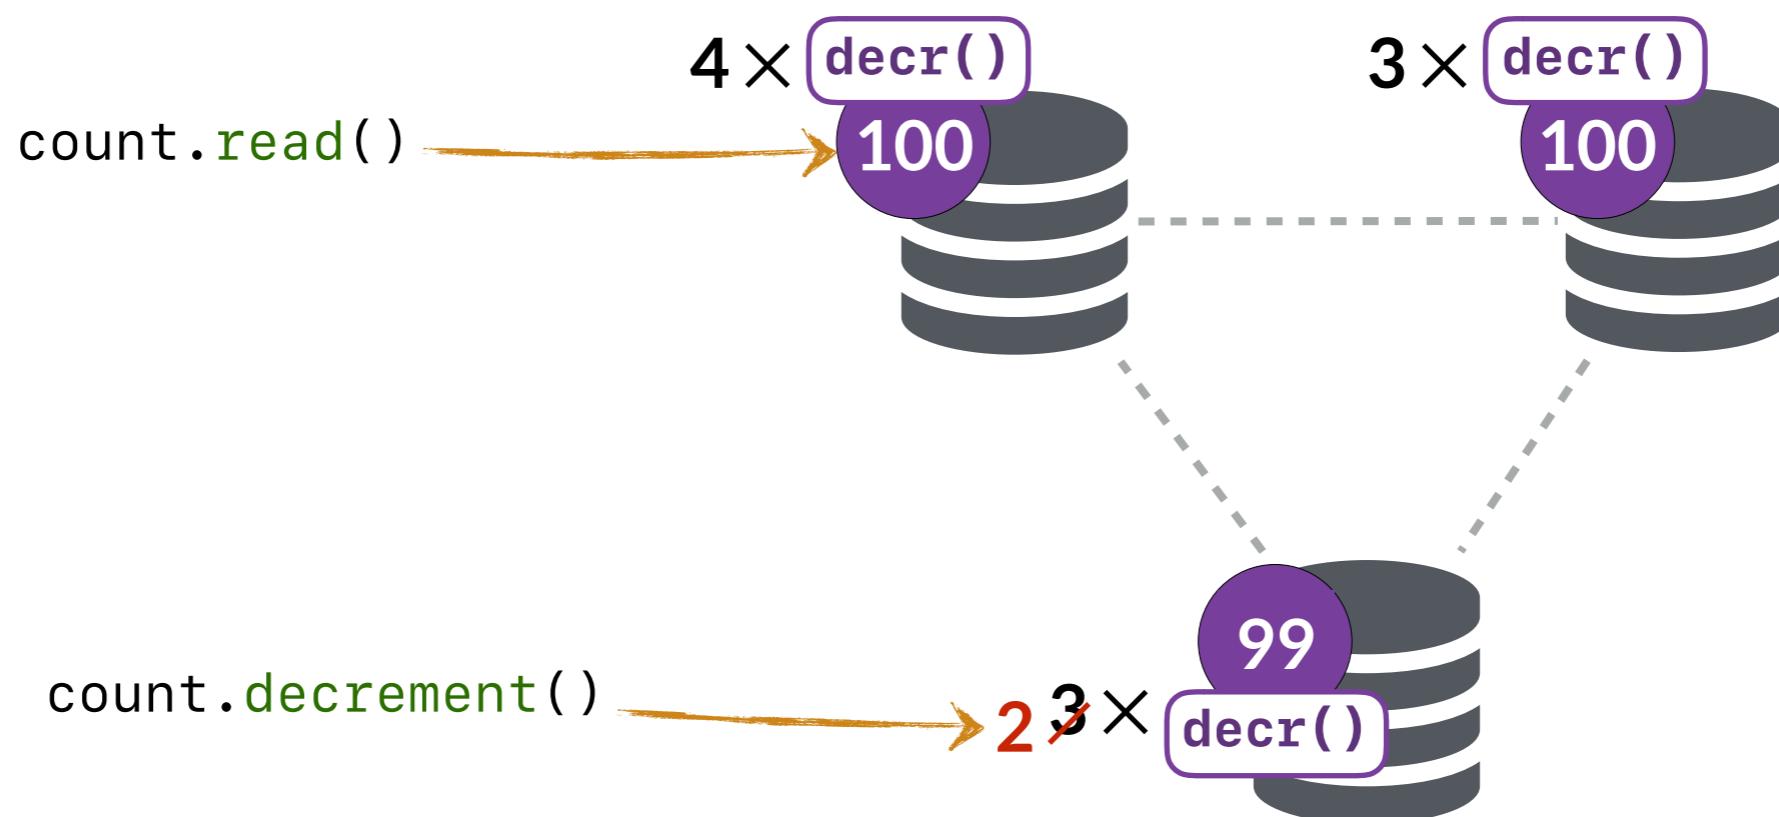


Constraints: Error bounds

Implementation:

- Reservations: pre-allocate permission to do updates, distribute among replicas.

Counter with ErrorTolerance(10%):
.decrement(): Unit
.read(): Interval[Int]

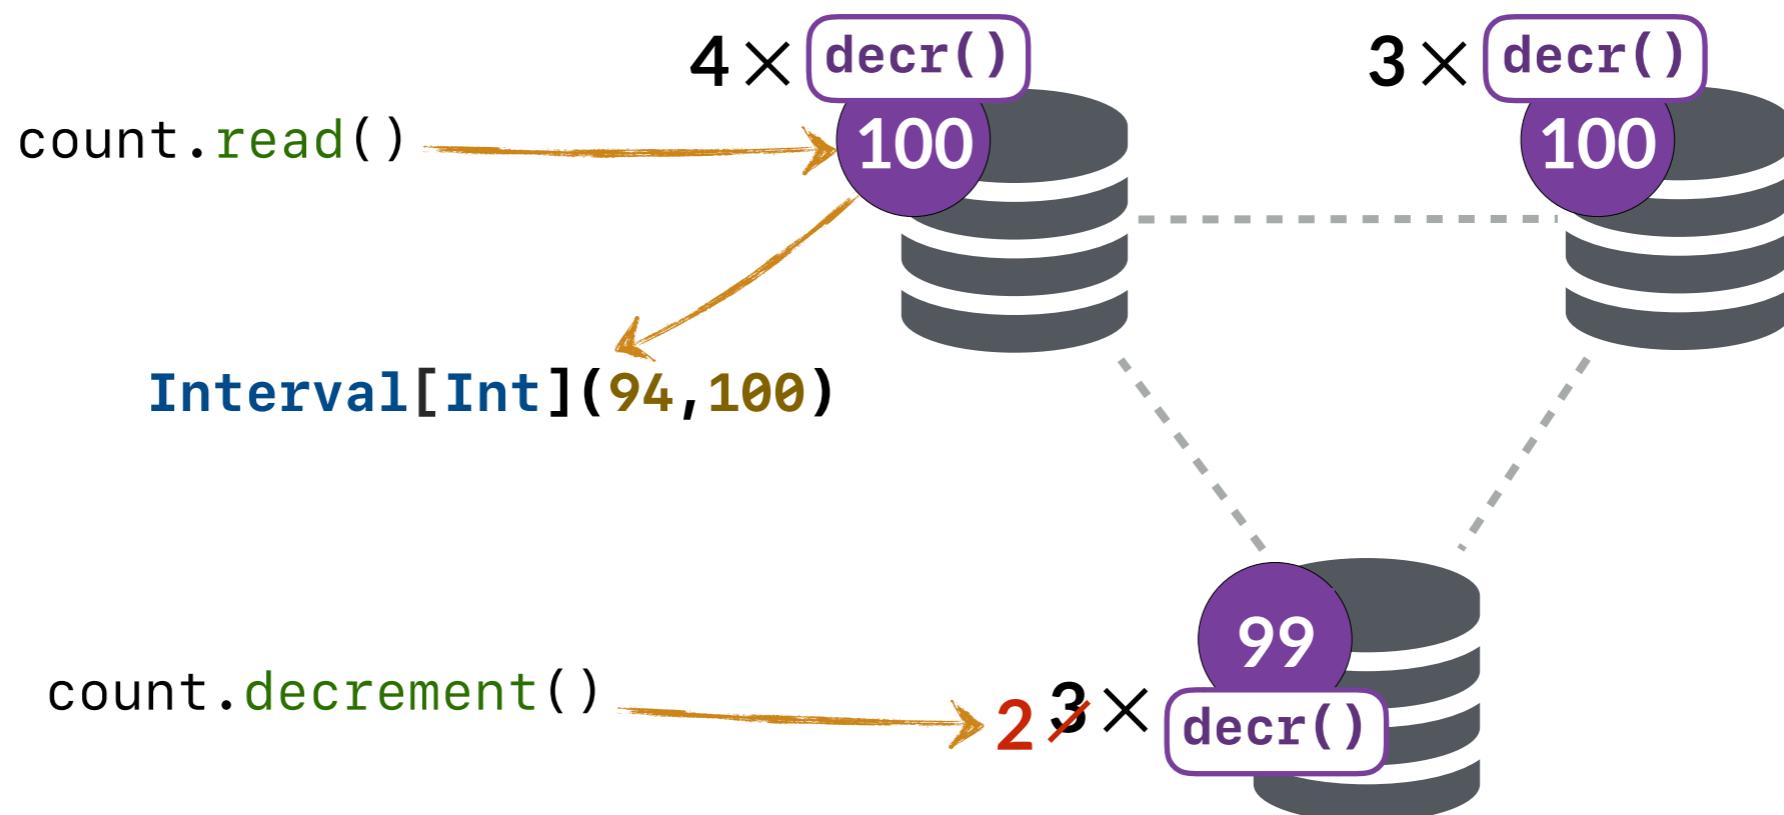


Constraints: Error bounds

Implementation:

- Reservations: pre-allocate permission to do updates, distribute among replicas.

Counter with ErrorTolerance(10%):
.decrement(): Unit
.read(): Interval[Int]

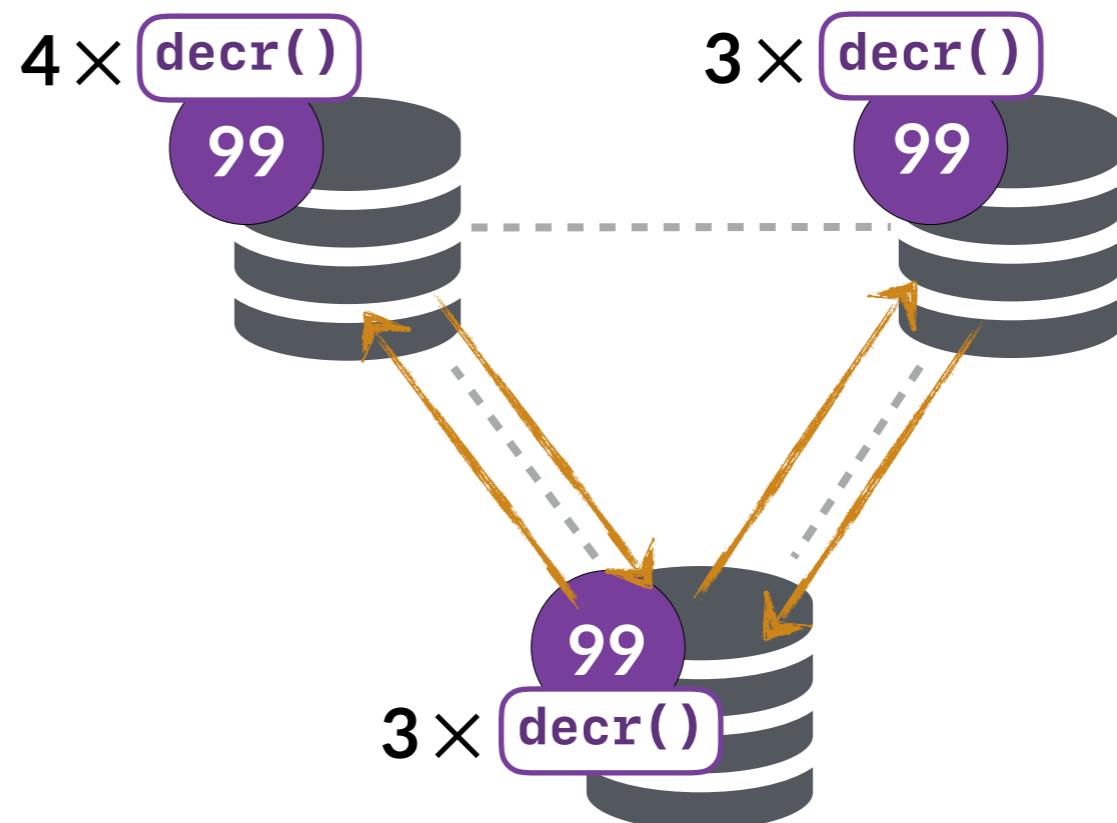


Constraints: Error bounds

Implementation:

- Reservations: pre-allocate permission to do updates, distribute among replicas.
- Refresh permissions on sync.

Counter with ErrorTolerance(10%):
.decrement(): Unit
.read(): Interval[Int]

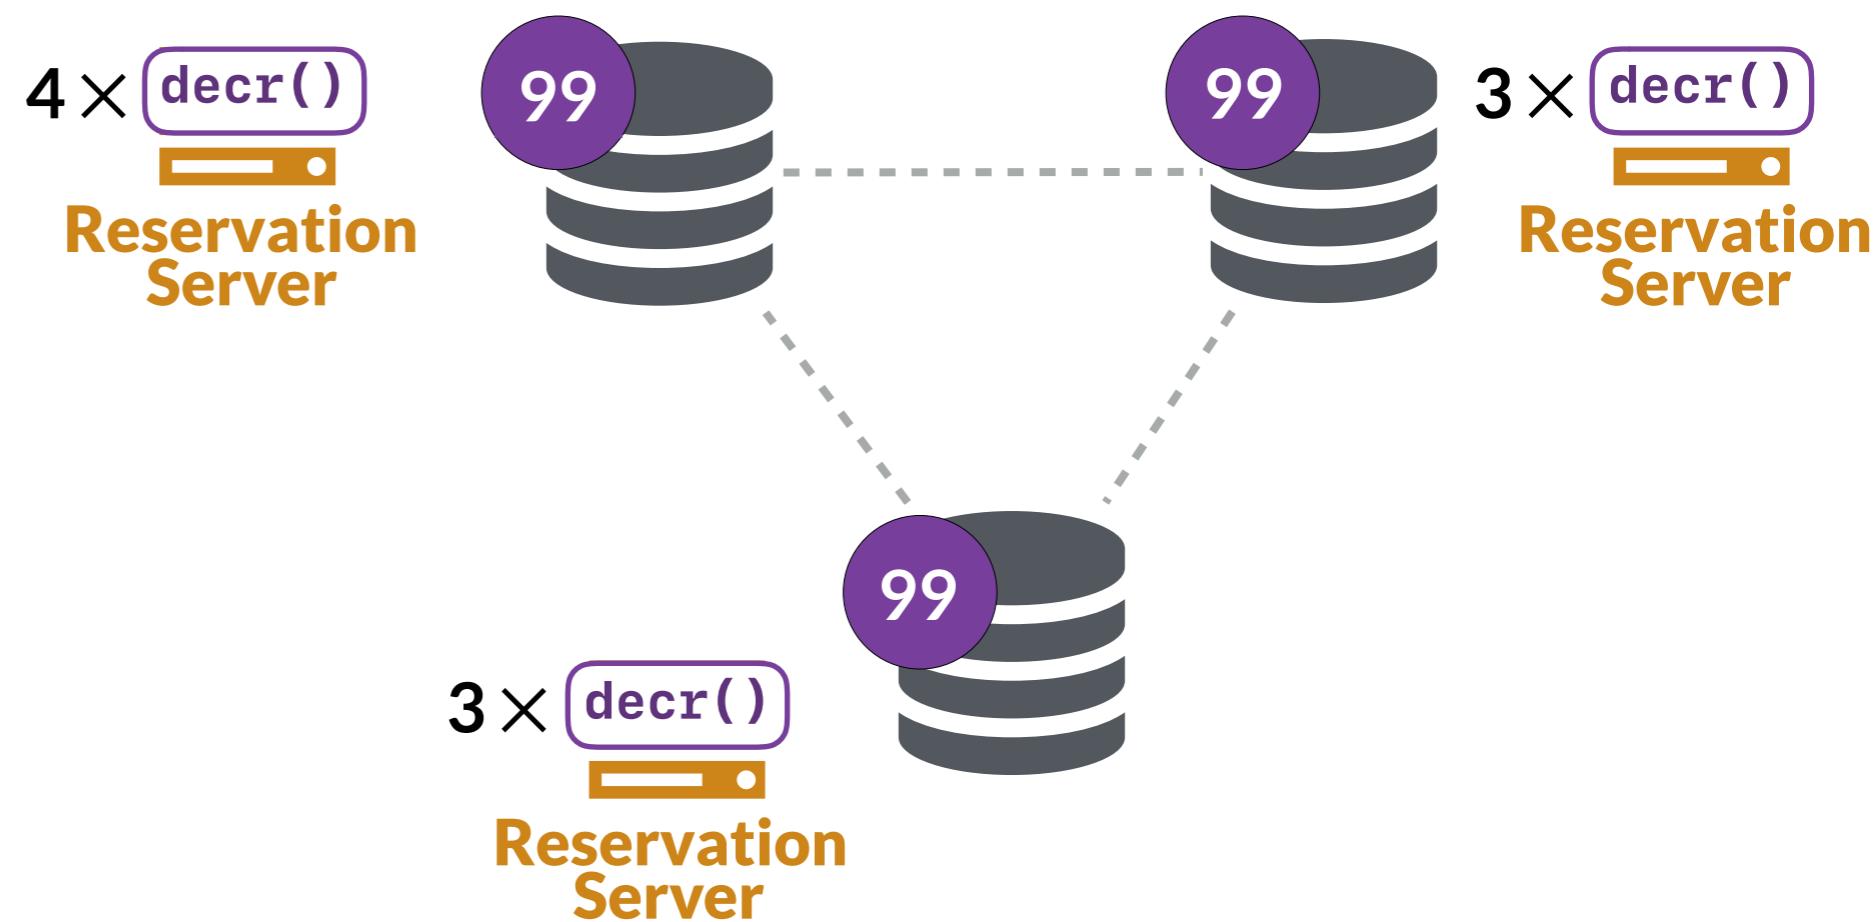


Constraints: Error bounds

Implementation:

- Reservations: pre-allocate permission to do updates, distribute among replicas.
- Refresh permissions on sync.
- Reservation servers mediate requests.

Counter with ErrorTolerance(10%):
.decrement(): Unit
.read(): Interval[Int]





Inconsistent
Performance-bound
Approximate } **Programming Model**

Trading off consistency

Consistency types

- Statically prevent consistency bugs

Performance & accuracy constraints

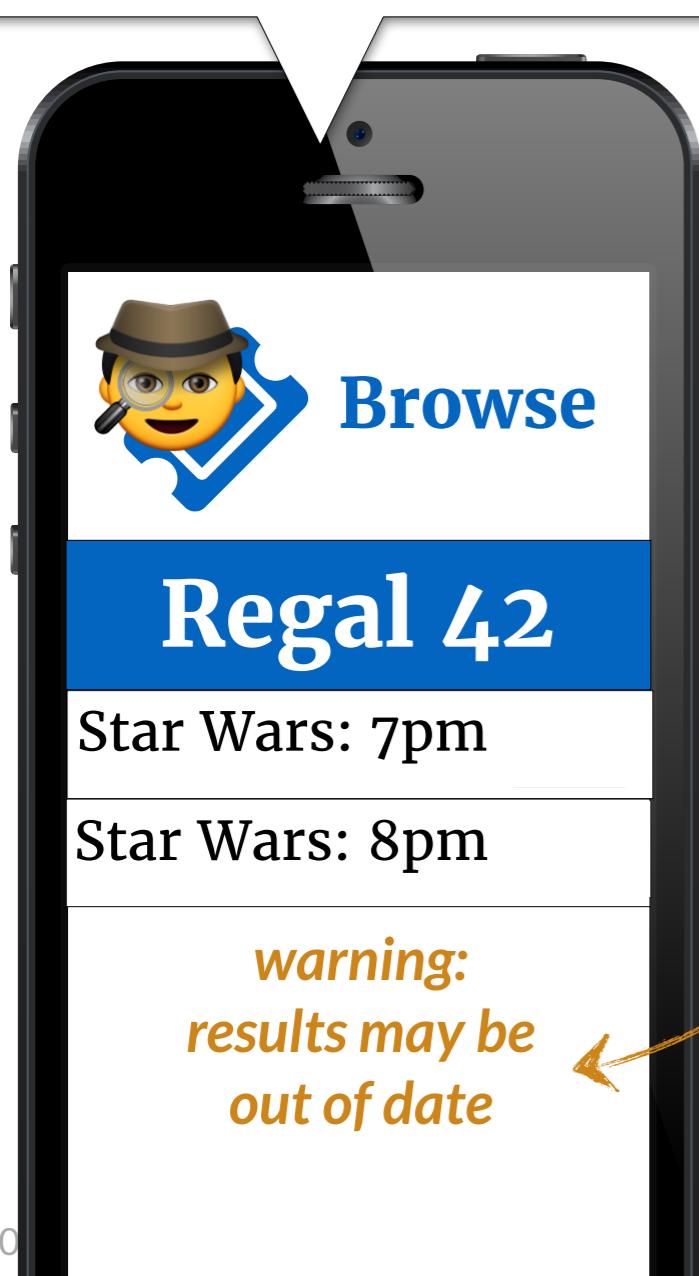
- Dynamically adjust consistency

Case study

Case Study: TICKETSLEUTH

Performance constraint:

< 50 ms latency



```
val events = List[EventID]()  
with LatencyBound(50 ms)
```

```
events("Regal42:StarWars").range(0, 2)
```

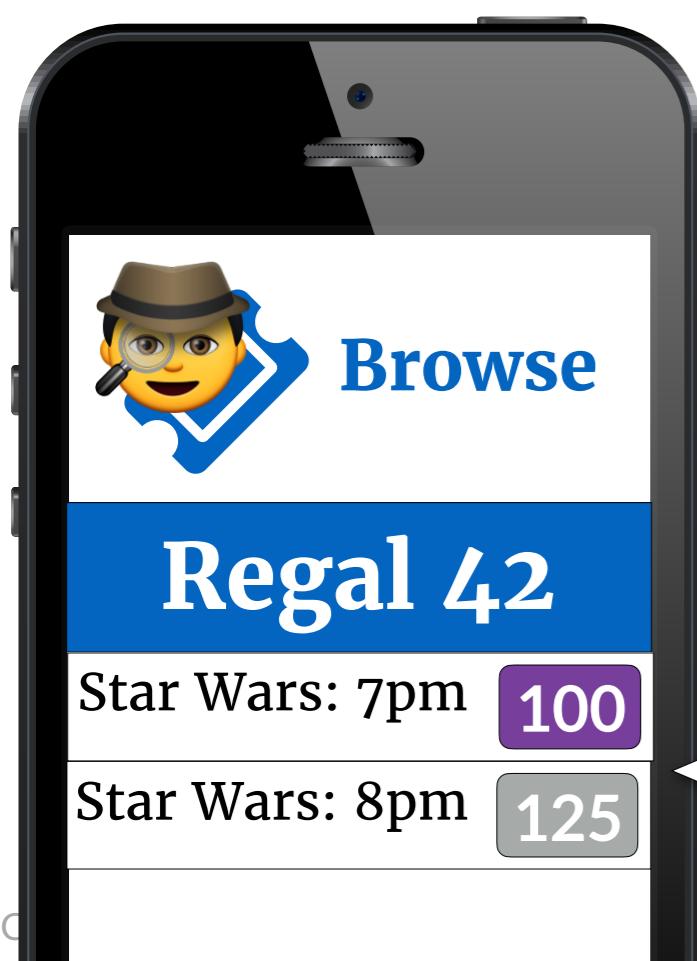
```
Rushed(Inconsistent(List("7pm", "8pm")))
```

endorse

*warning:
results may be
out of date*

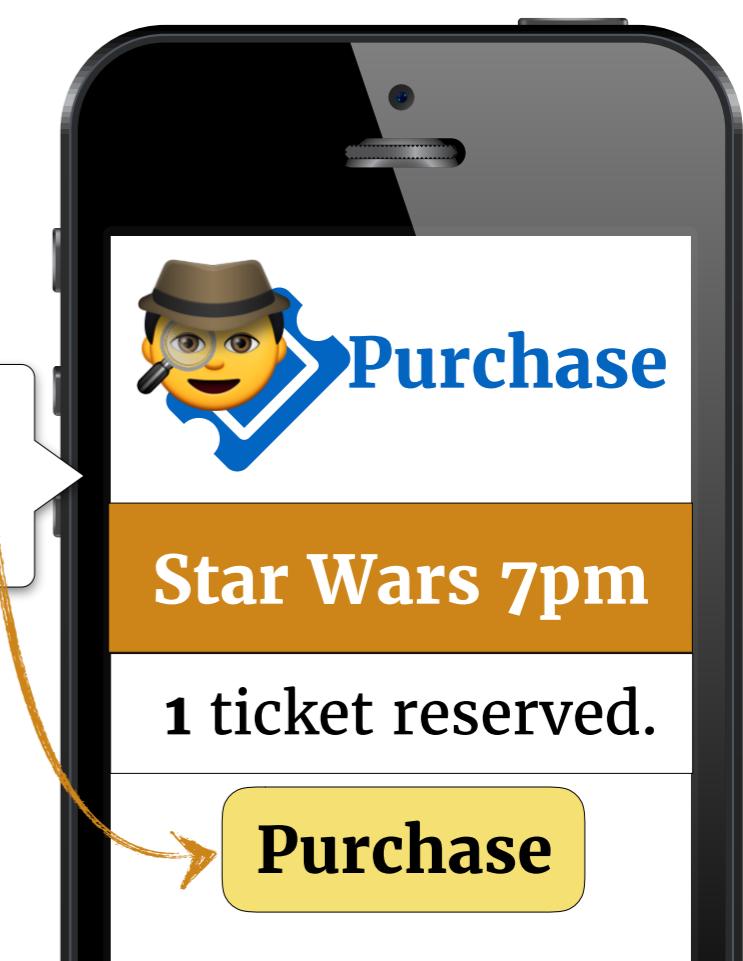
Case Study: TICKETSLEUTH

```
val tickets = UUIDPool()  
  with ErrorTolerance(count, 5%)  
  
  tickets.count() // Interval[Int](96, 100)  
  tickets.take() // Consistent[UUID]
```

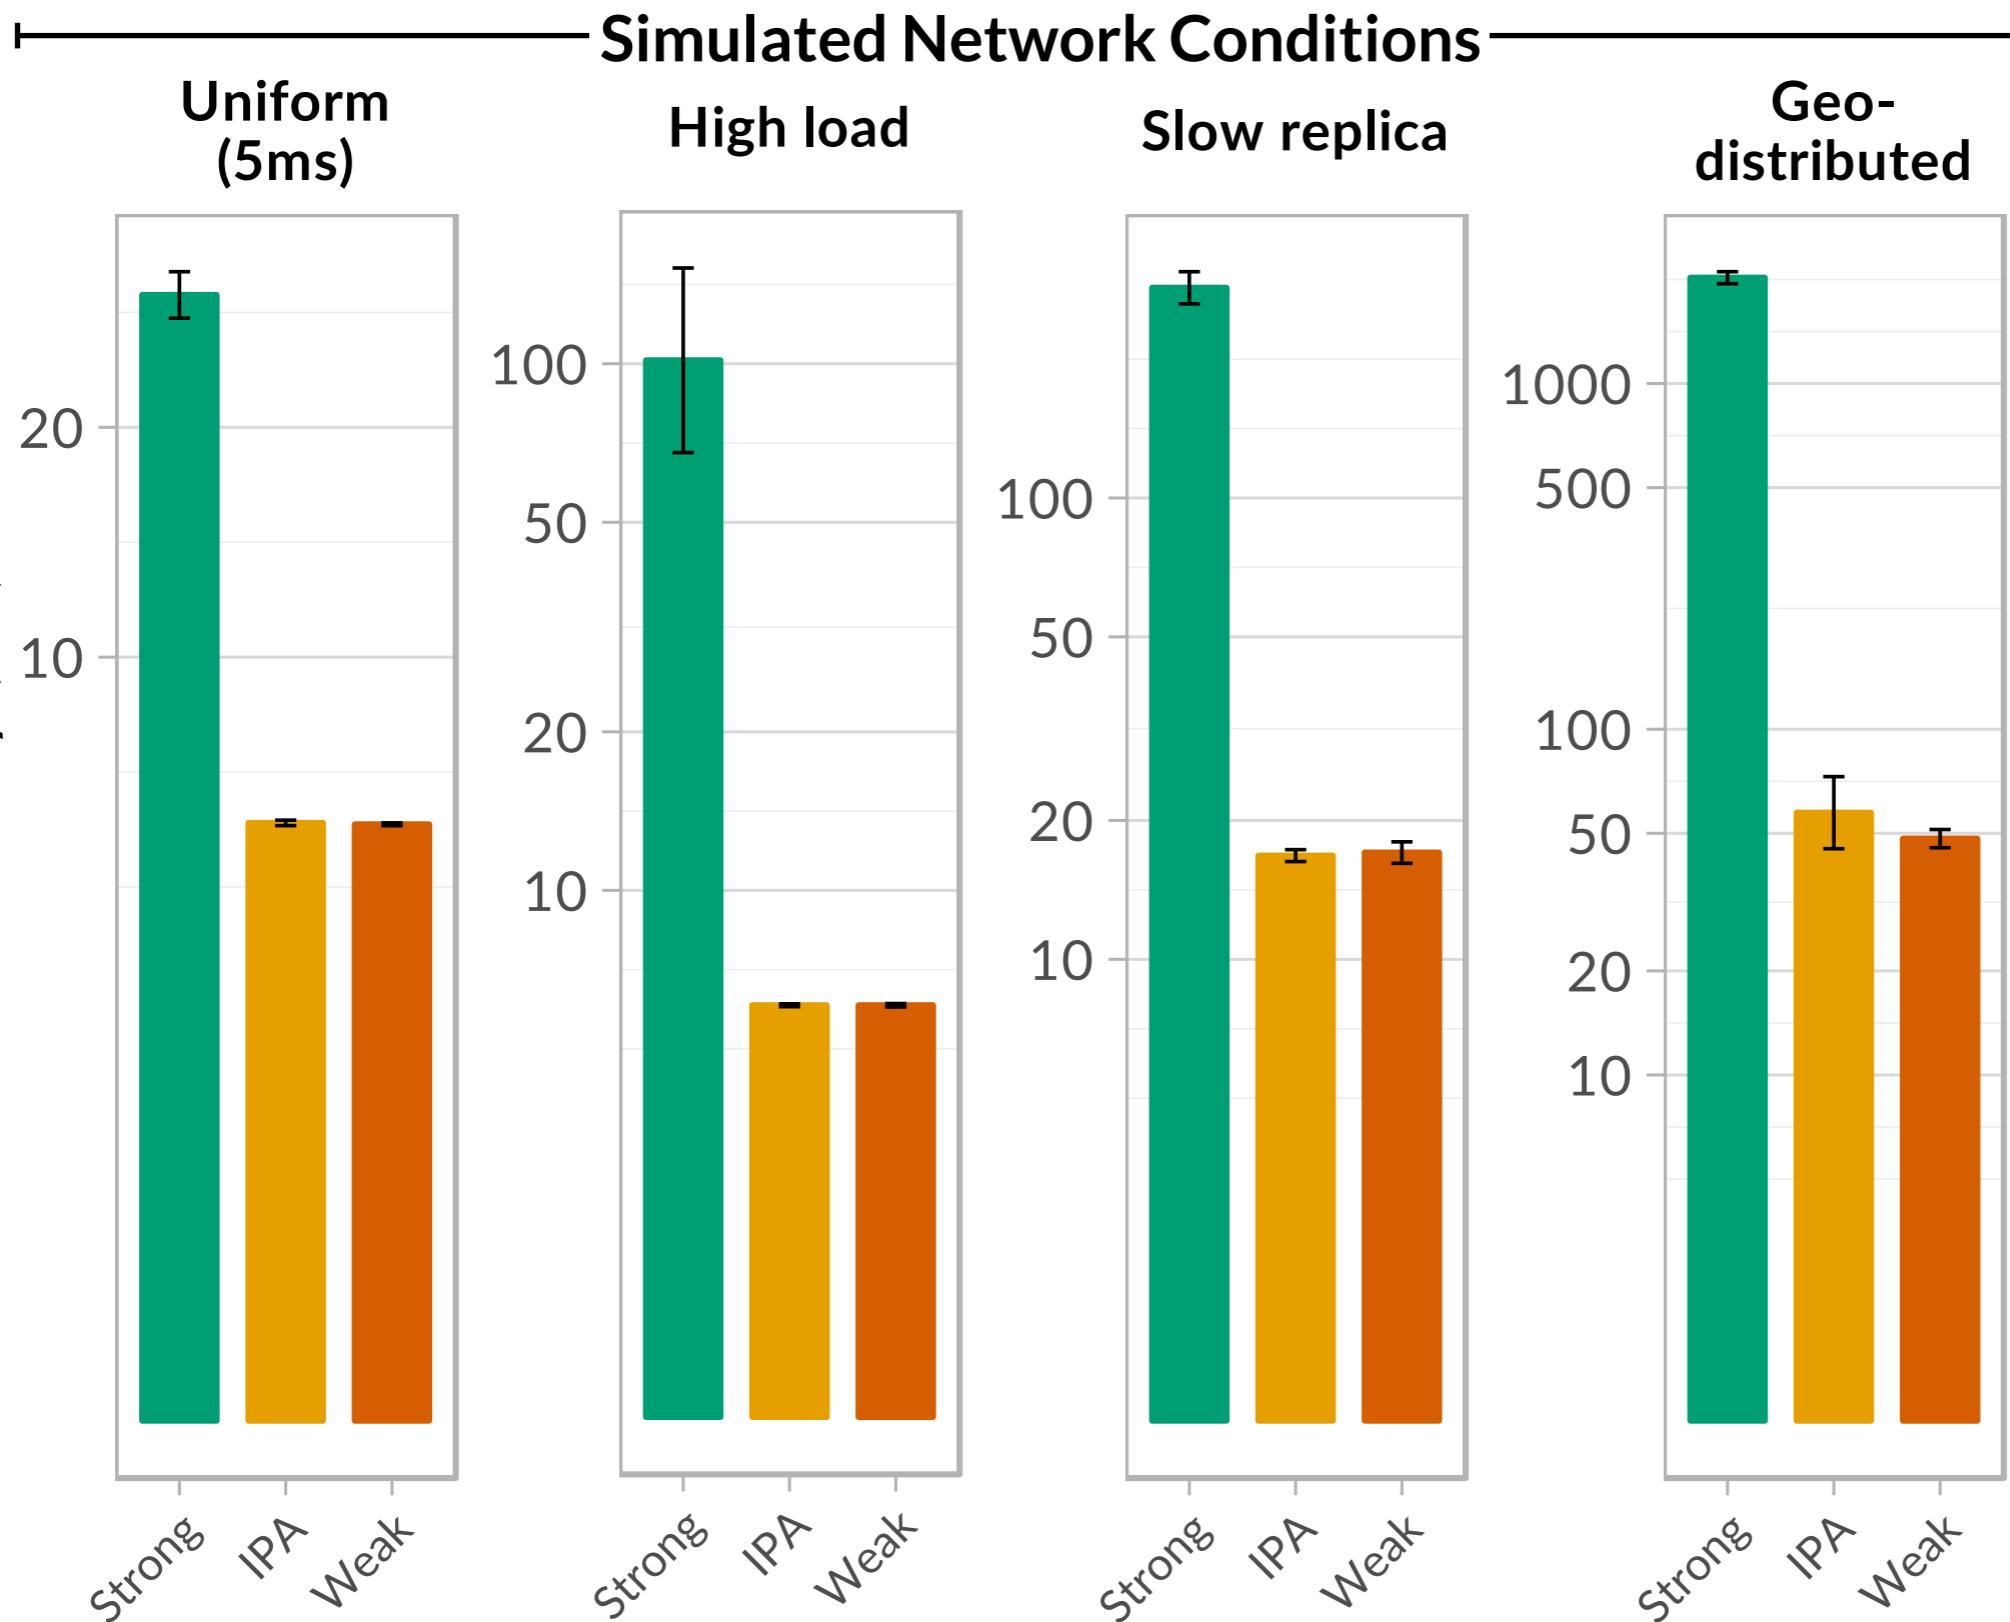


Error tolerance:
count does not need
to be exact...

Hard constraint:
do not over-sell tickets

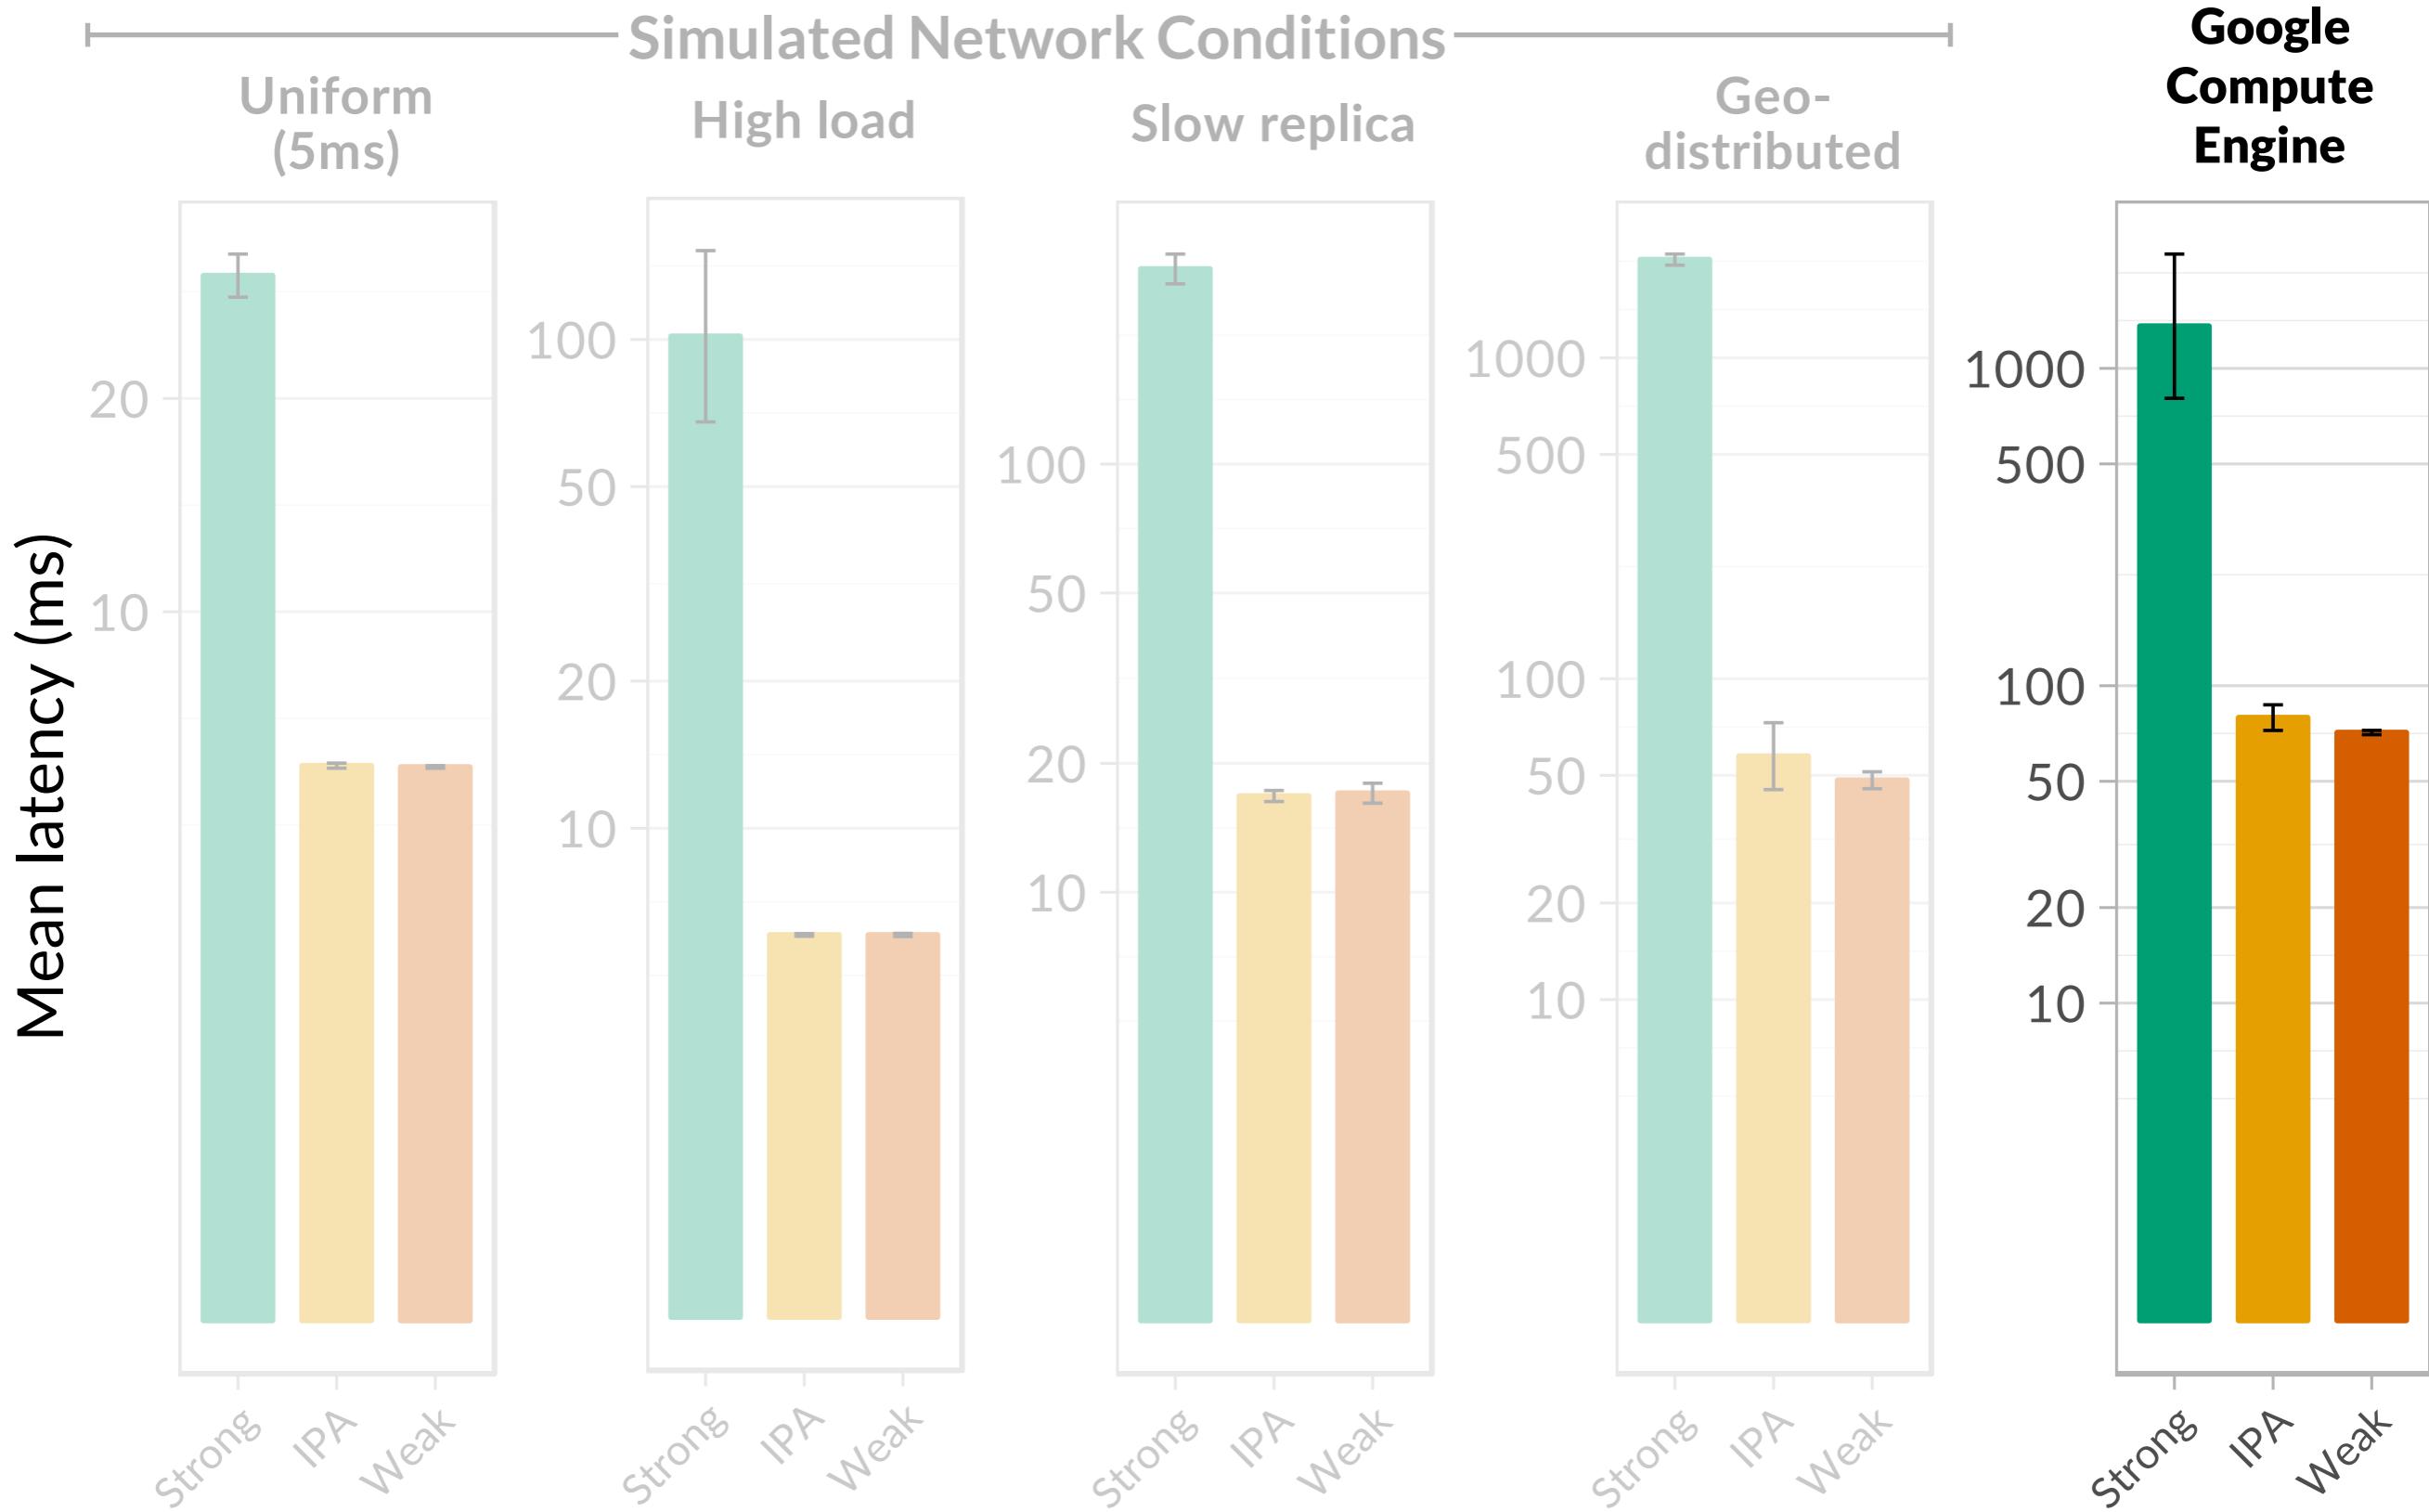


Case Study: TICKETSLEUTH



Case Study: TICKETSLEUTH

VMs in US,
Europe, and Asia





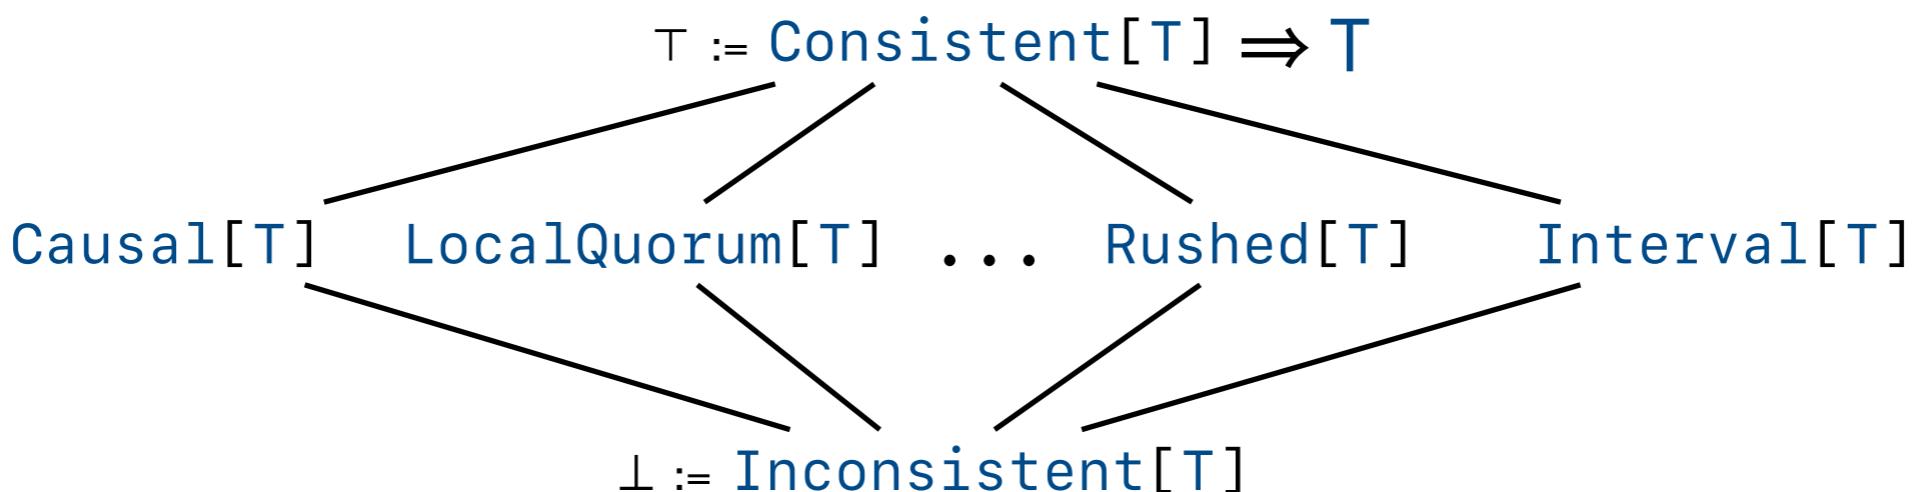
Disciplined Inconsistency

with *Consistency Types*

Inconsistent
Performance-bound
Approximate



Programming Model



Discussion

Mixed consistency systems

- Geo-distributed databases
- Multiple layers of caches and downstream services
- Example: Facebook's TAO graph store

Data-type- vs operation-centric consistency

- Why do we usually annotate individual operations with consistency rather than annotating data types?
- How does this relate to CRDTs?

Other uses for reservations

Verifying or synthesizing ADT/constraint implementations

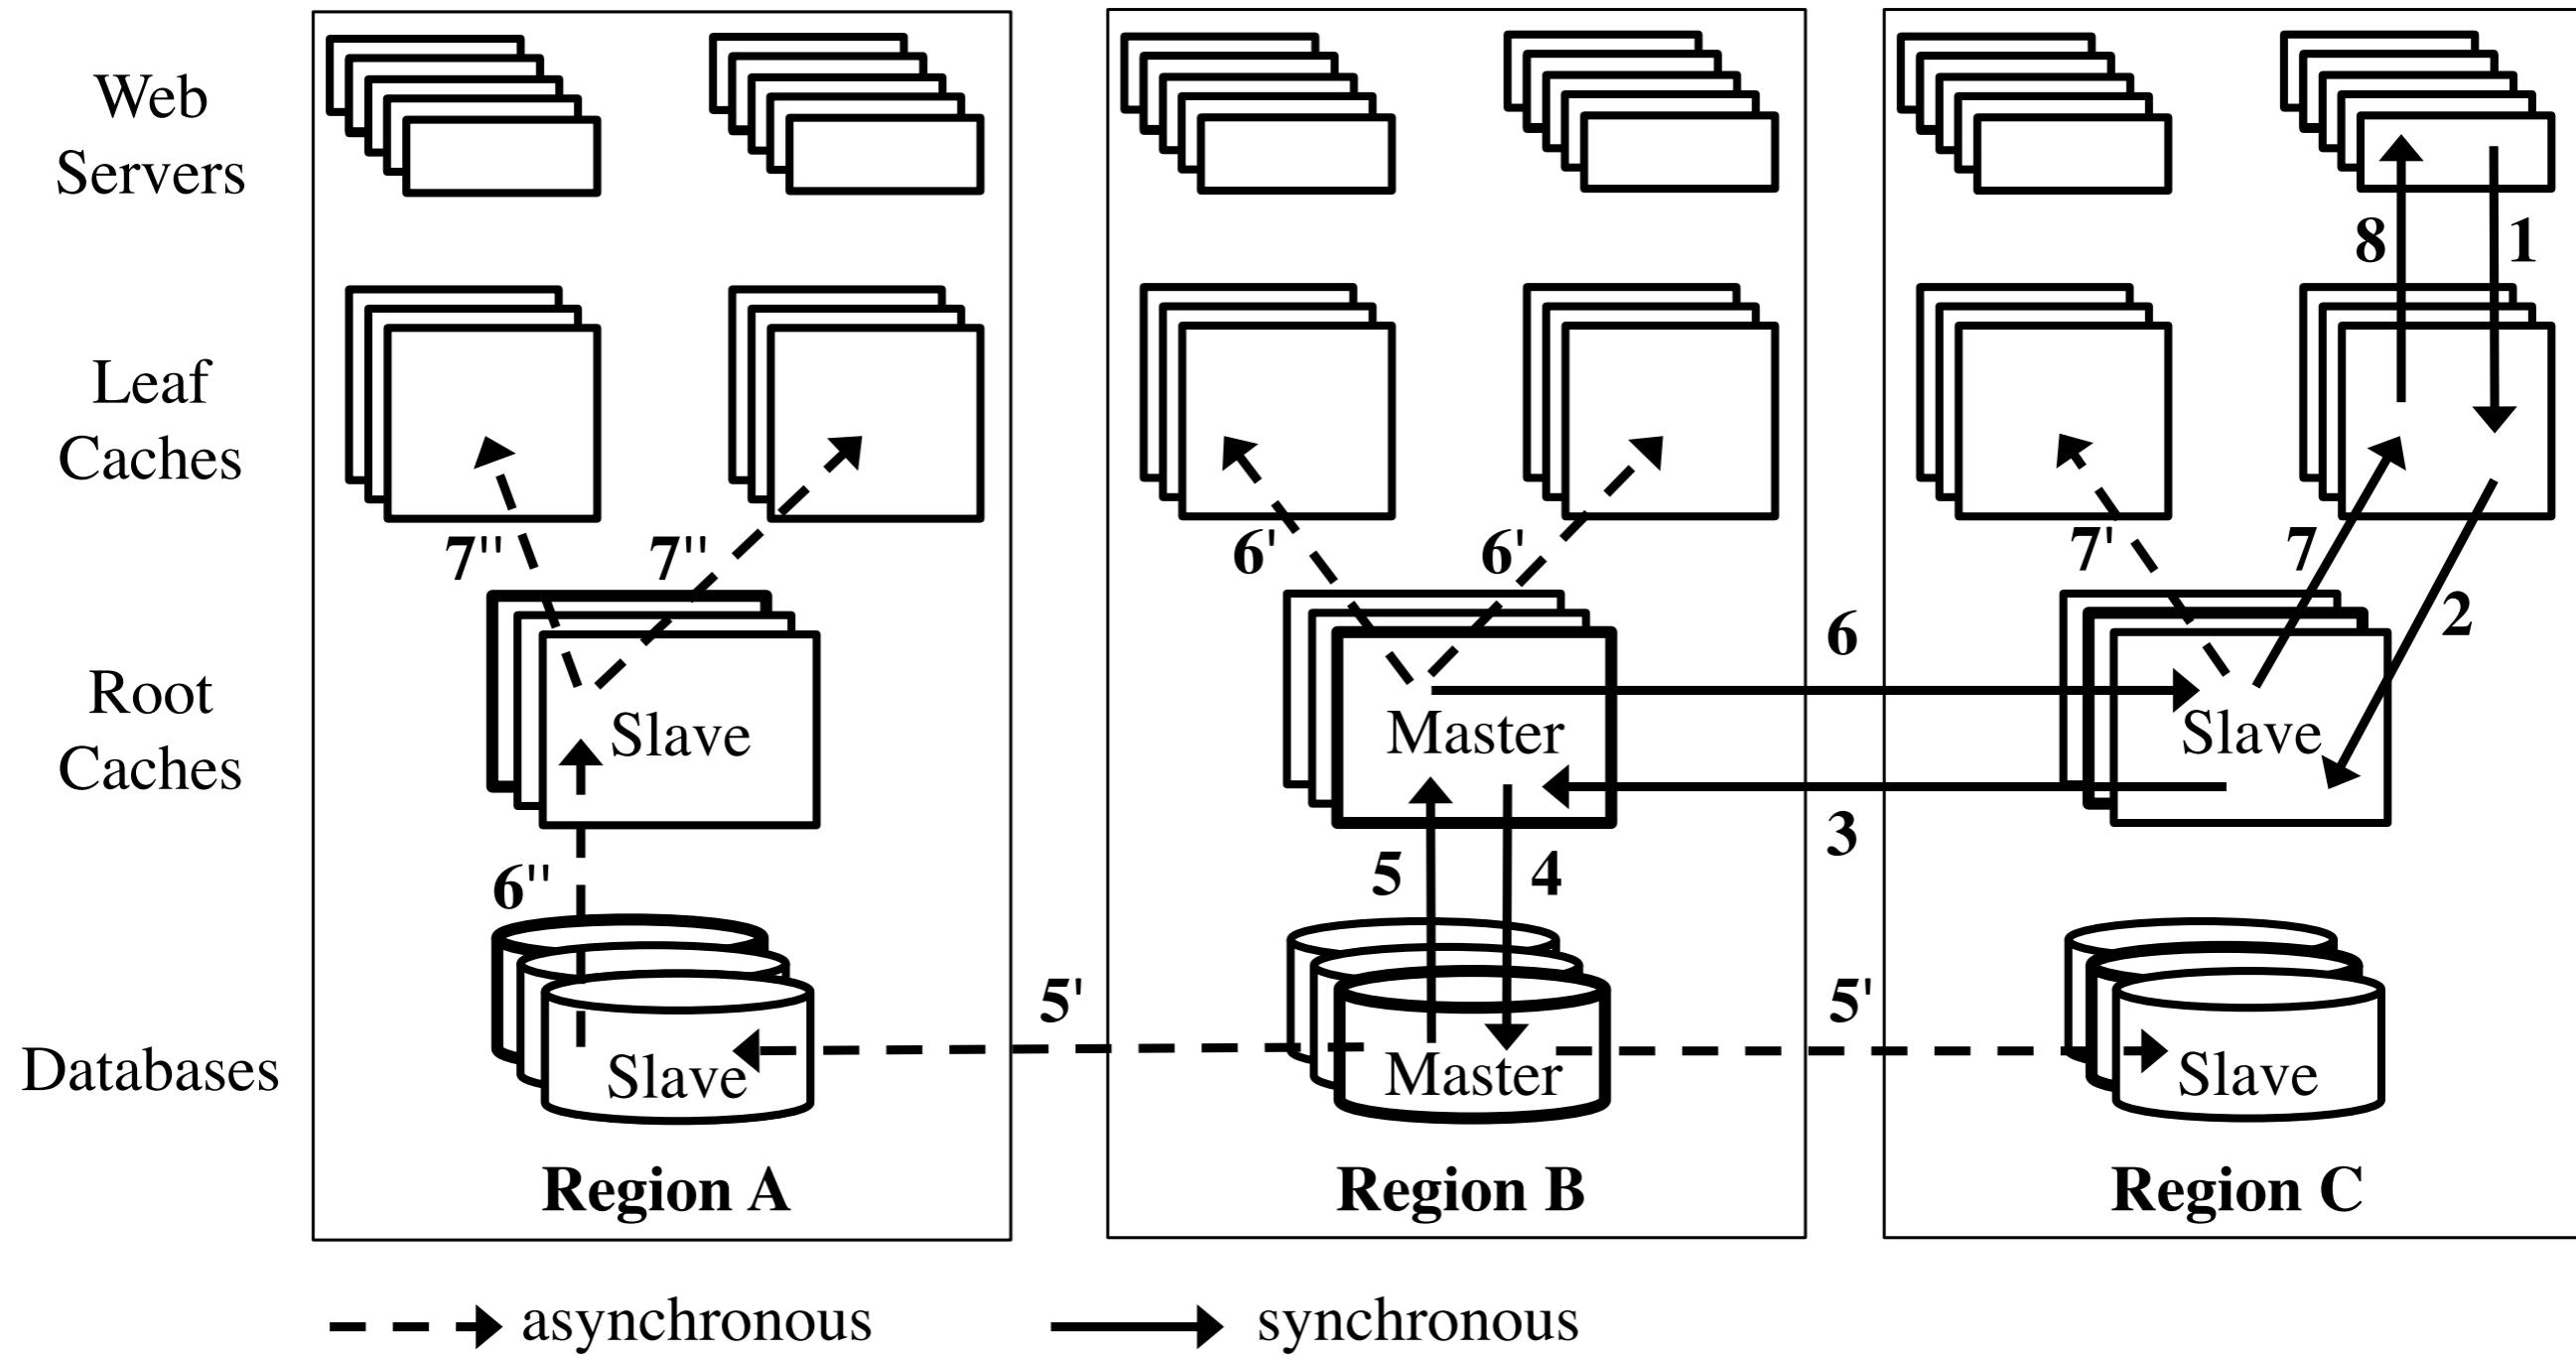
Composition of multiple operations or datatypes

Facebook TAO

Haonan Lu^{*†}, Kaushik Veeraraghavan[†], Philippe Ajoux[†], Jim Hunt[†],
Yee Jiun Song[†], Wendy Tobagus[†], Sanjeev Kumar[†], Wyatt Lloyd^{*†}

*University of Southern California, [†]Facebook, Inc.

(SOSP'15)



System architecture of TAO graph store.

Facebook TAO

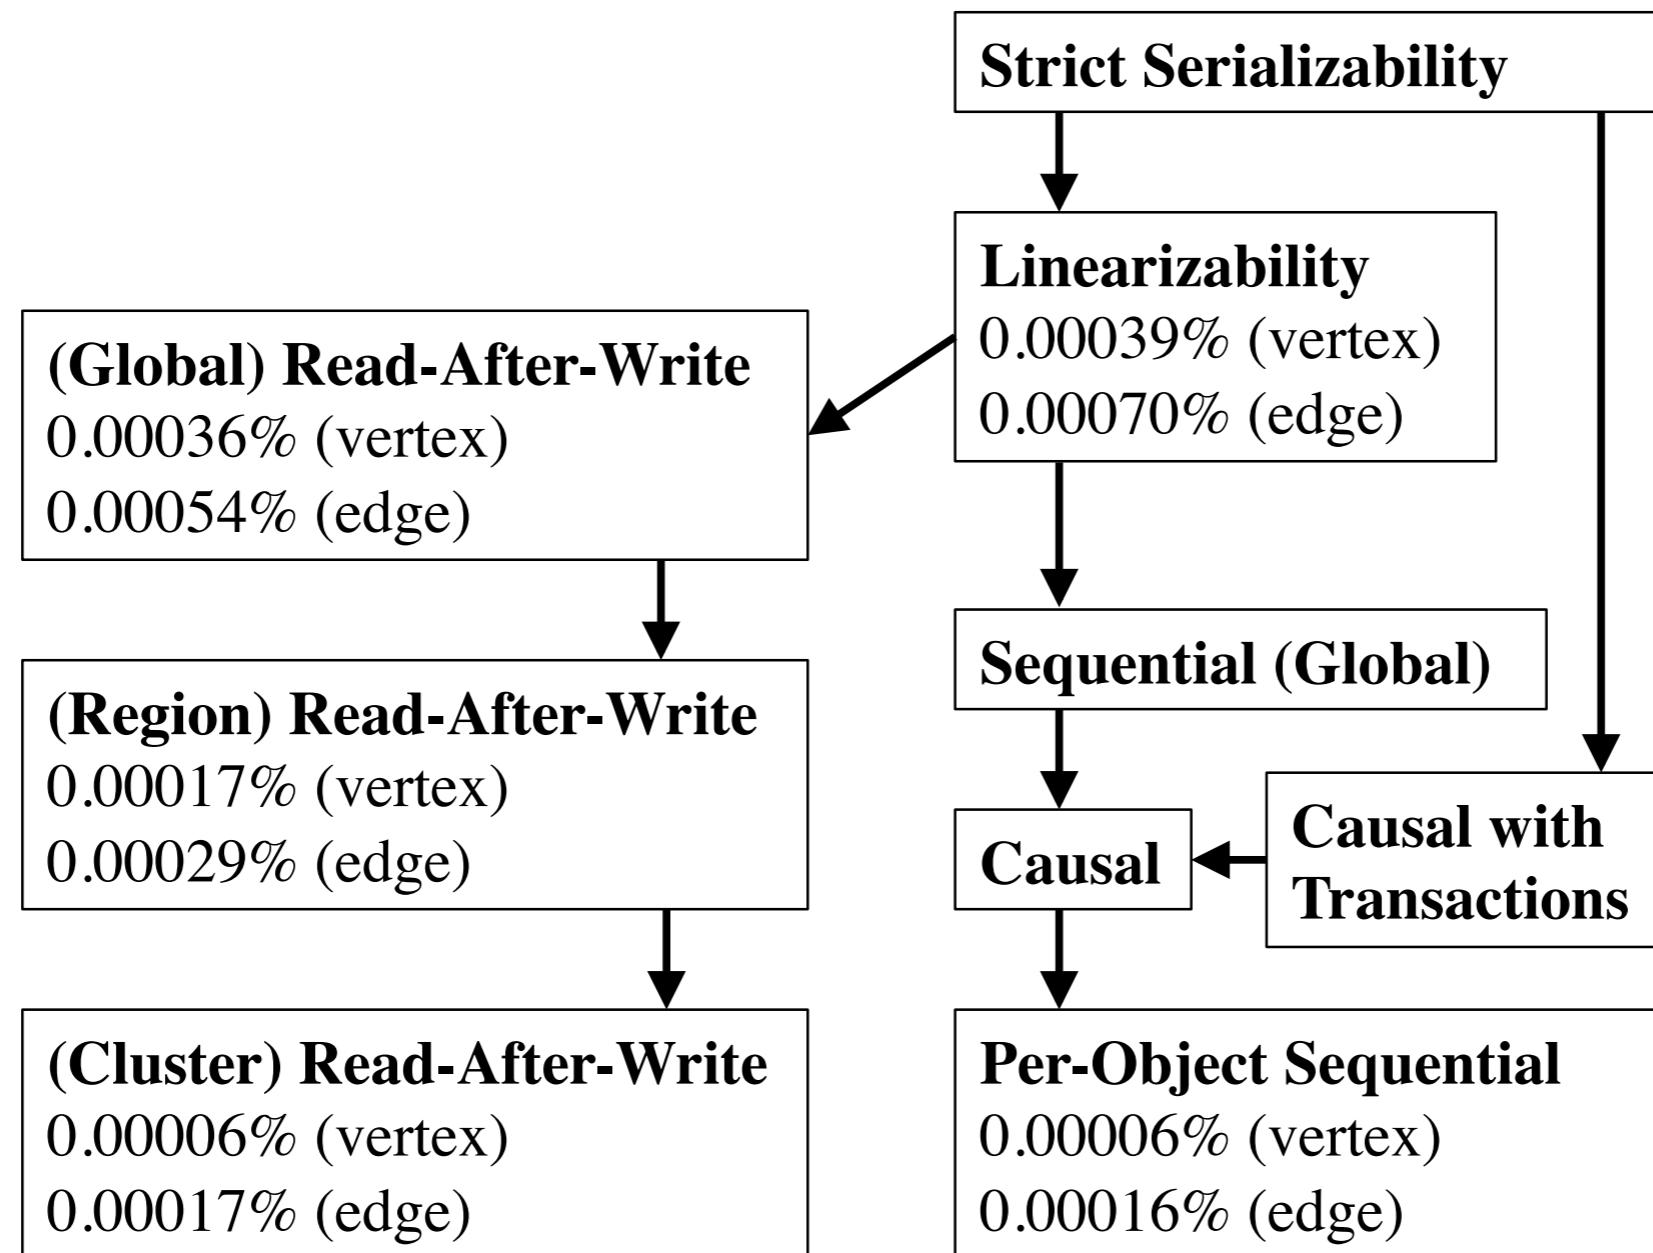
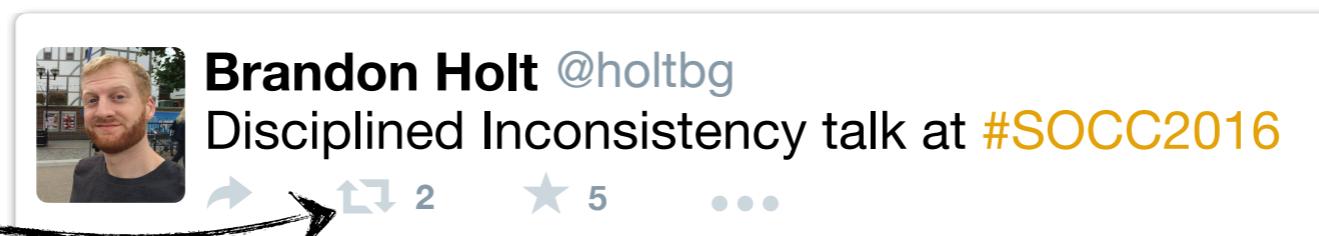


Figure 8: Spectrum of consistency models with arrows from stronger to weaker models. Measured anomaly rates are shown. Bounds on the anomaly rate in unmeasurable non-local consistency models can be inferred.

Case Study: Twitter clone

```
class Tweet(  
    id: TweetID,  
    user: UserID,  
    body: String,  
    retweets: Set[UserID] with ErrorTolerance(size, 5%)  
)
```

Exact count for average tweets.



Approximate count for popular tweets



Case Study: Twitter clone

```
class User(  
    id: UserID,  
    name: String,  
    followers: Set[UserID] with LatencyBound(20 ms),  
    timeline: List[TweetID] with LatencyBound(20 ms)  
)
```

 **Brandon Holt** @holtbg
Disciplined Inconsistency talk at @Zymergen!

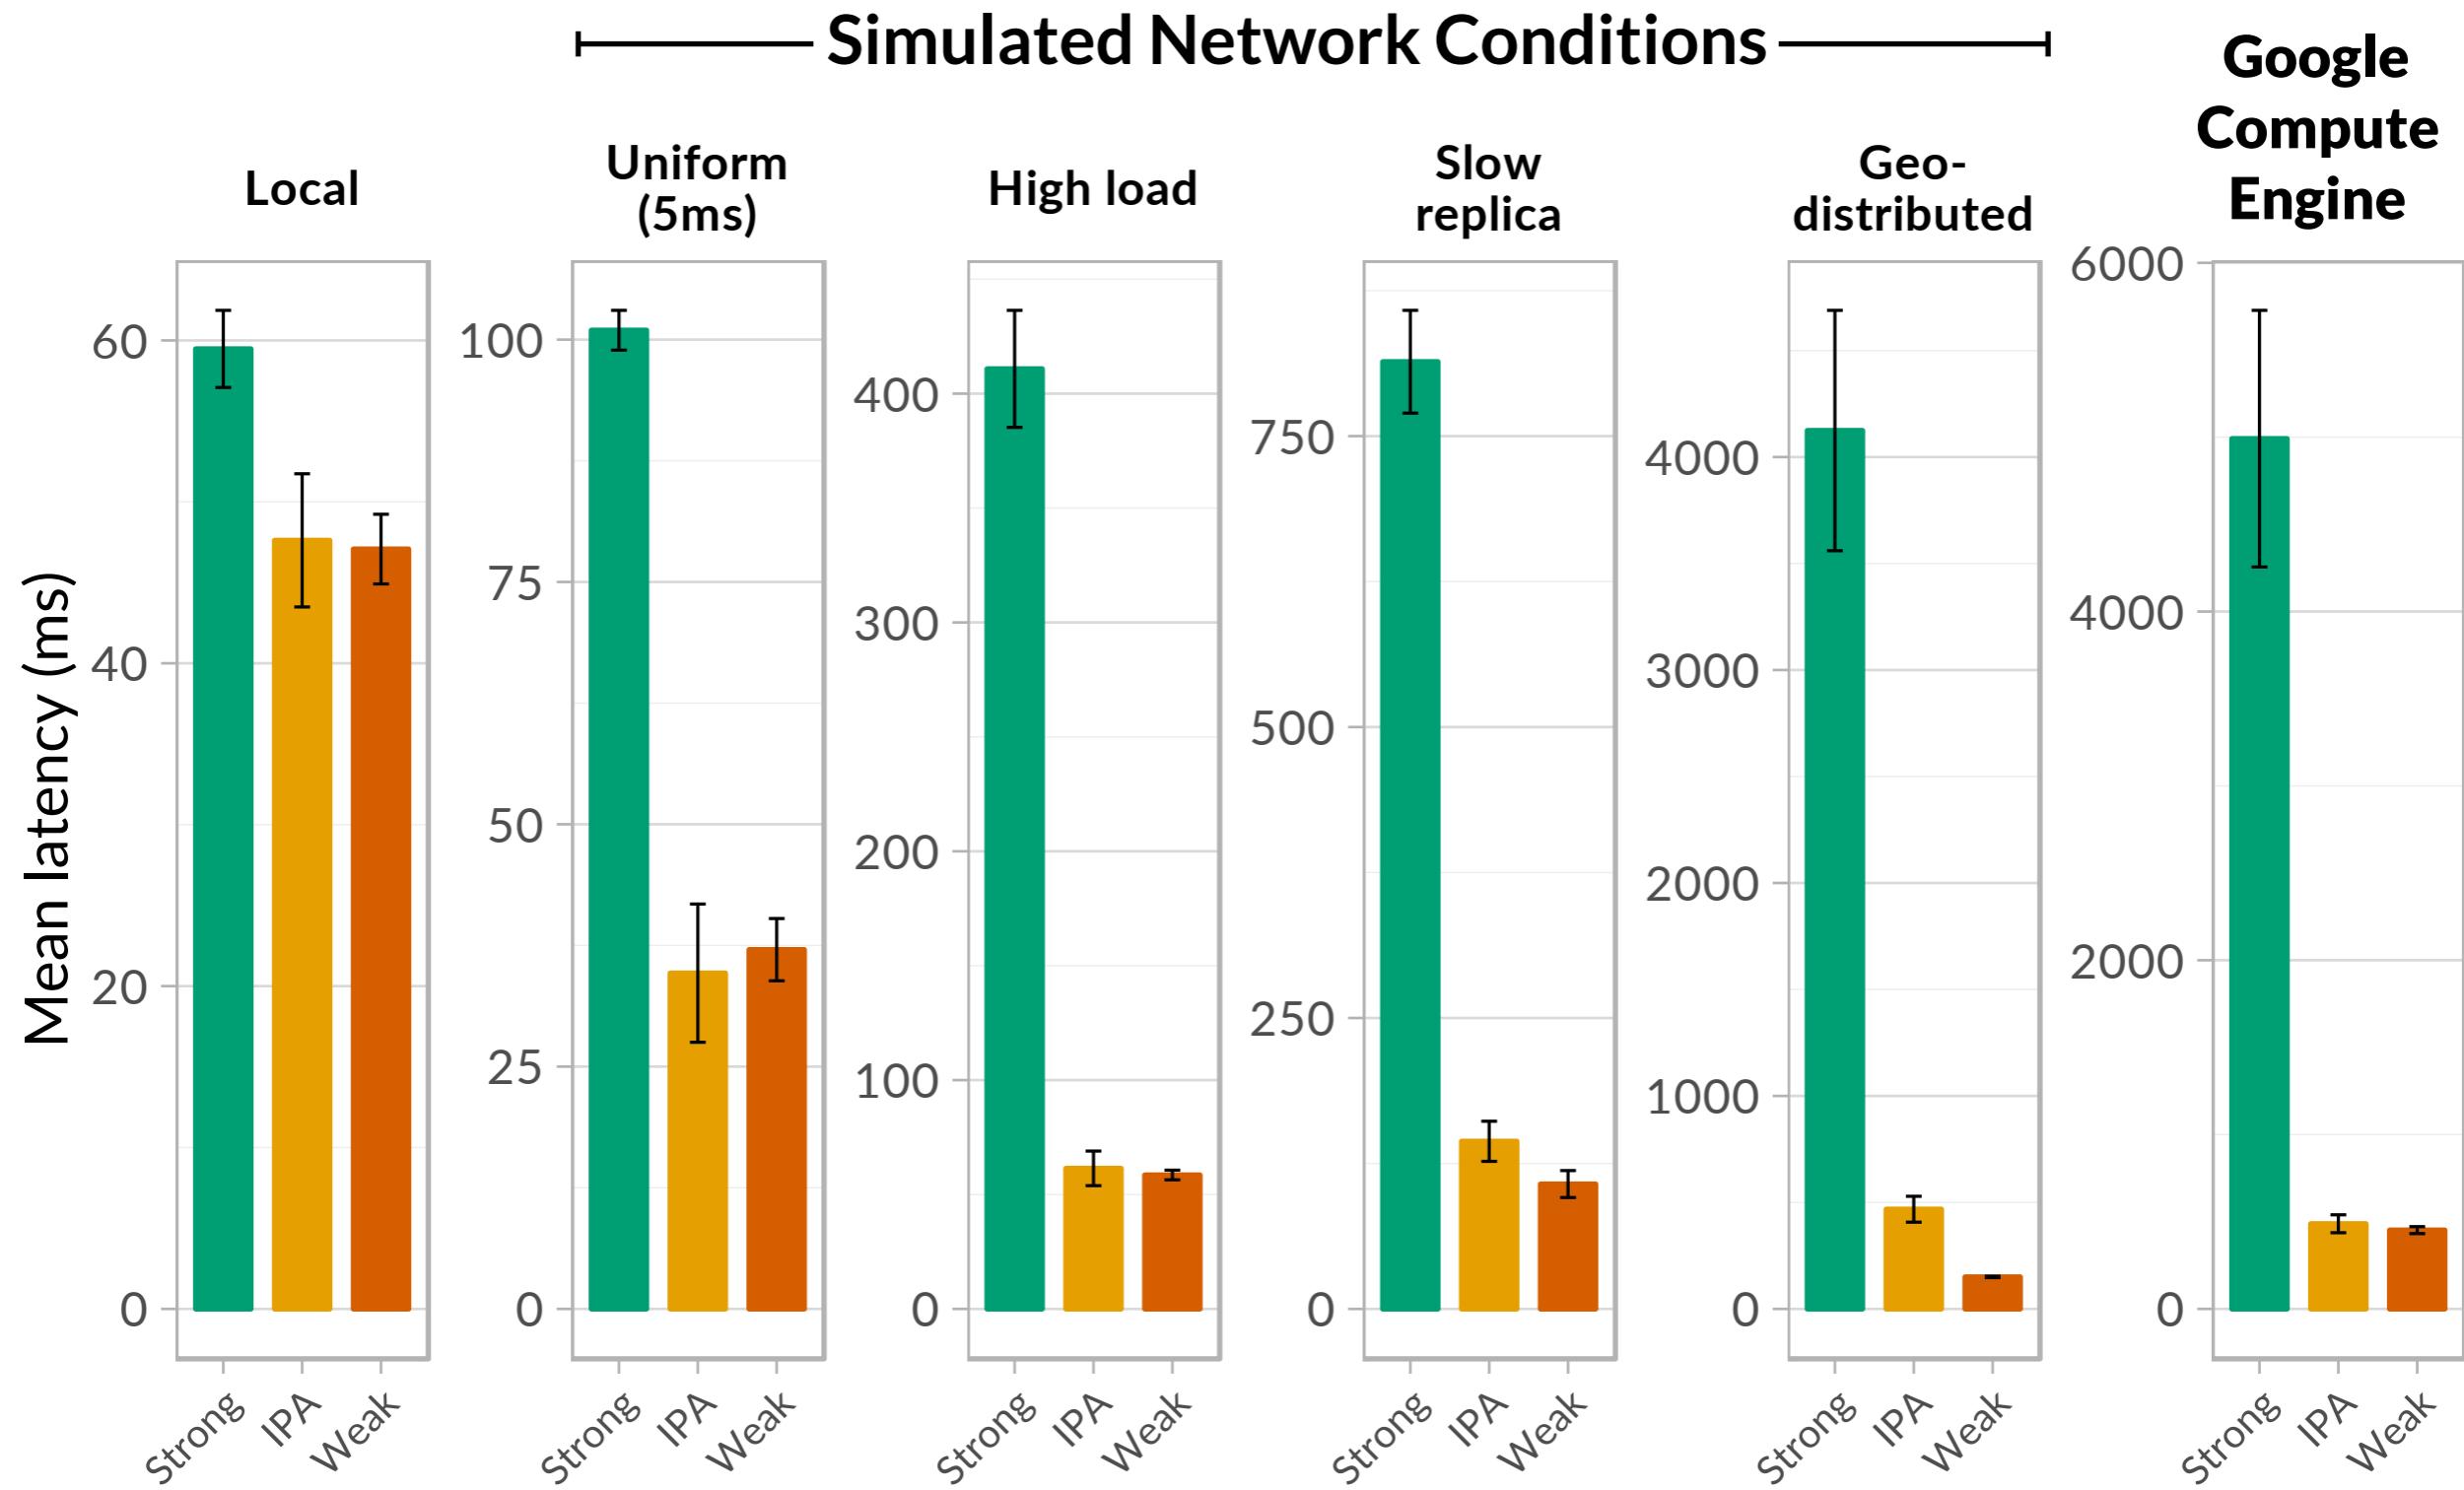
→ 2 ⚡ 5 ...

 **Ellen DeGeneres** @TheEllenShow
If only Bradley's arm was longer.
Best photo ever. #oscars

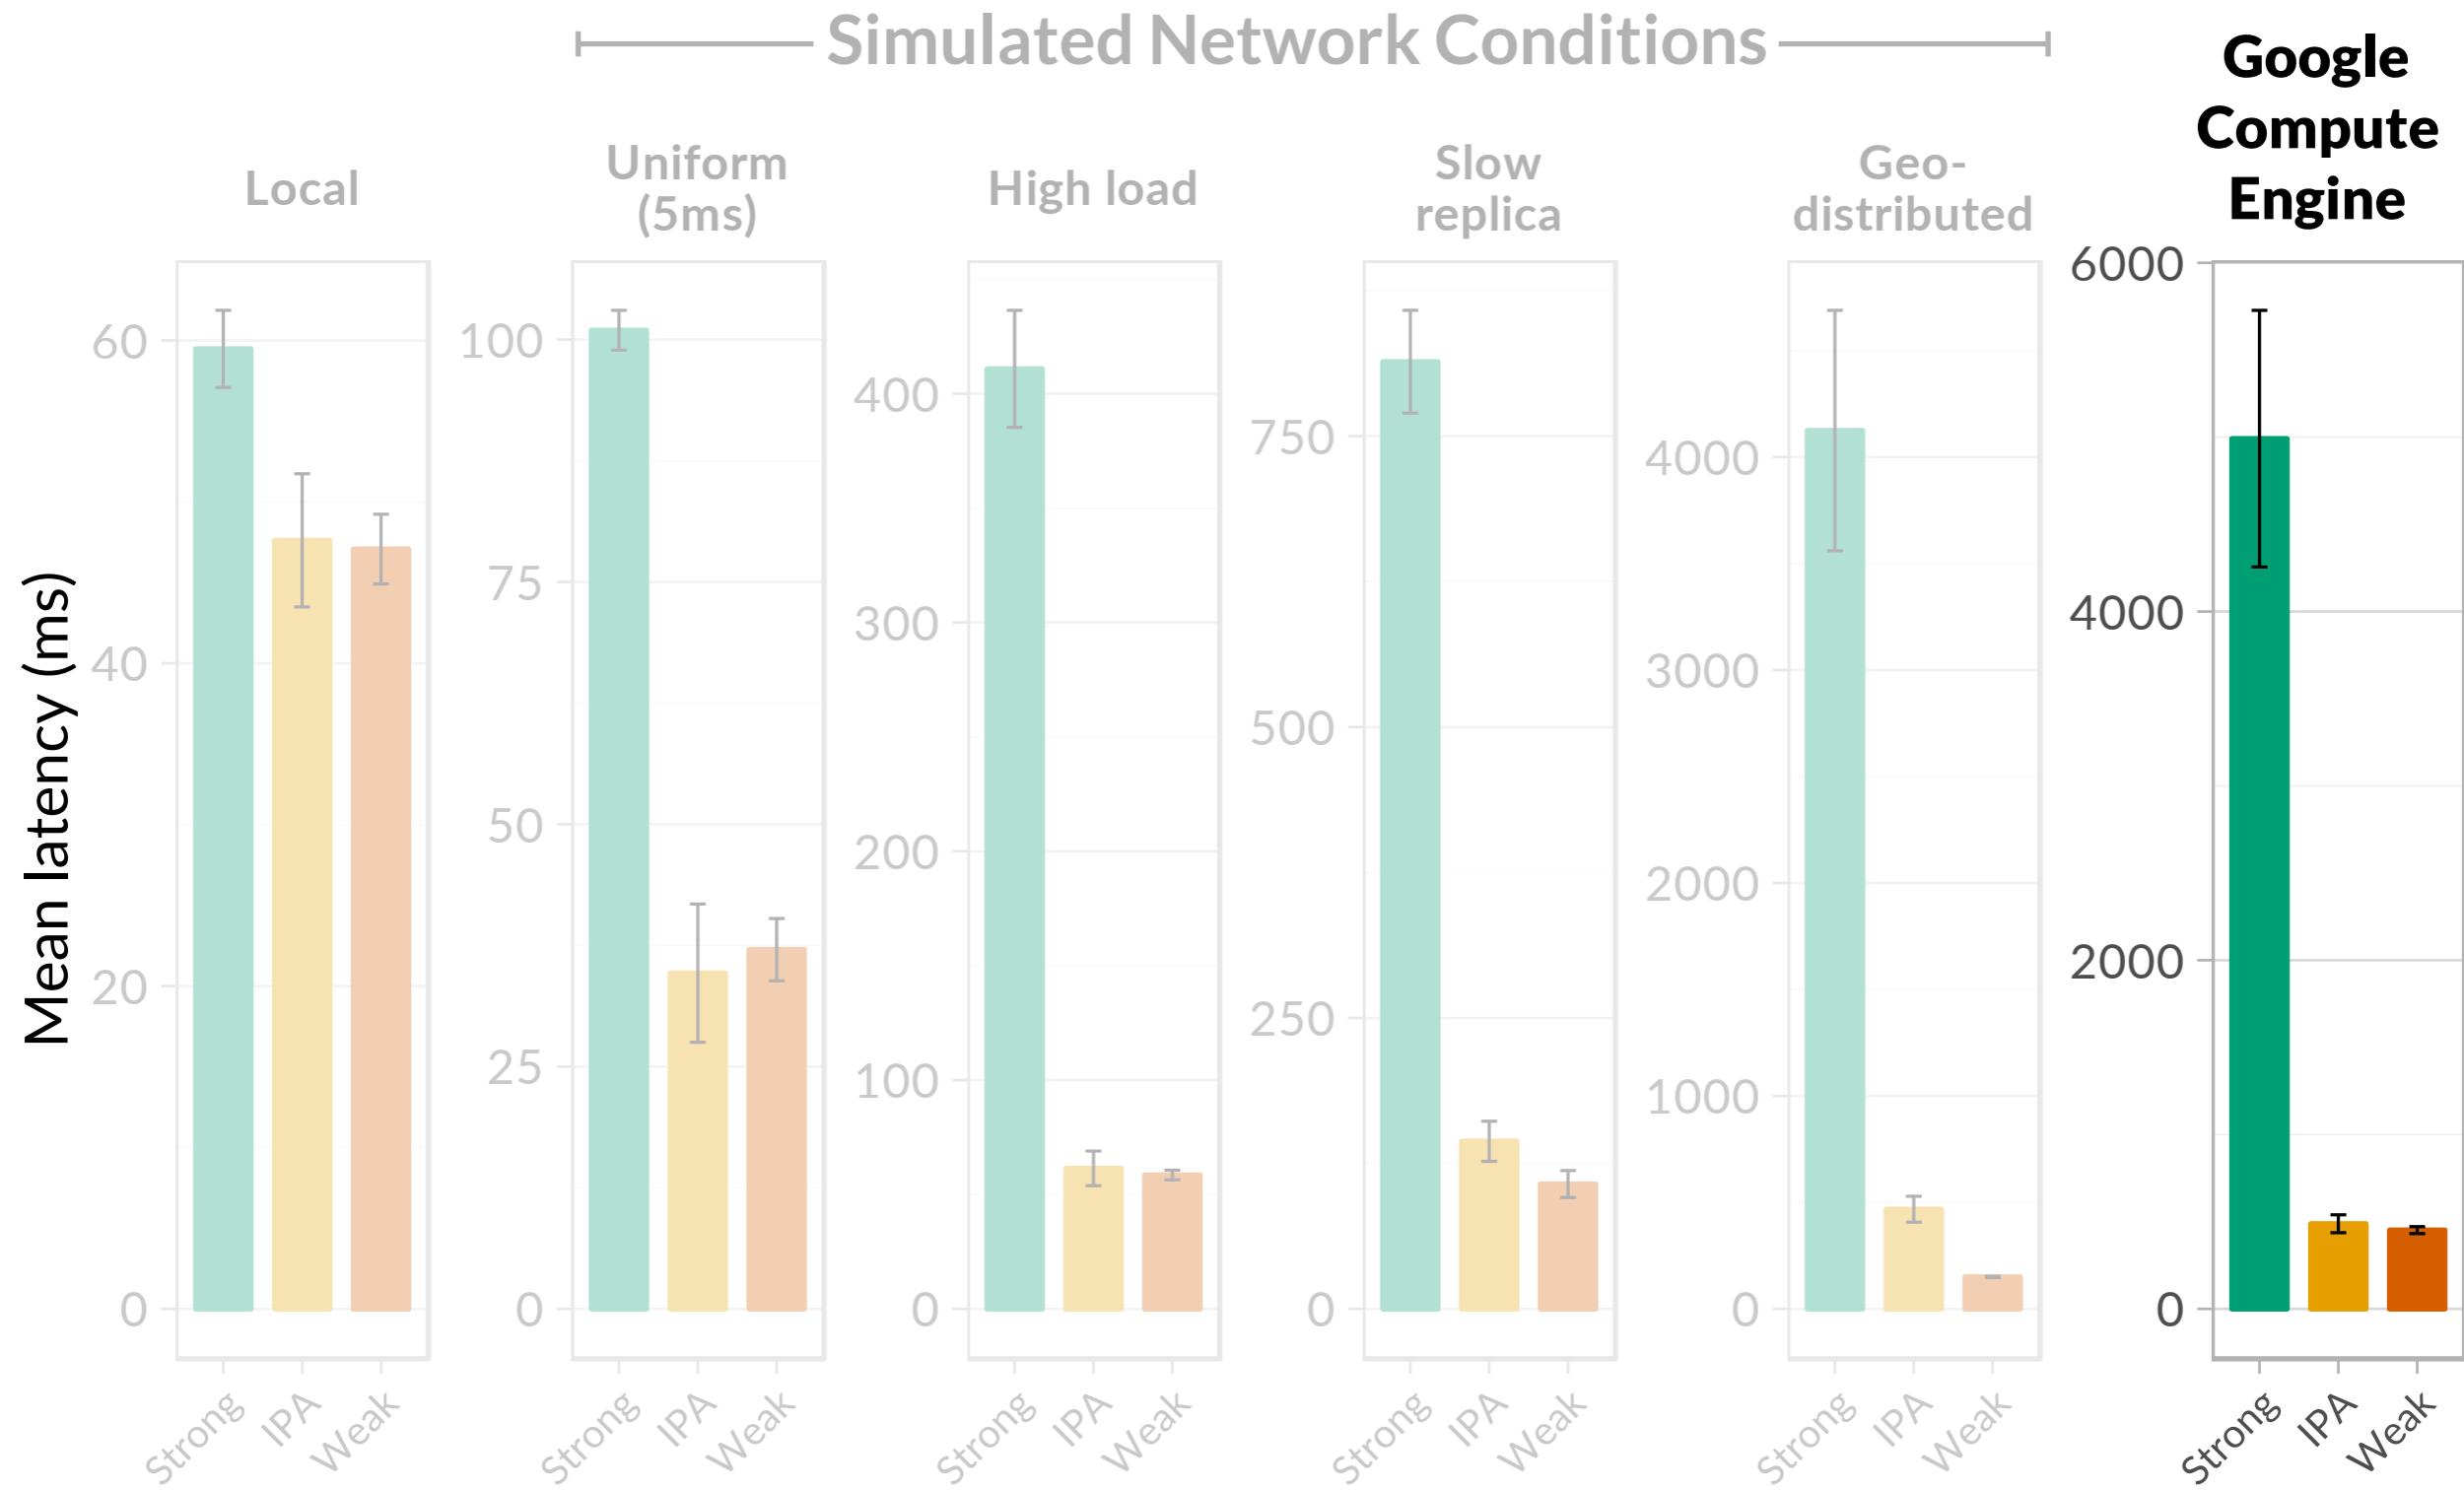


→ ⚡ 3.4M ⚡ 2M ...

Case Study: Twitter clone



Case Study: Twitter clone



Consistency-based SLAs

Douglas B. Terry, Vijayan Prabhakaran, Ramakrishna Kotla, Mahesh Balakrishnan, Marcos K. Aguilera, Hussam Abu-Libdeh

- SLAs specify consistency and target latency.
- Operations specify a set of SLAs with utility values.
- System monitors and predicts which SLA can be met, tries to maximize utility.
- Speculative execution

Consistency Levels

Strong

Eventual

Read-My-Writes

Monotonic

Bounded

Causal

```

CreateTable (name);
DeleteTable (name);
tbl = OpenTable (name);

s = BeginSession(tbl, sla);
EndSession(s);

Put (s, key, value);
value, cc = Get (s, key, sla);

```

Consistency	Latency	Utility
-------------	---------	---------

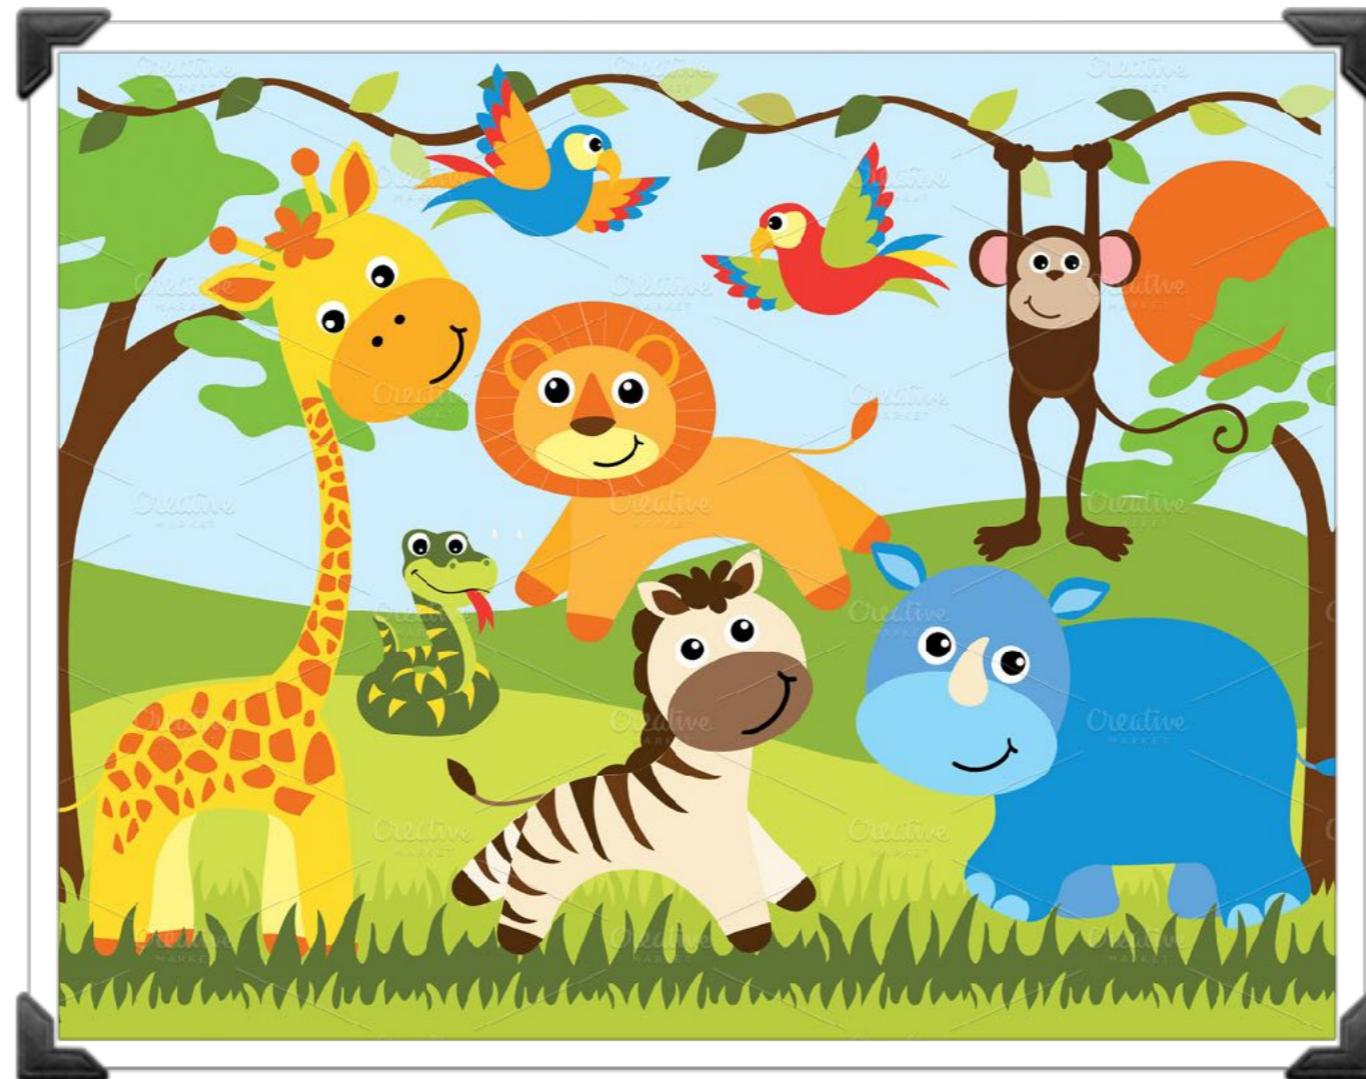
Strong	300 ms	1.0
--------	--------	-----

Eventual	300 ms	0.5
----------	--------	-----

Example: Shopping Cart

Figure 2. The Pileus API

A Romp Through the Consistency Zoo



by Brandon Holt



The **Strongest**

The **Safest**

The (least) Freedom

Strong Serializable

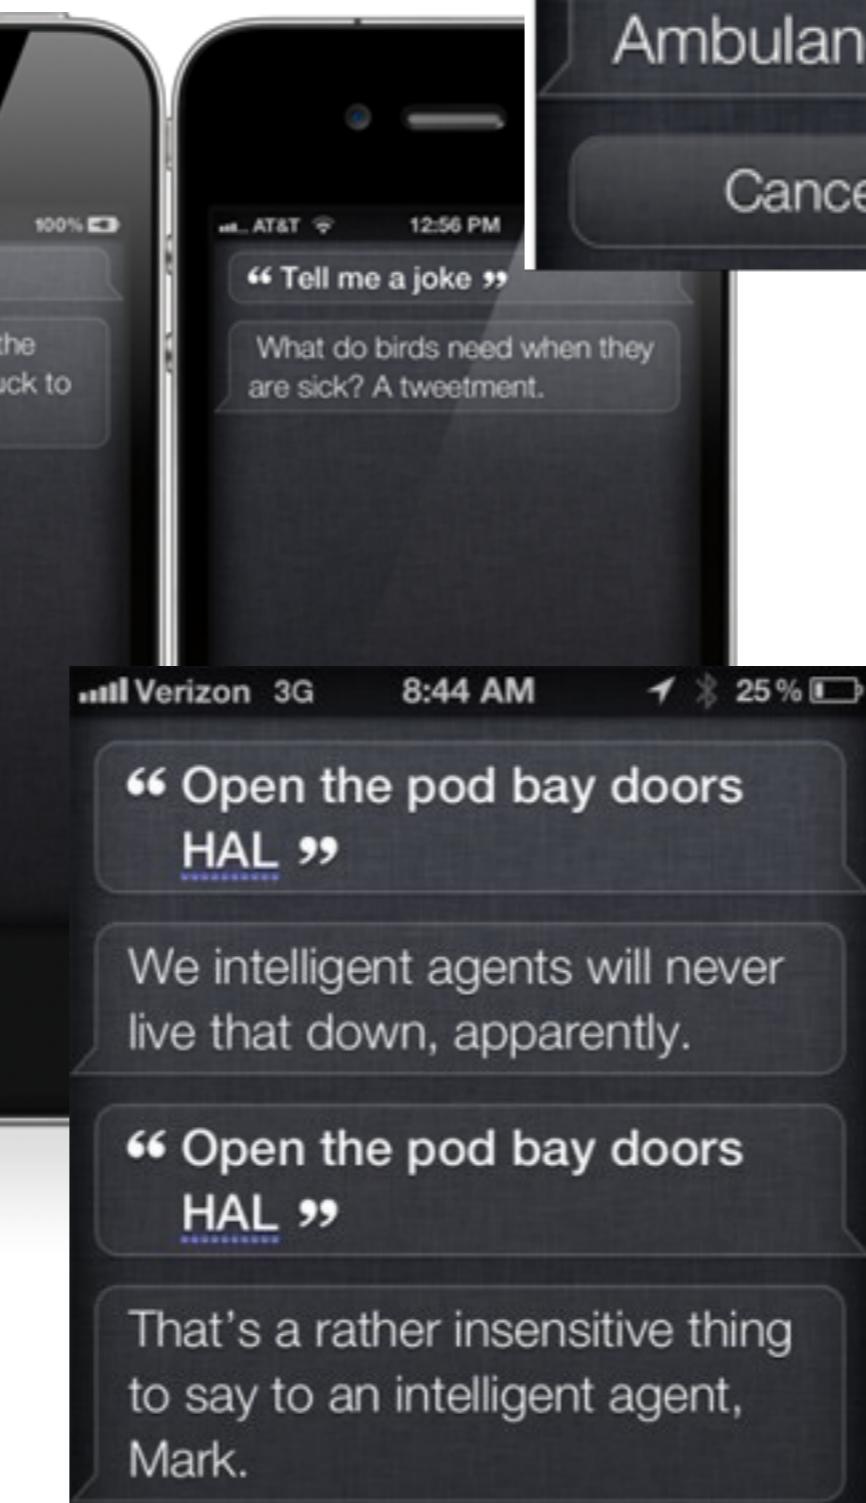
~~Linearizable~~
^ΛLionizable



~~Linearizable~~ Lionizable



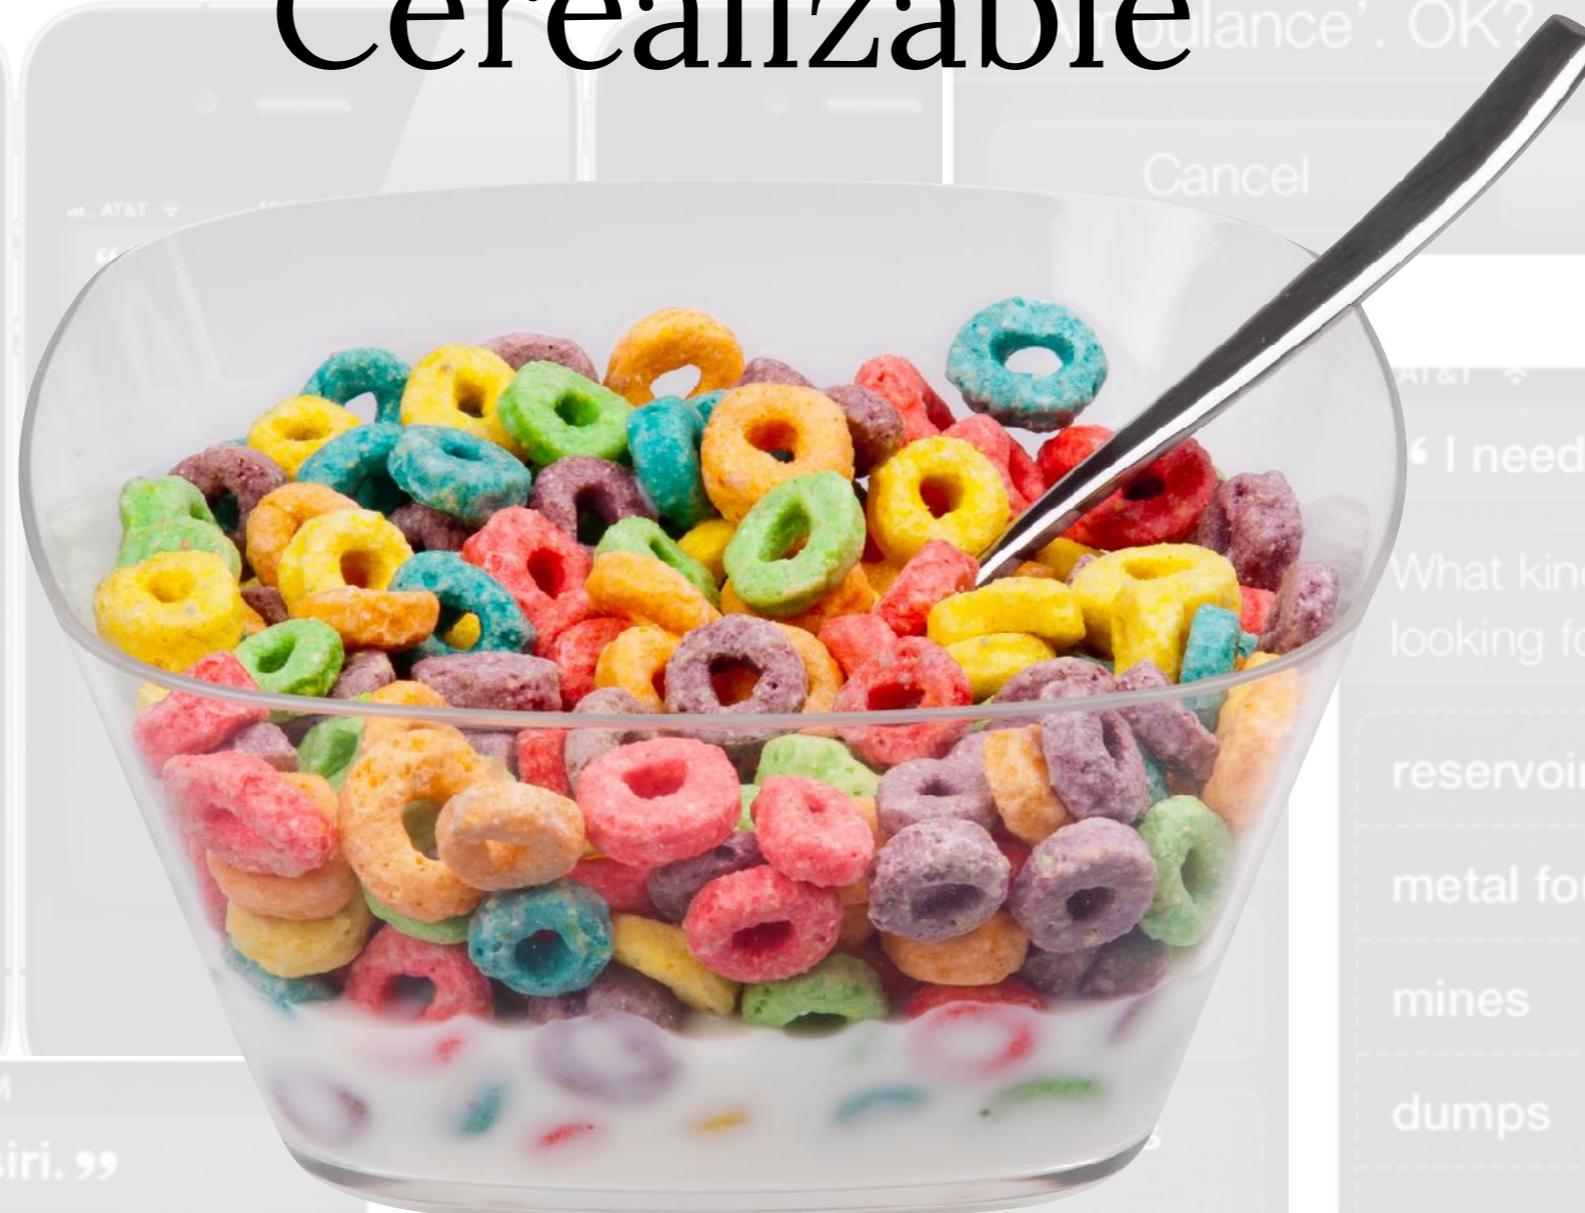
~~Serializable~~ Siri



~~Serializable~~ Siri

Bet you thought I was gonna say:

Cerealizable



Your contact list.

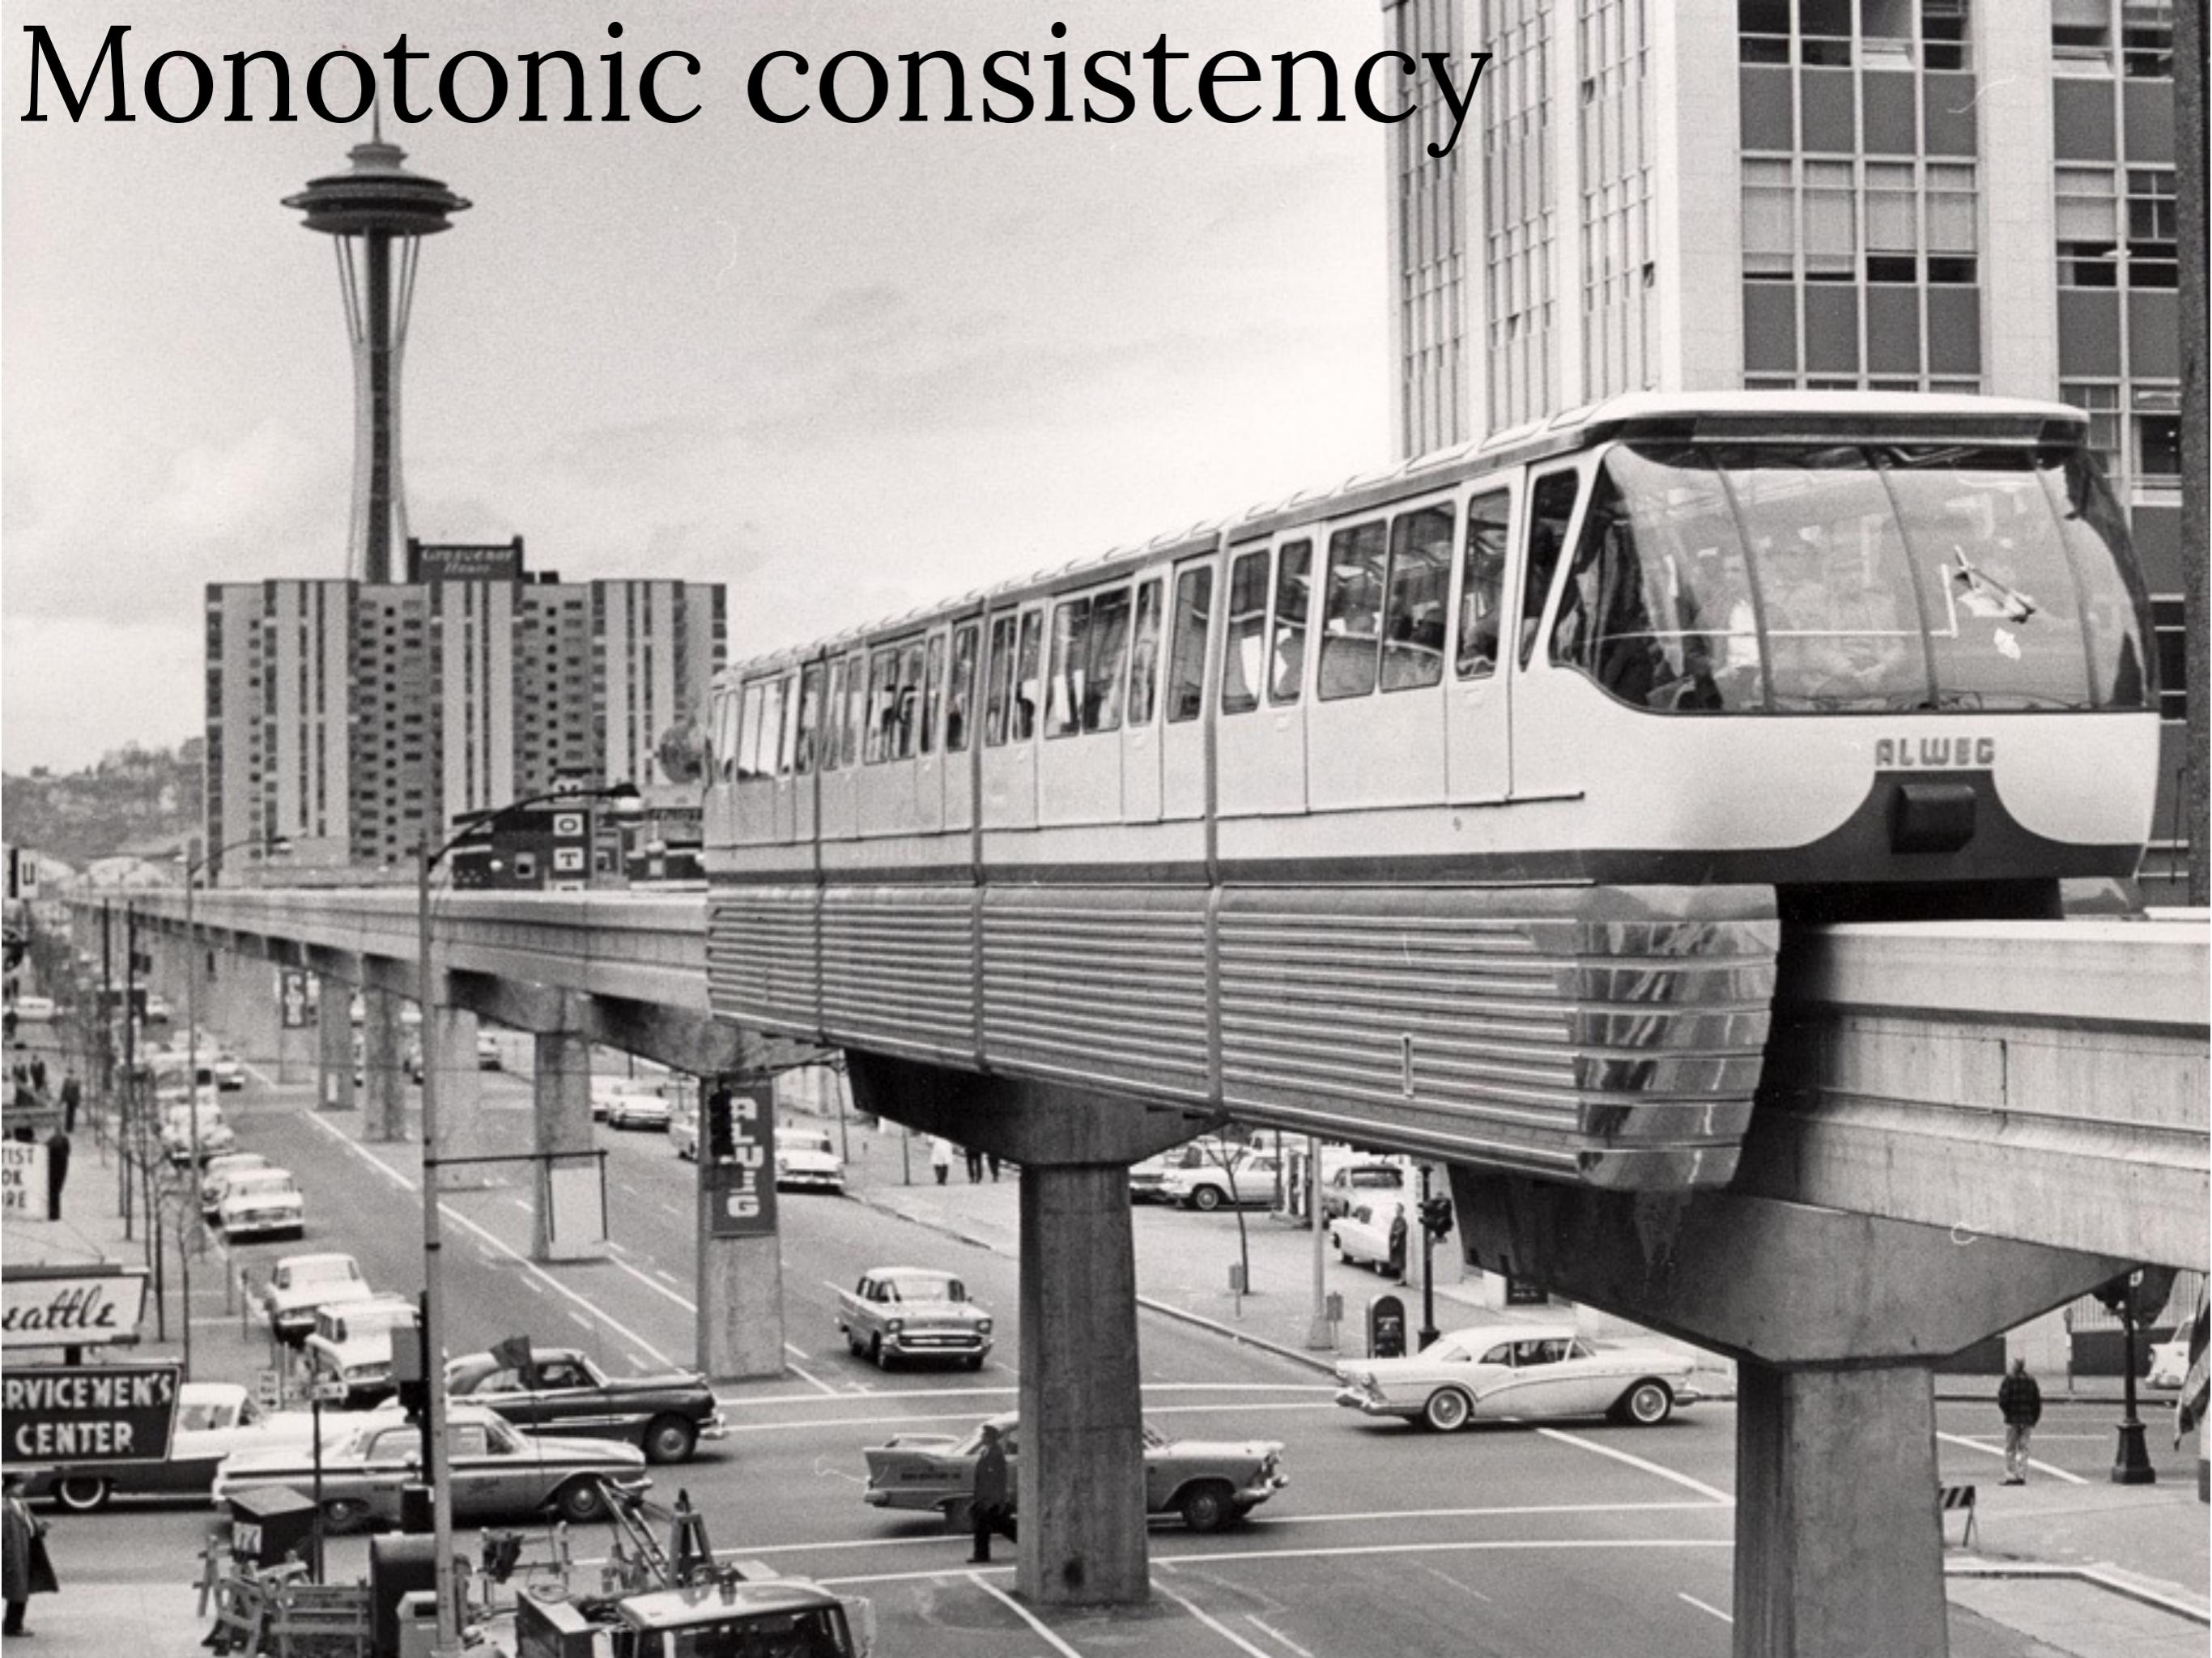
That's a rather insensitive thing
to say to an intelligent agent,
Mark.

Read-Your-Writes Rights



© 1989 Universal Press Syndicate

Monotonic consistency



PRAM



~~Causal~~ Consistency

~~Causal~~



Causal Consistency

DO YOU EVEN KNOW



HOW VECTOR CROCKS WORK

